

---

# Dynamic Privilege

---

A Thesis Defense by Thomas Unger

---

# Overview

- ❖ Ch 1: Rigidity Prevents Optimization
- ❖ Ch 2: Introducing Dynamic Privilege
- ❖ Ch 3: Demonstrating Low-Level and OS-Level Access
- ❖ Ch 4: Case Studies
- ❖ Conclusion



There's a lot more to say here

---

# Overview

---

- ❖ Ch 1: Rigidity Prevents Optimization
- ❖ Ch 2: Introducing Dynamic Privilege
- ❖ Ch 3: Demonstrating Low-Level and OS-Level Access
- ❖ Ch 4: Case Studies
- ❖ Conclusion

---

# Ch 1: Rigidity Prevents Optimization

---

- ❖ General Purpose Operating Systems (GPOSs) rule the world
- ❖ Specialized OSs can outperform GPOSs by making **trade-offs**
- ❖ Can we bring some flexibility GPOS?
  - ❖ By enabling application specialization

# Should We Mind the Gap?

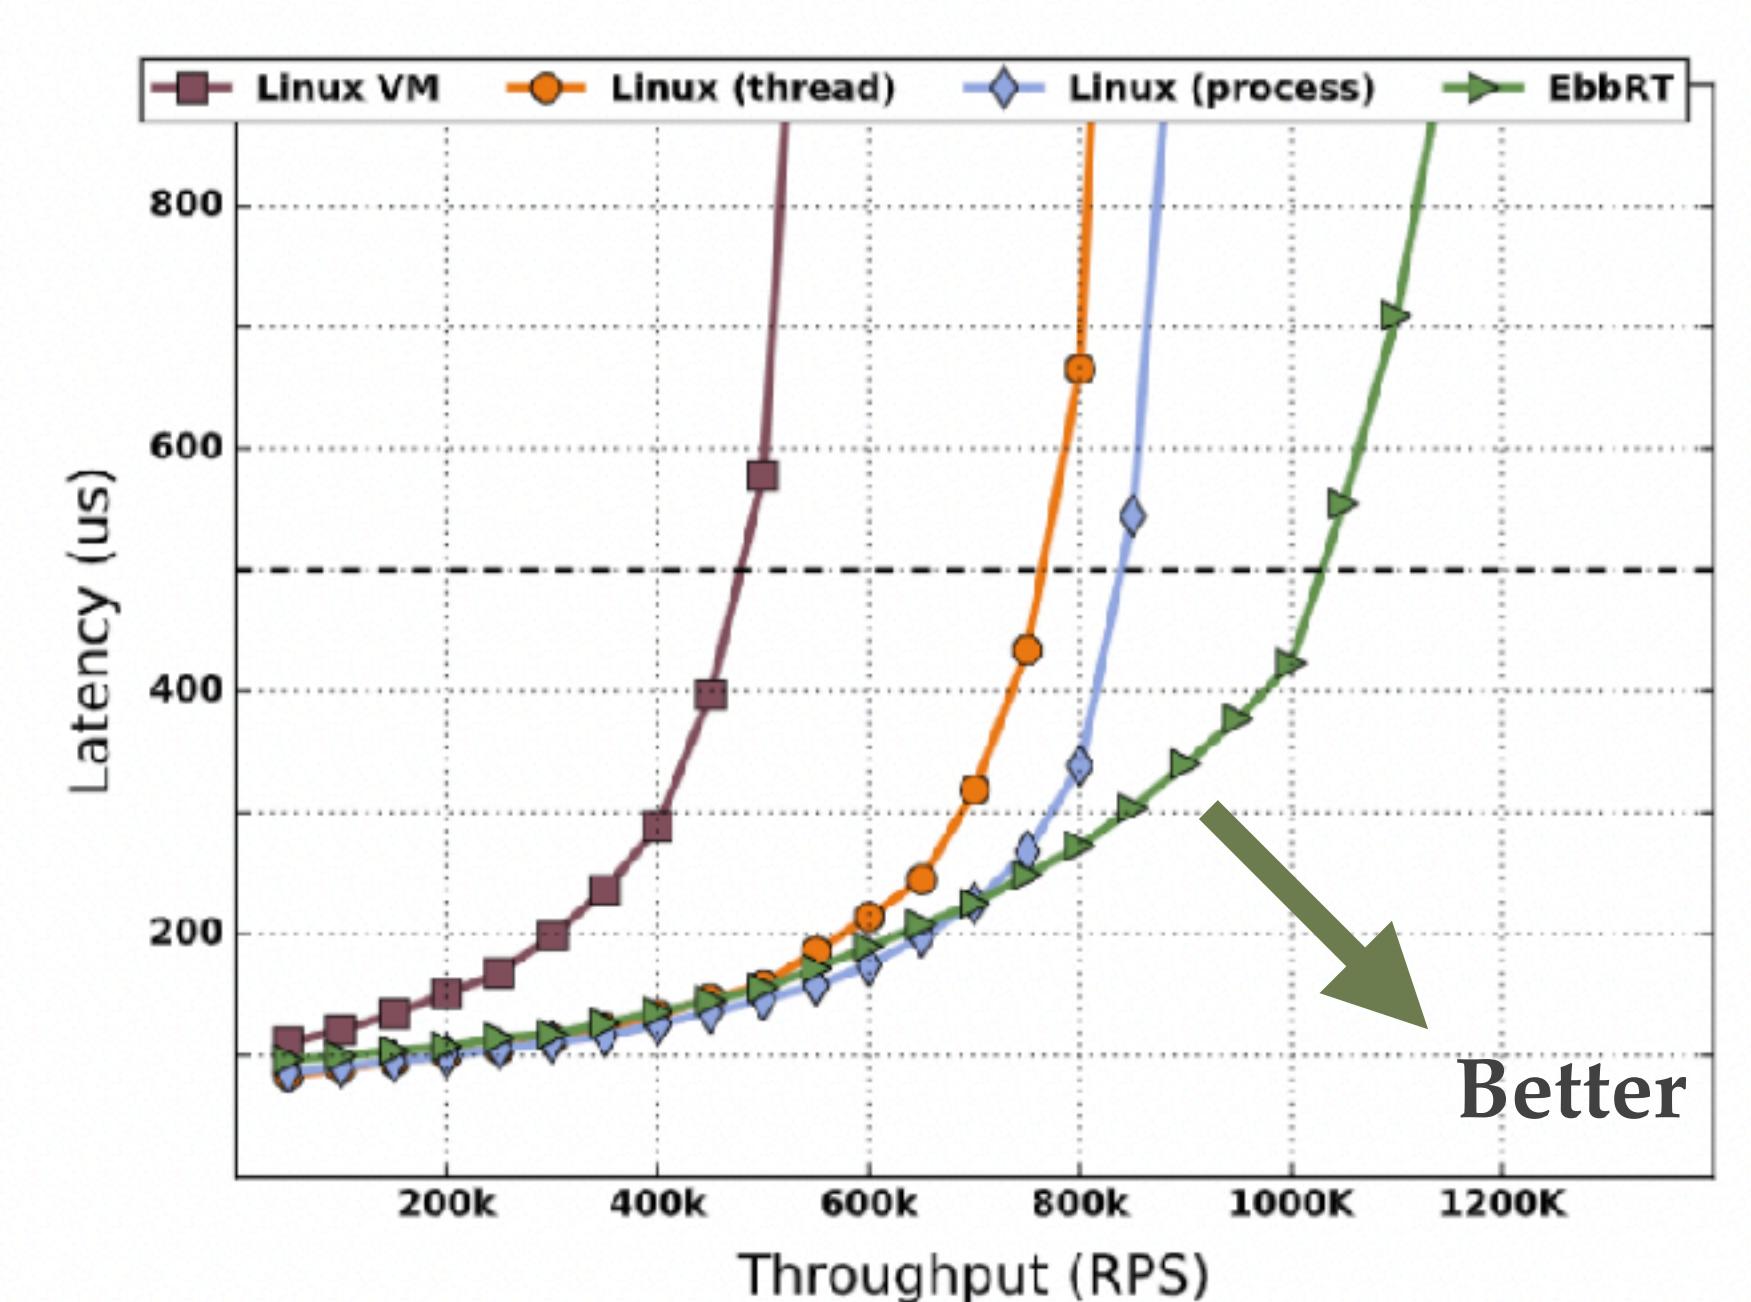
- ❖ Is there a gap between what **computing hardware is capable of**,  
and **what is actually achieved** on modern operating systems?
- ❖ Specialized OSs can serve as a lower bound
  - ❖ Are we within 5%



# Should We Mind the Gap?

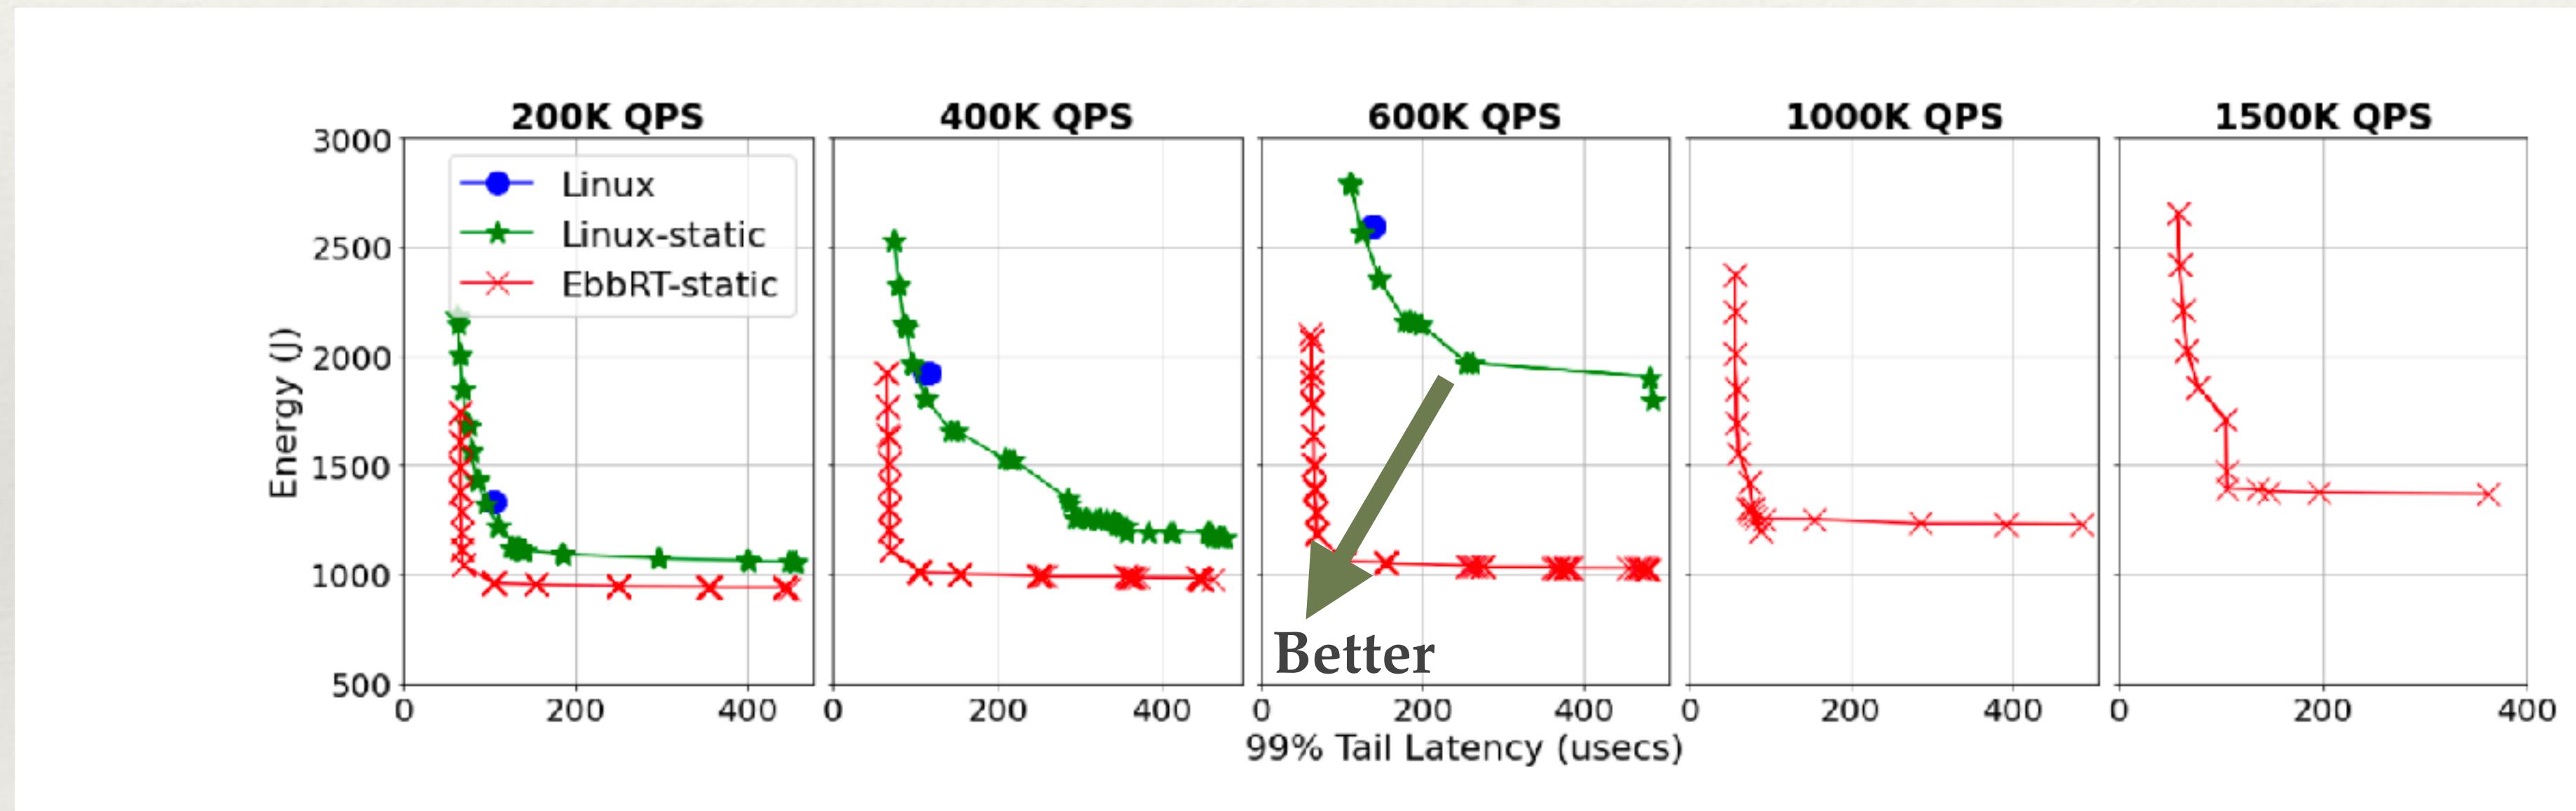
- ❖ Doubled the throughput of memcached, a heavily optimized networked key-value store
- ❖ representative workload
- ❖ well tuned baseline

EbbRT: 2.1x better throughput than Linux on a multi-core memcached



# Should We Mind the Gap?

EbbRT 50% Energy Savings on memcached



# Should We Mind the Gap?

- ❖ Is there a gap between what computing hardware is capable of, and what is actually achieved on modern operating systems?
- ❖ Specialized OSs can serve as a lower bound
  - ❖ Are we within 10%?



---

# Rigidity Prevents Tradeoffs

---

- ❖ A static divide between application and kernel introduces rigidity
  - ❖ Applications face immutable system interfaces, can't trade off
  - ❖ Tools only apply to one domain or the other
  - ❖ Programmers get siloed

---

# Start with GPOS Whittle Down

---

- ❖ Start with the full feature application context
- ❖ Extend it with OS-Level Access
- ❖ Modify the system to meet application requirements

---

# Ch 2: Introducing Dynamic Privilege

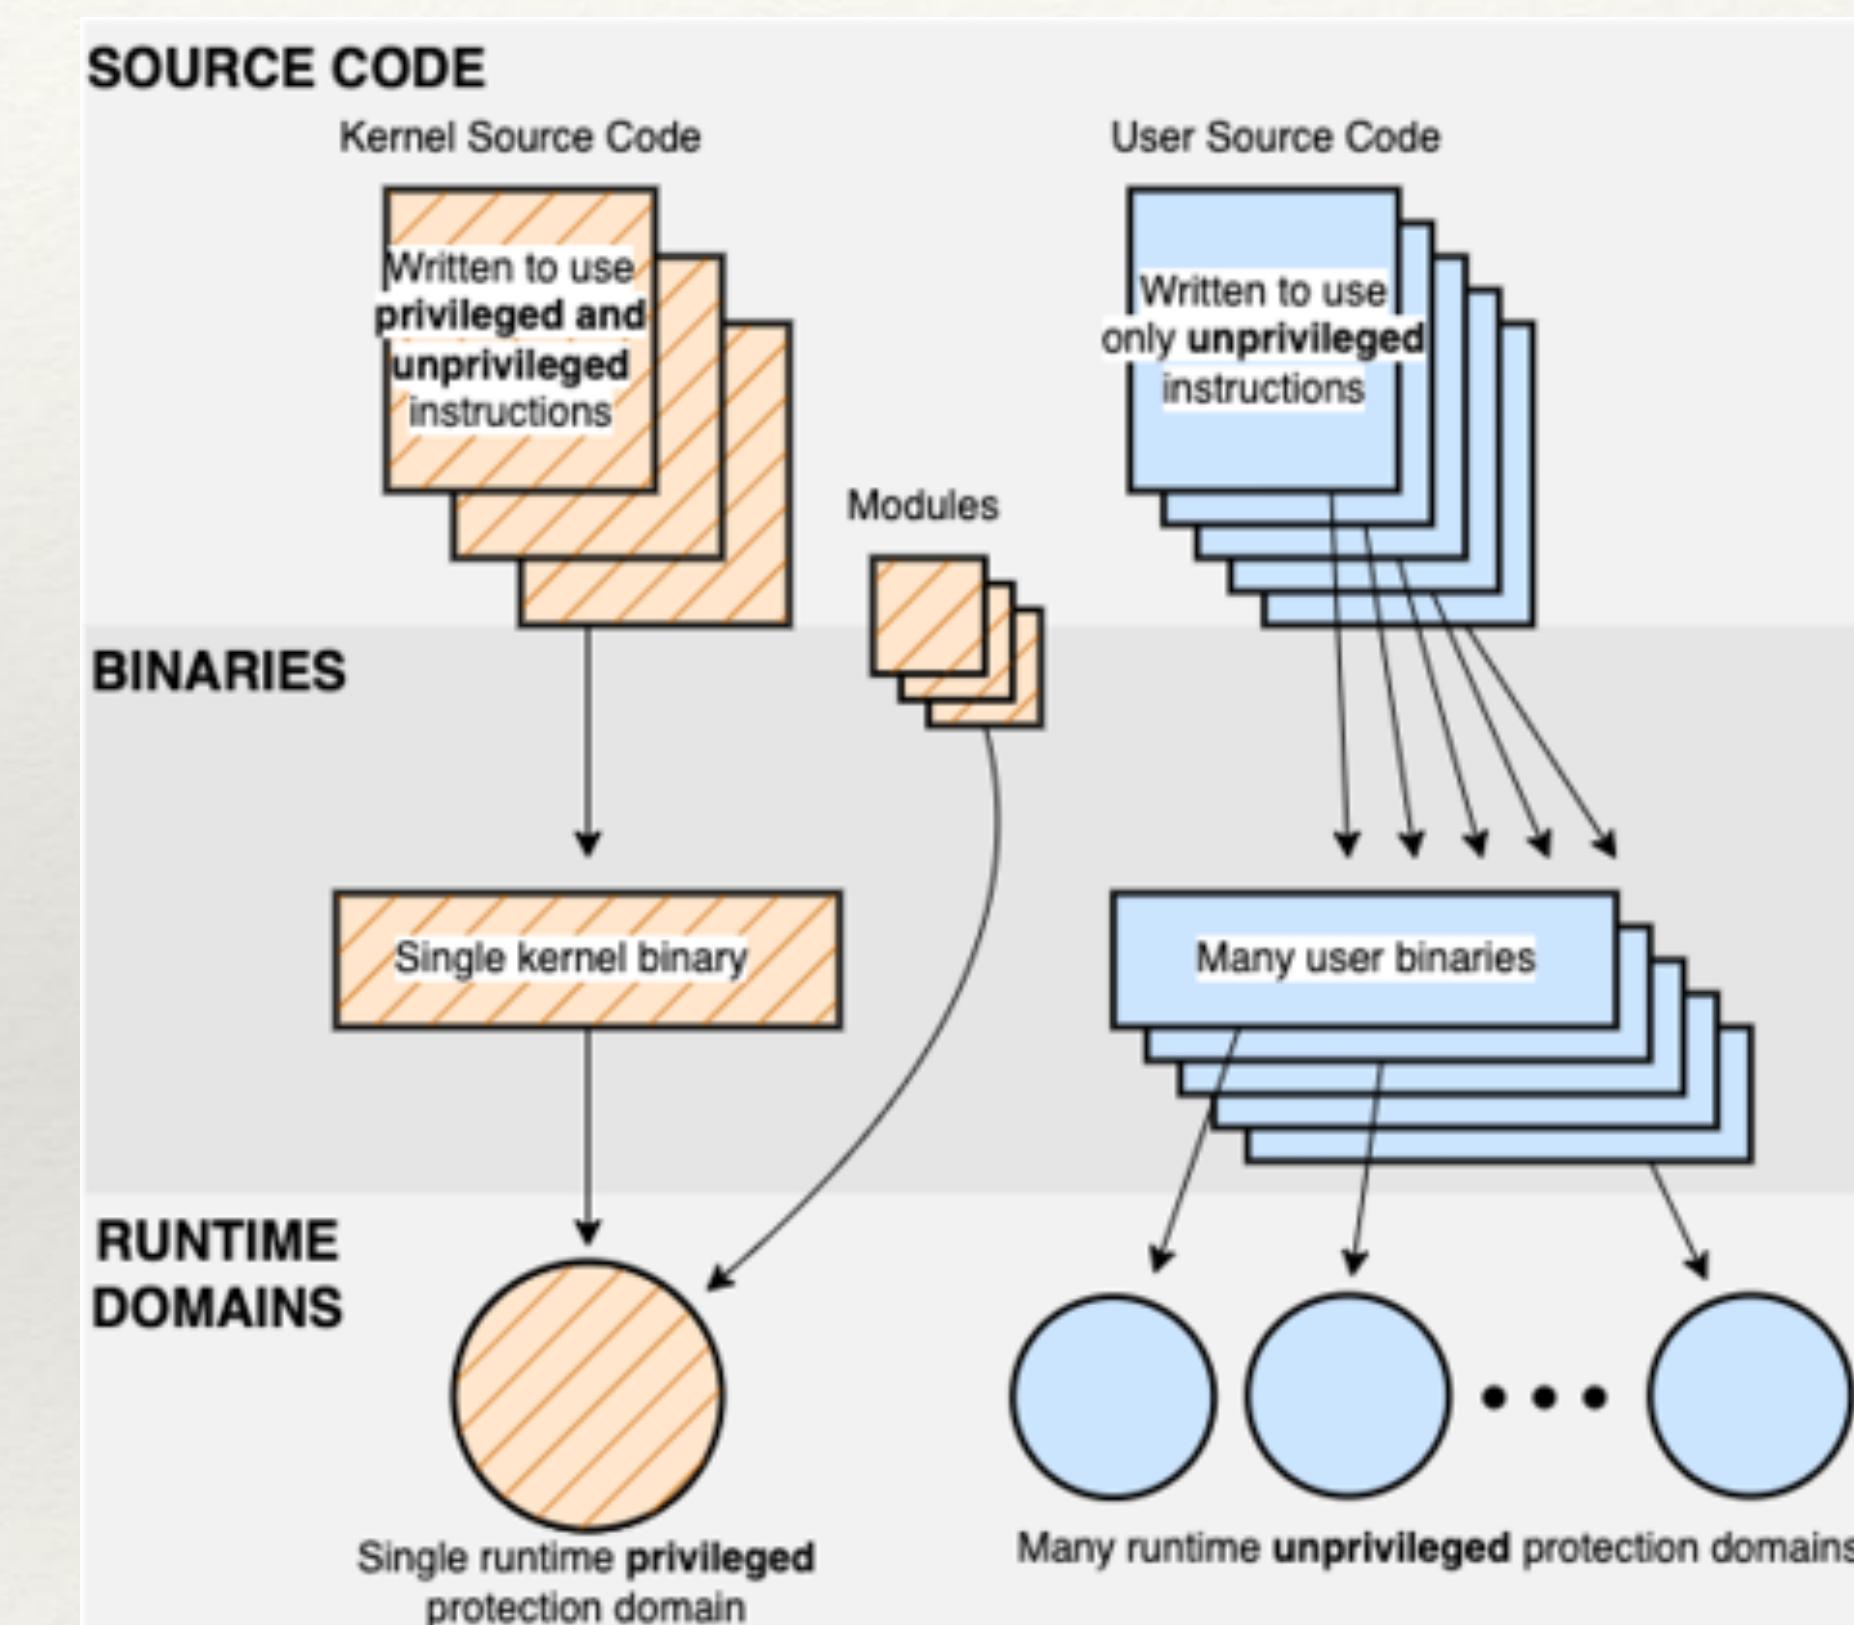
---

- ❖ The static divide
- ❖ It is built from hardware privilege
- ❖ Dynamic Privilege, a mechanism that orthogonalizes hardware privilege
- ❖ A thesis statement

# Software is Separated by Privilege



- ❖ Applications and kernels are divided at every stage of their lifetimes
  - ❖ Separate **code-bases**
  - ❖ Independent **binaries**
  - ❖ Isolated **runtimes**



Source to Runtime Spectrum

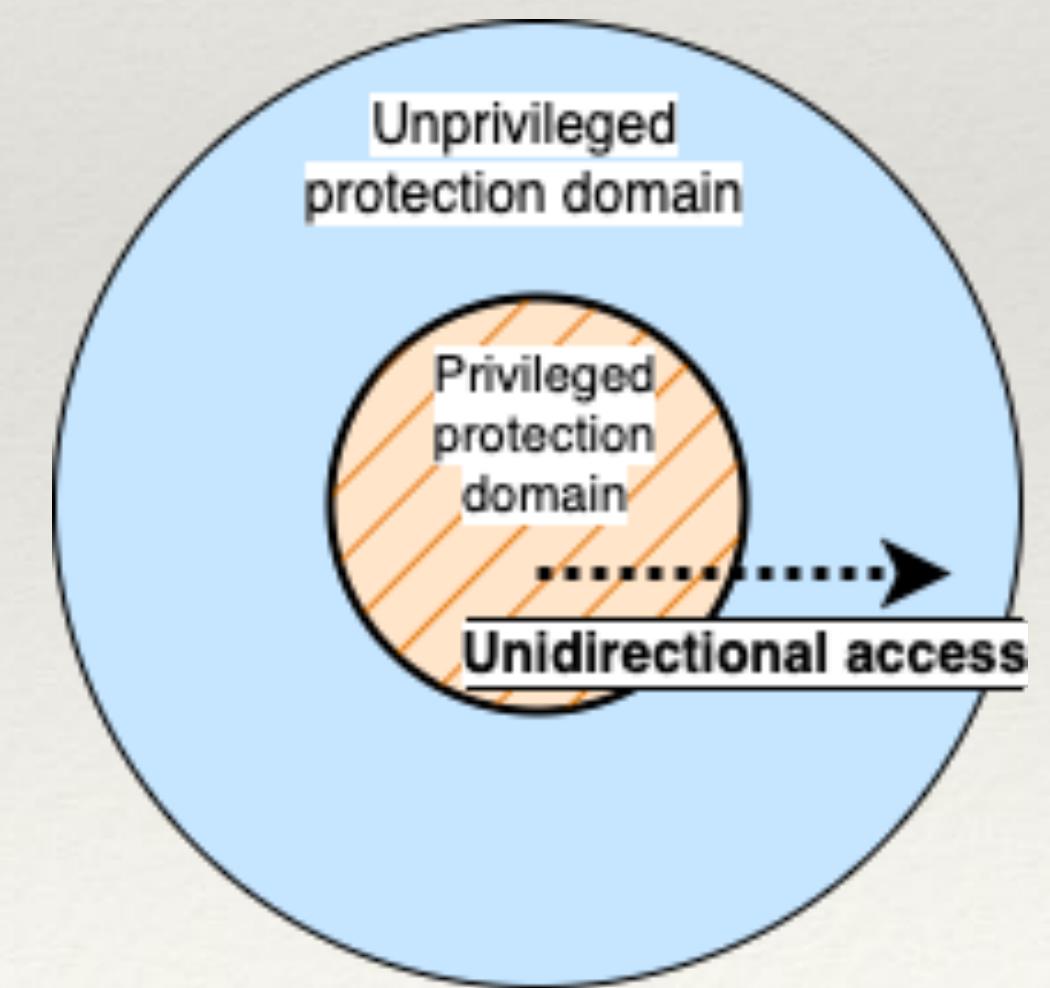
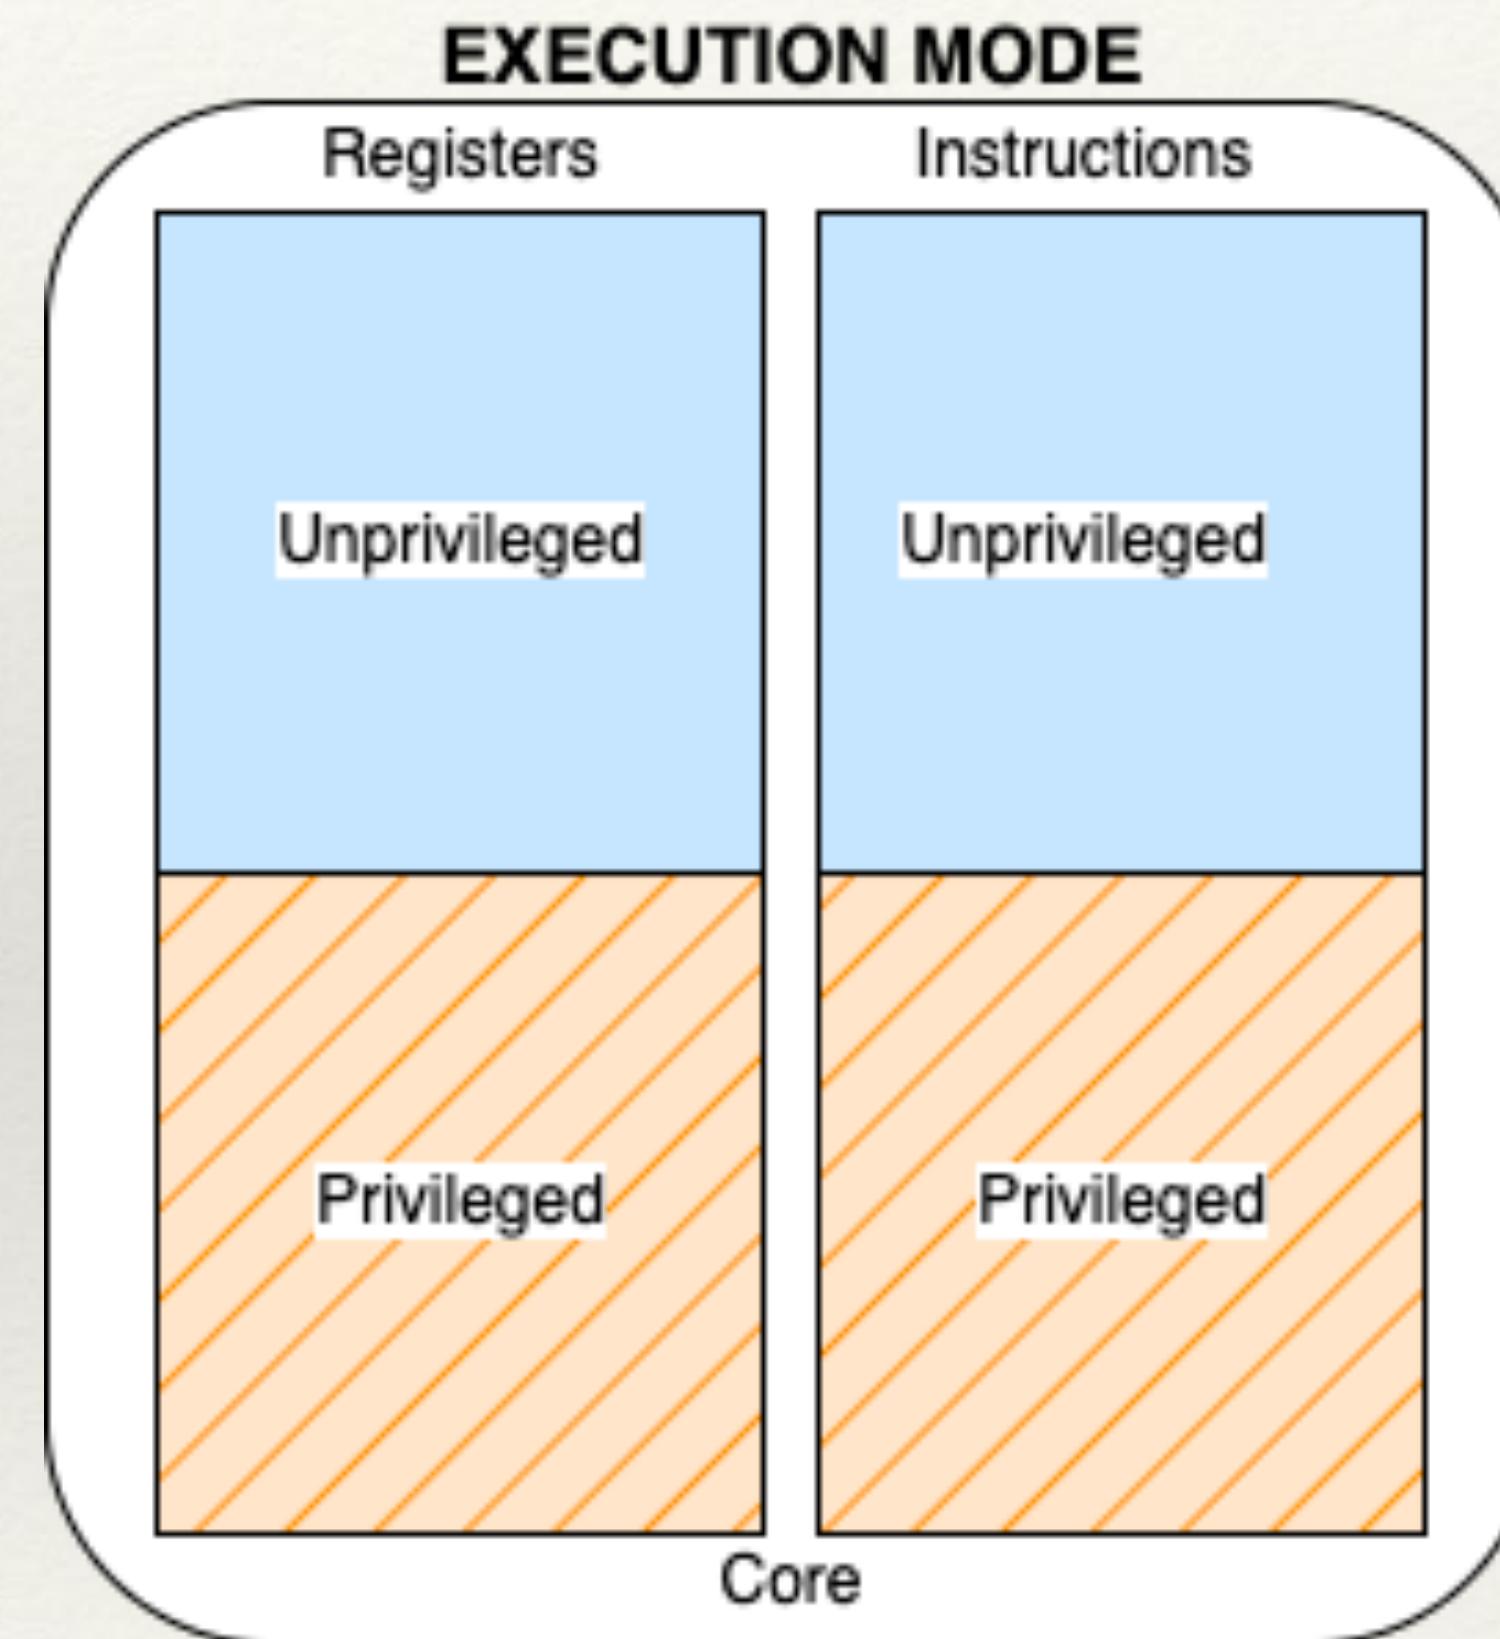
# Hardware Privilege

- ❖ Processors (Hardware) operate in distinct modes of execution
  - ❖ Unprivileged or user-mode
  - ❖ Privileged or supervisor mode



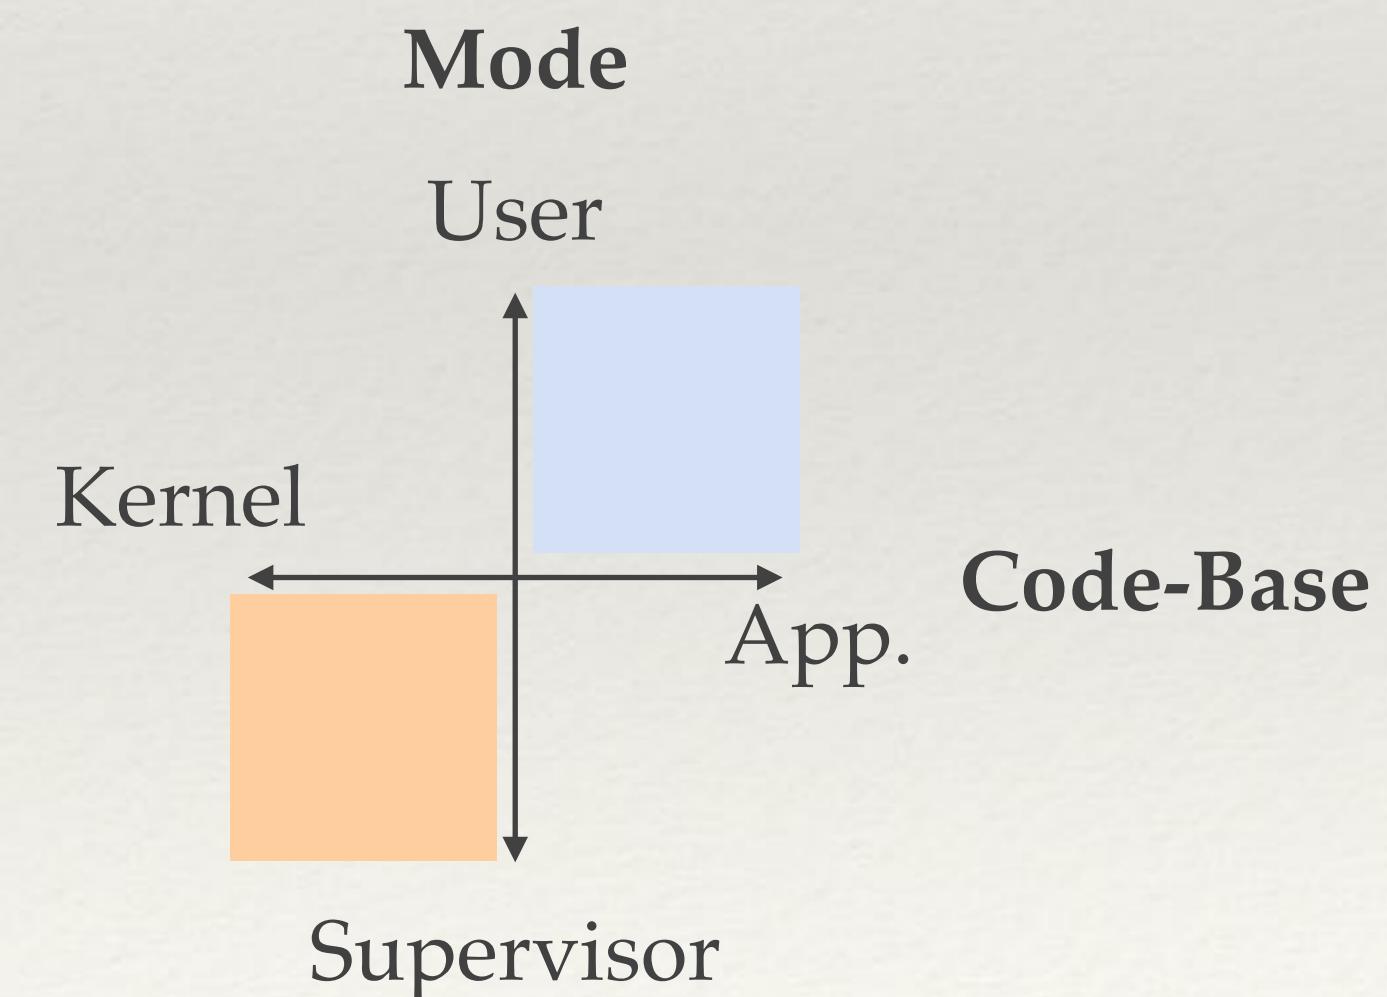
# Hardware Privilege

- ❖ What can you access?
  - ❖ Instructions
  - ❖ Registers
  - ❖ Memory



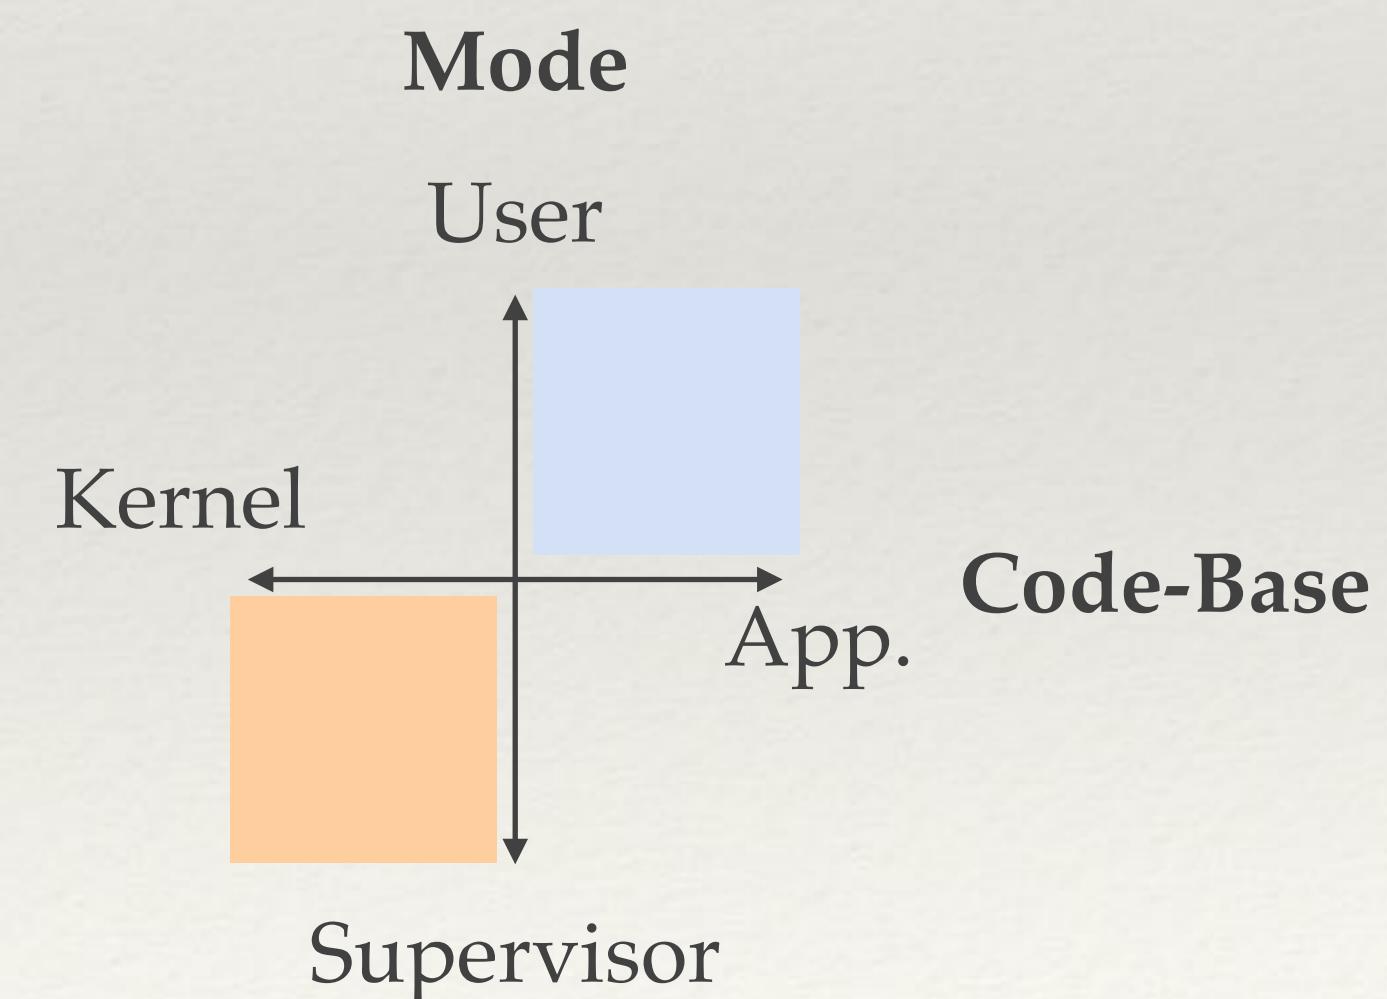
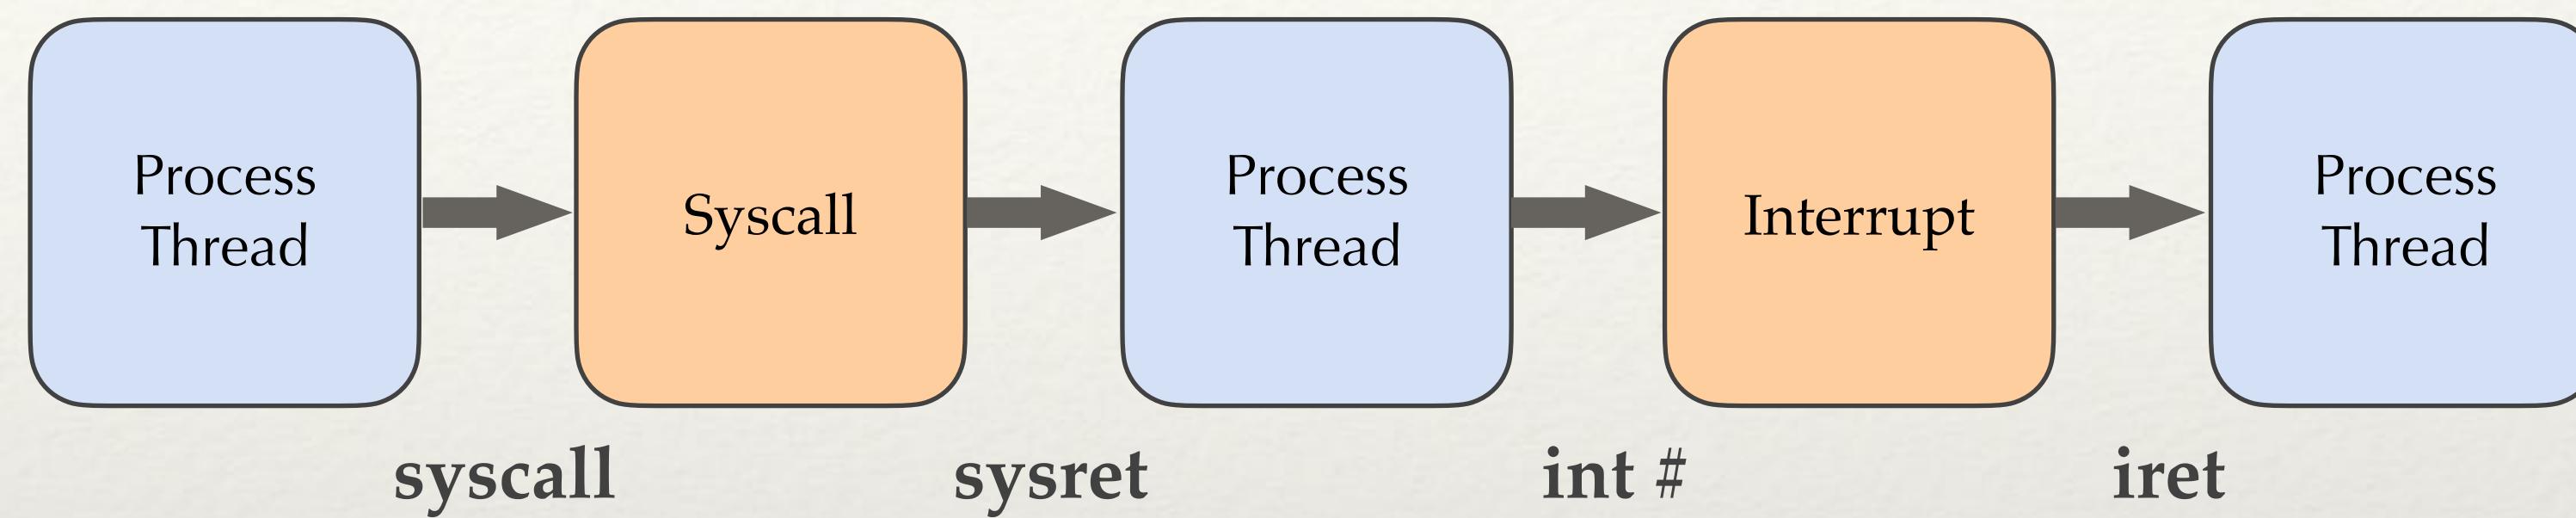
# Privilege Separates Execution

- ❖ Coupled transitions
  - ❖ Codebase Change <-> Privilege Level Change

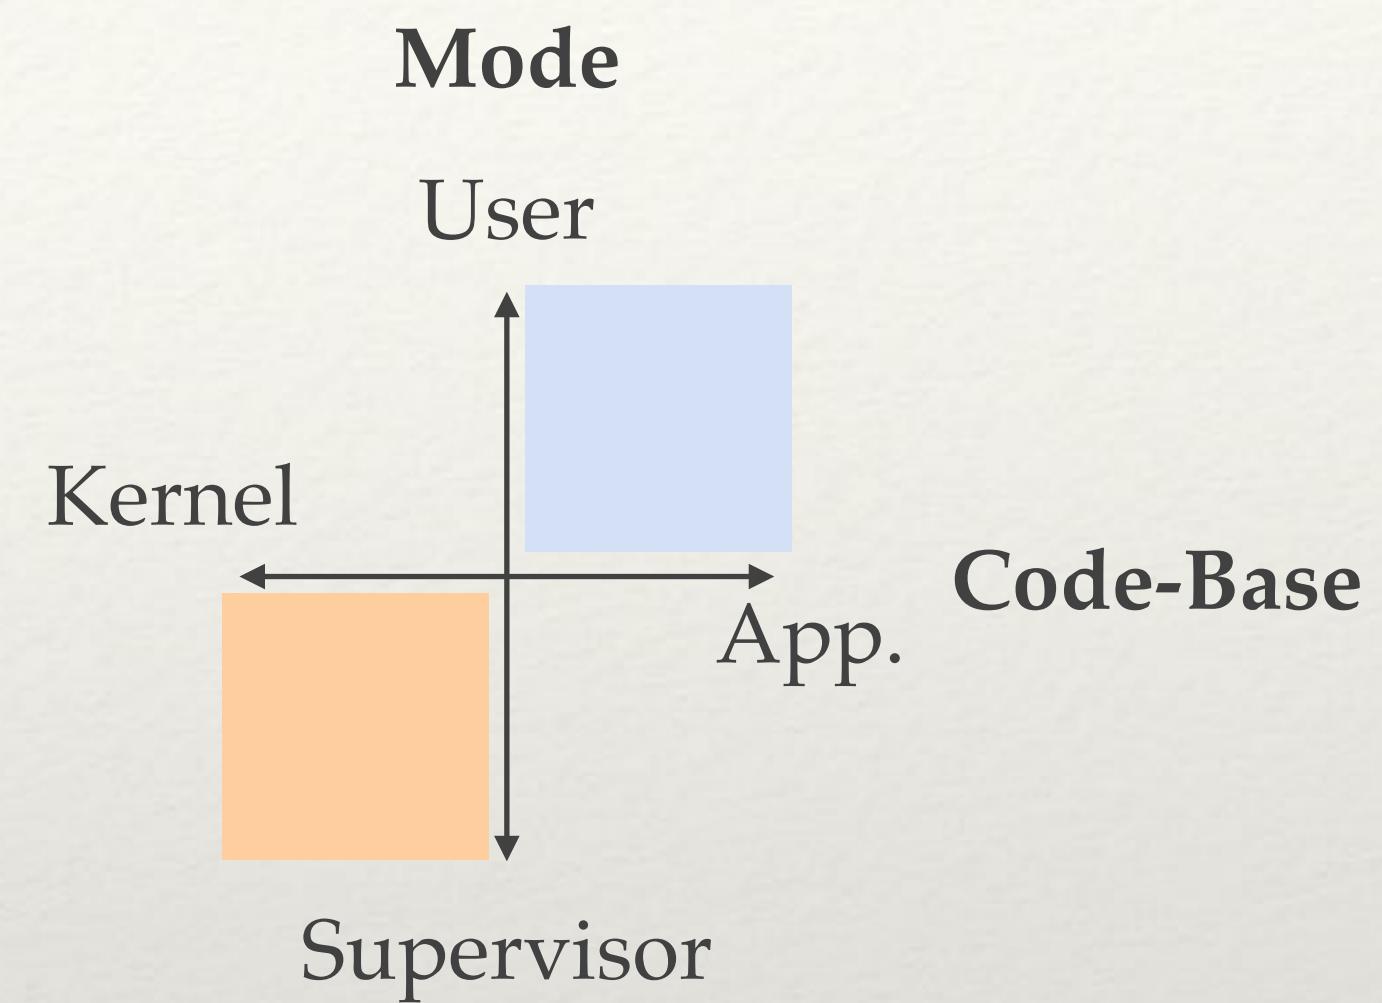
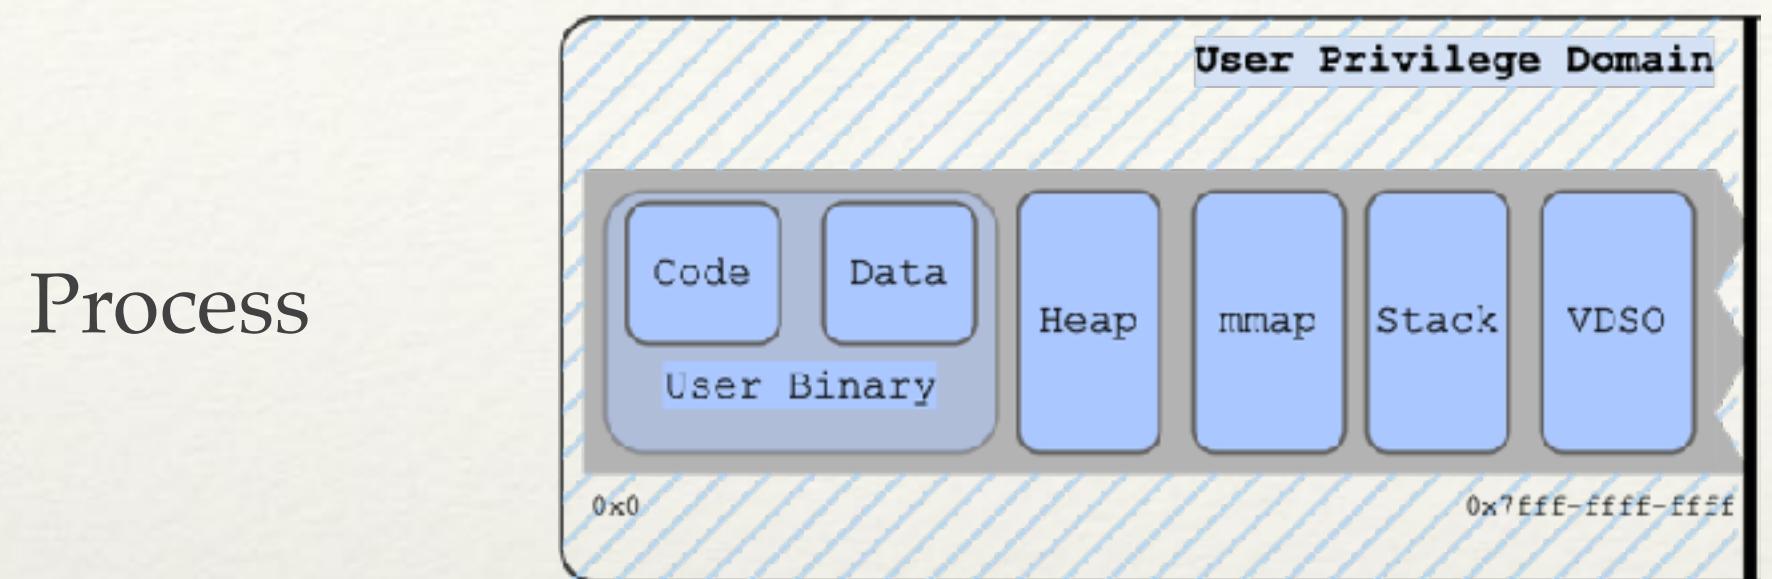


# Privilege Separates Execution

- ❖ Coupled transitions
- ❖ Codebase Change <-> Privilege Level Change

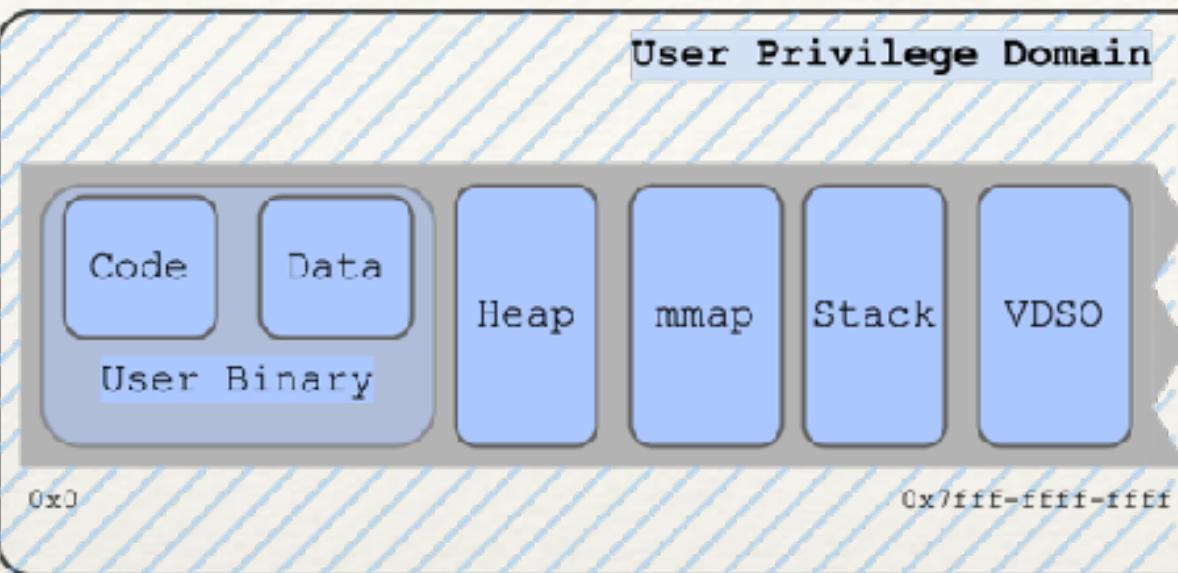


# Privilege Separates State (Memory)

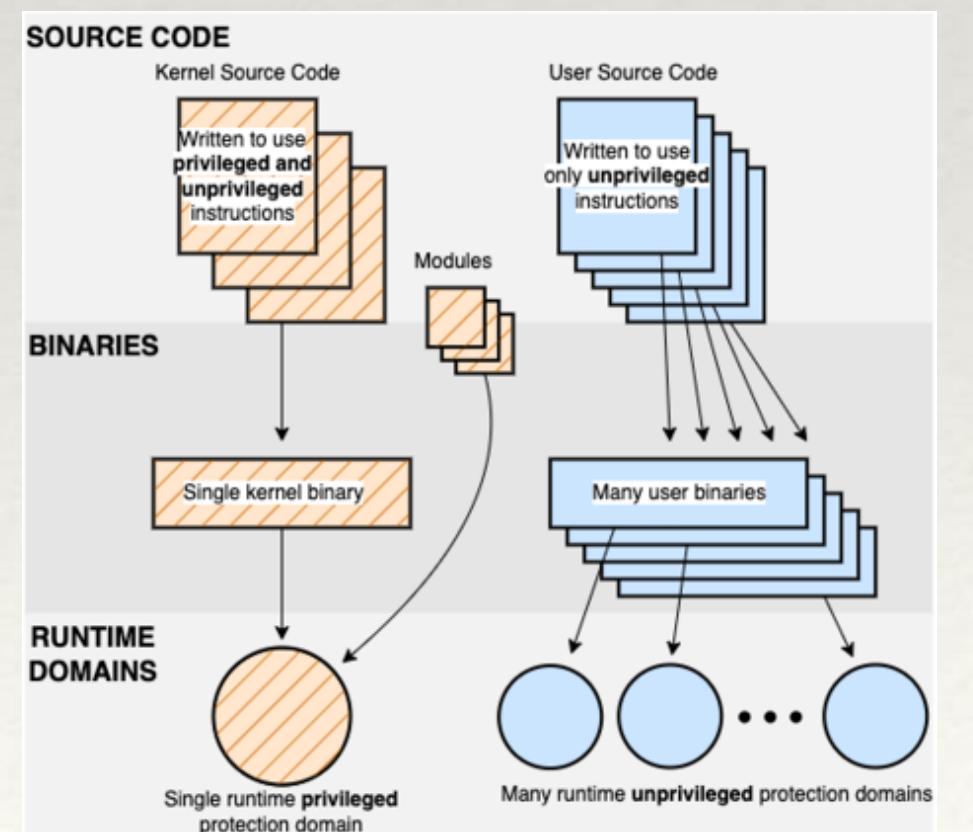
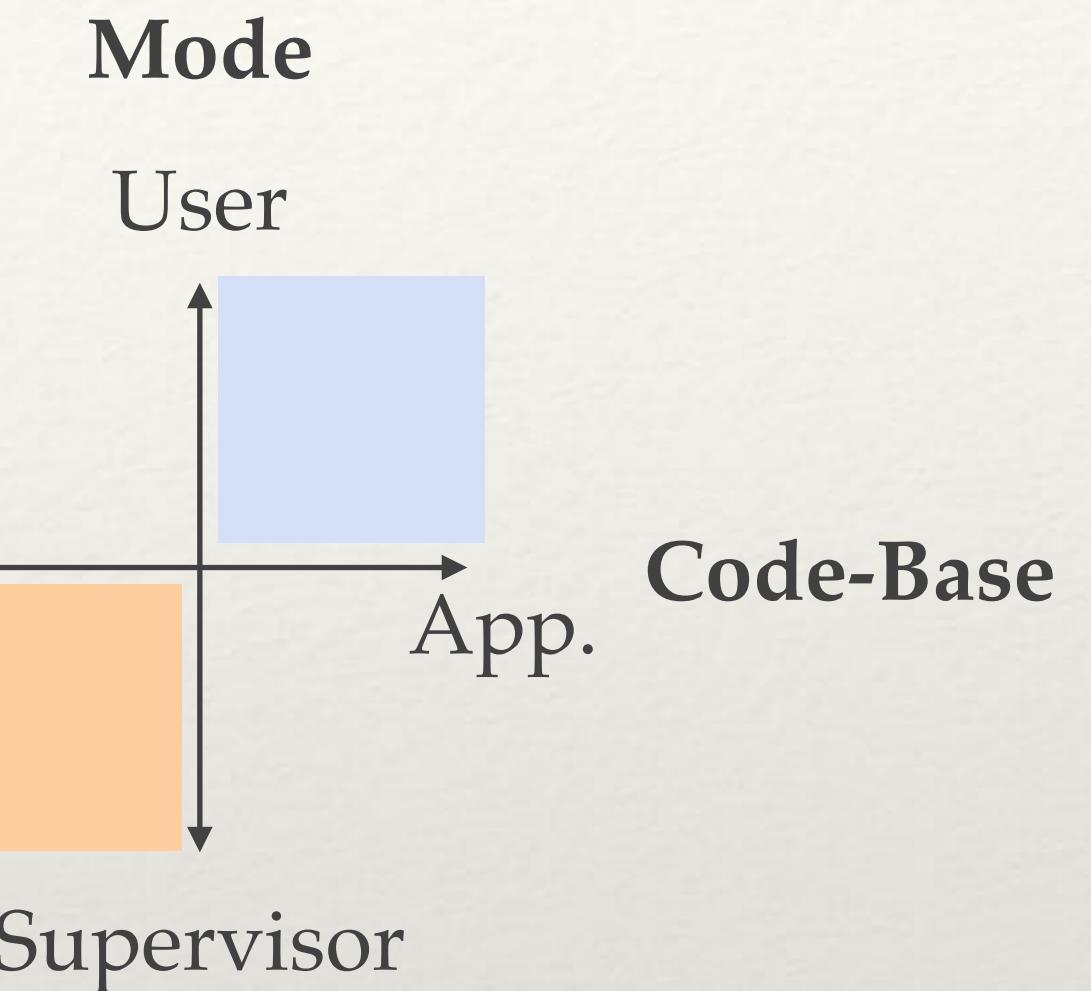
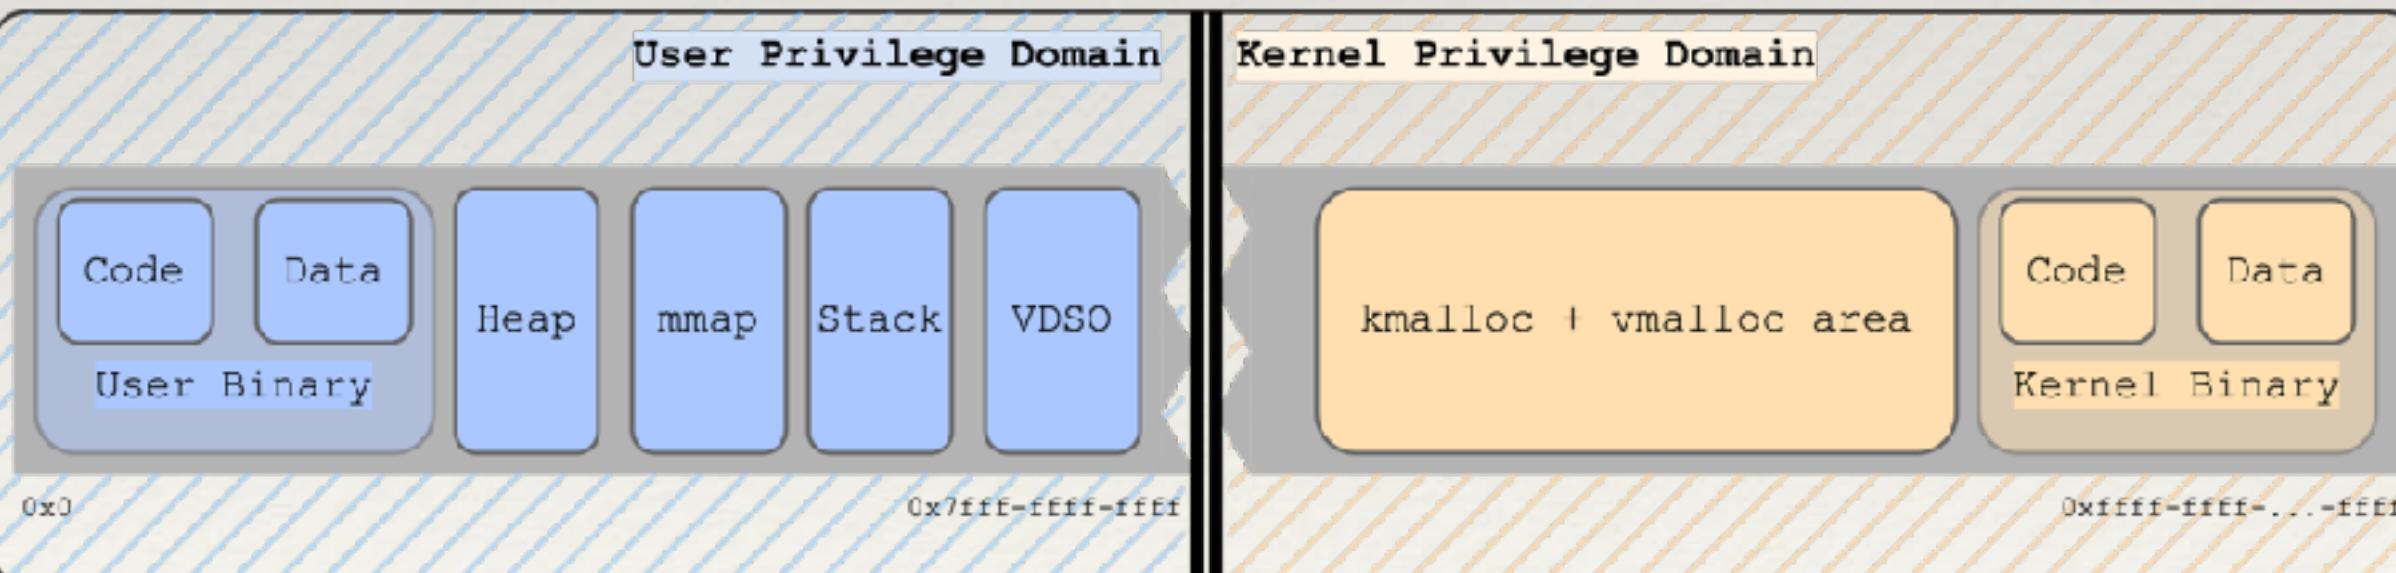


# Privilege Separates State (Memory)

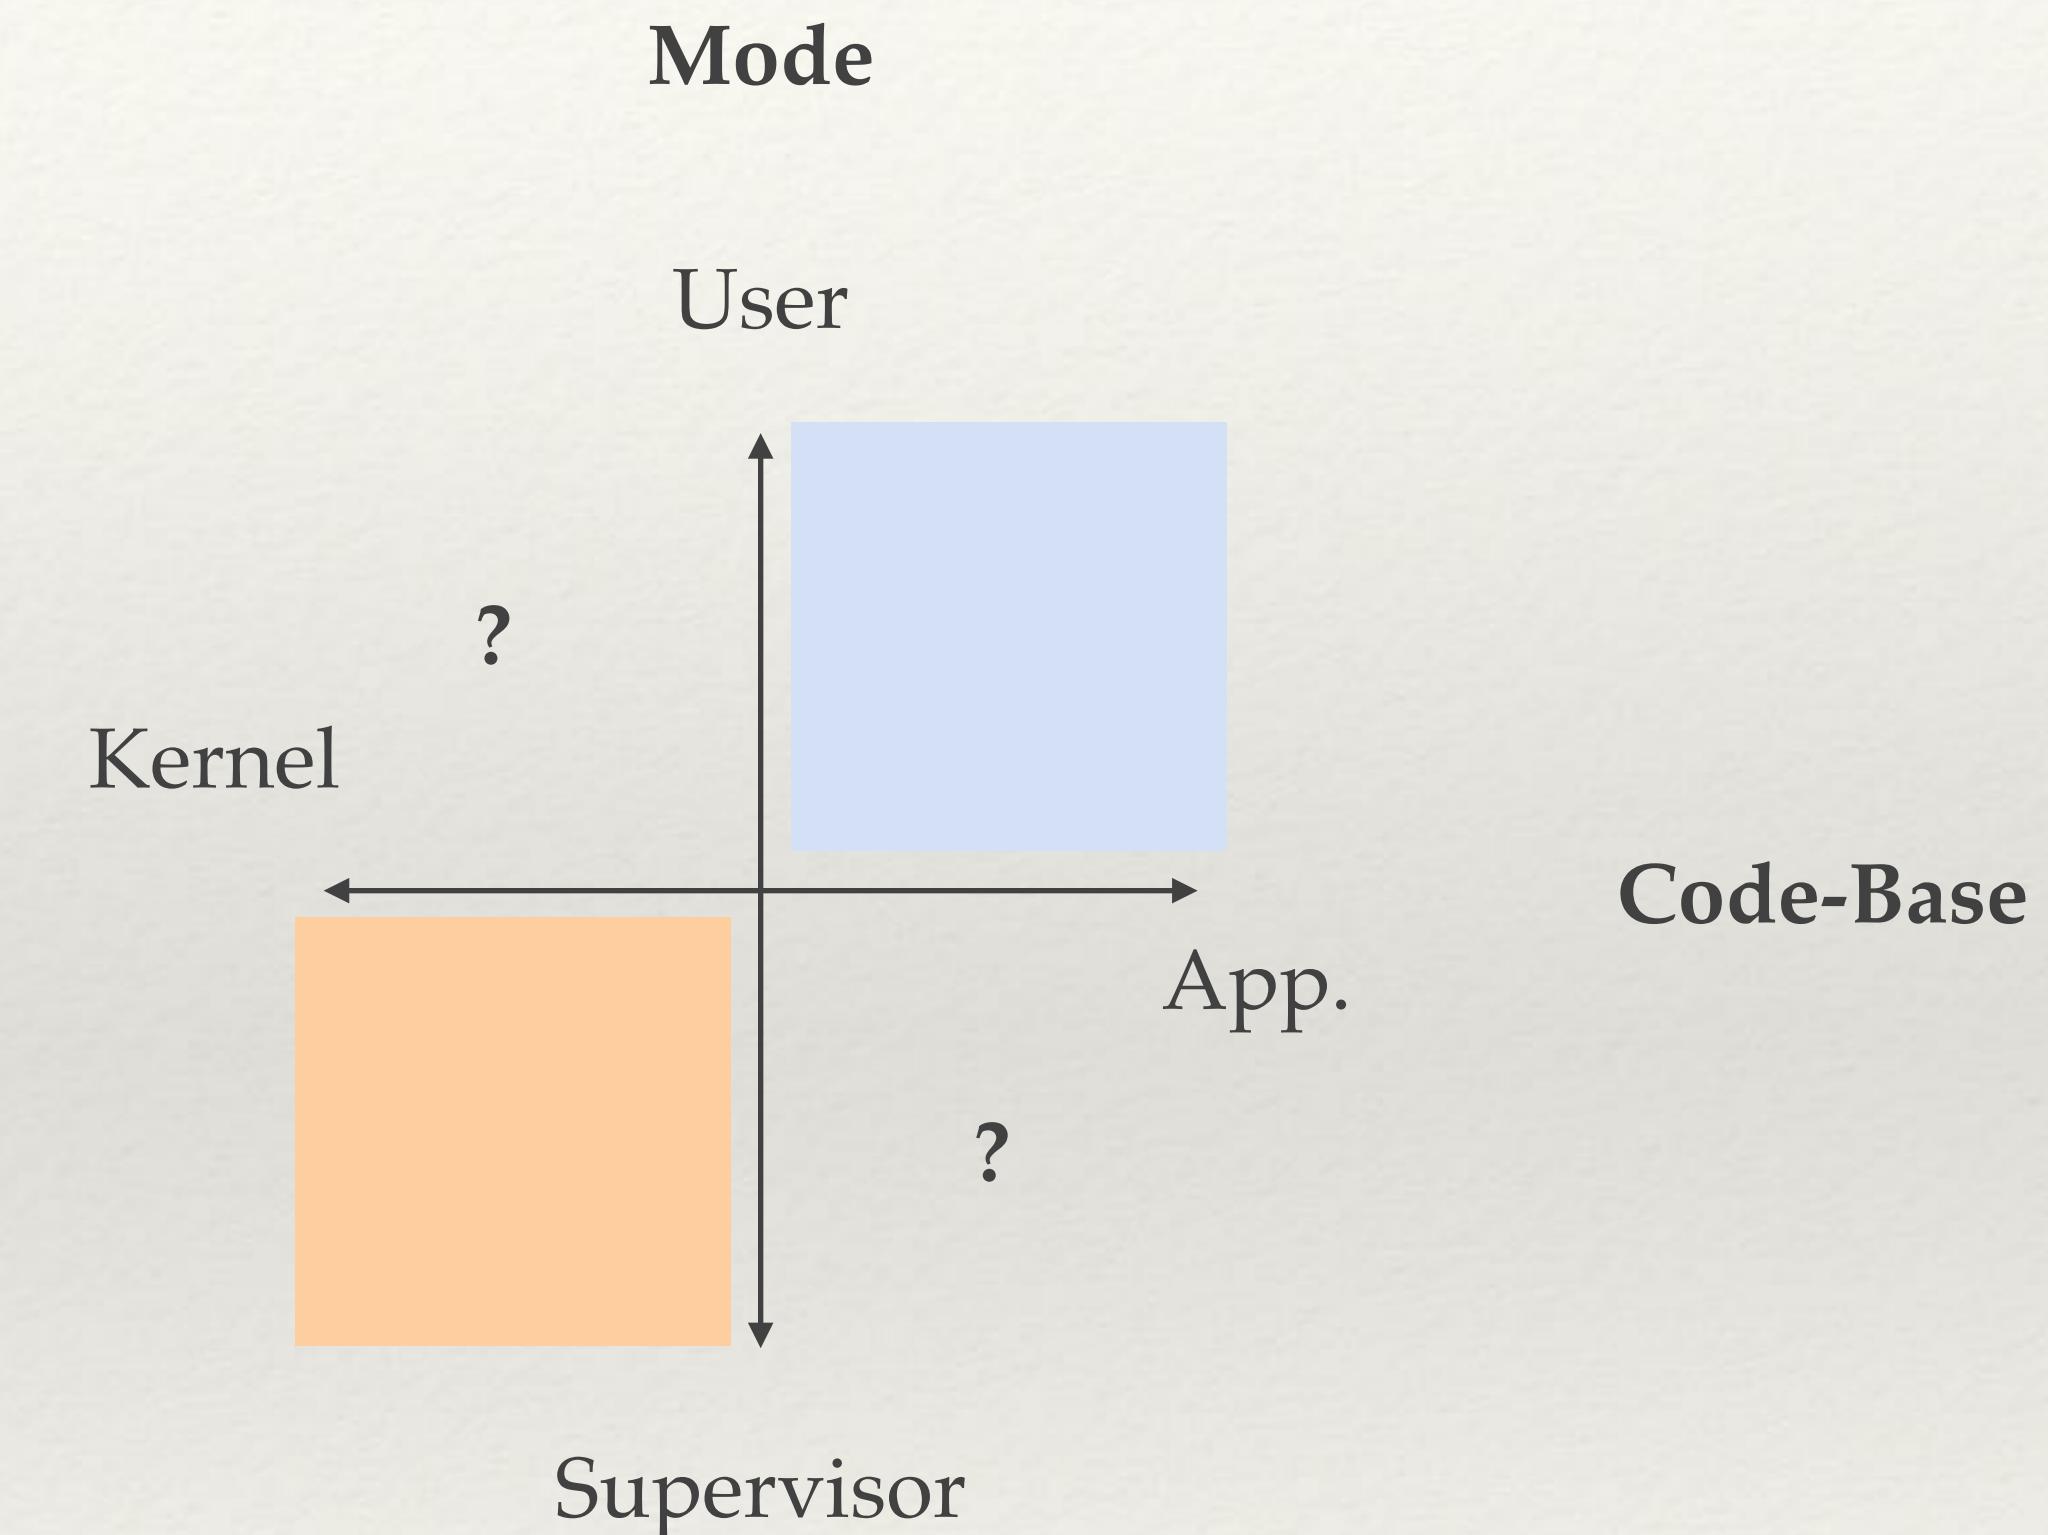
Process



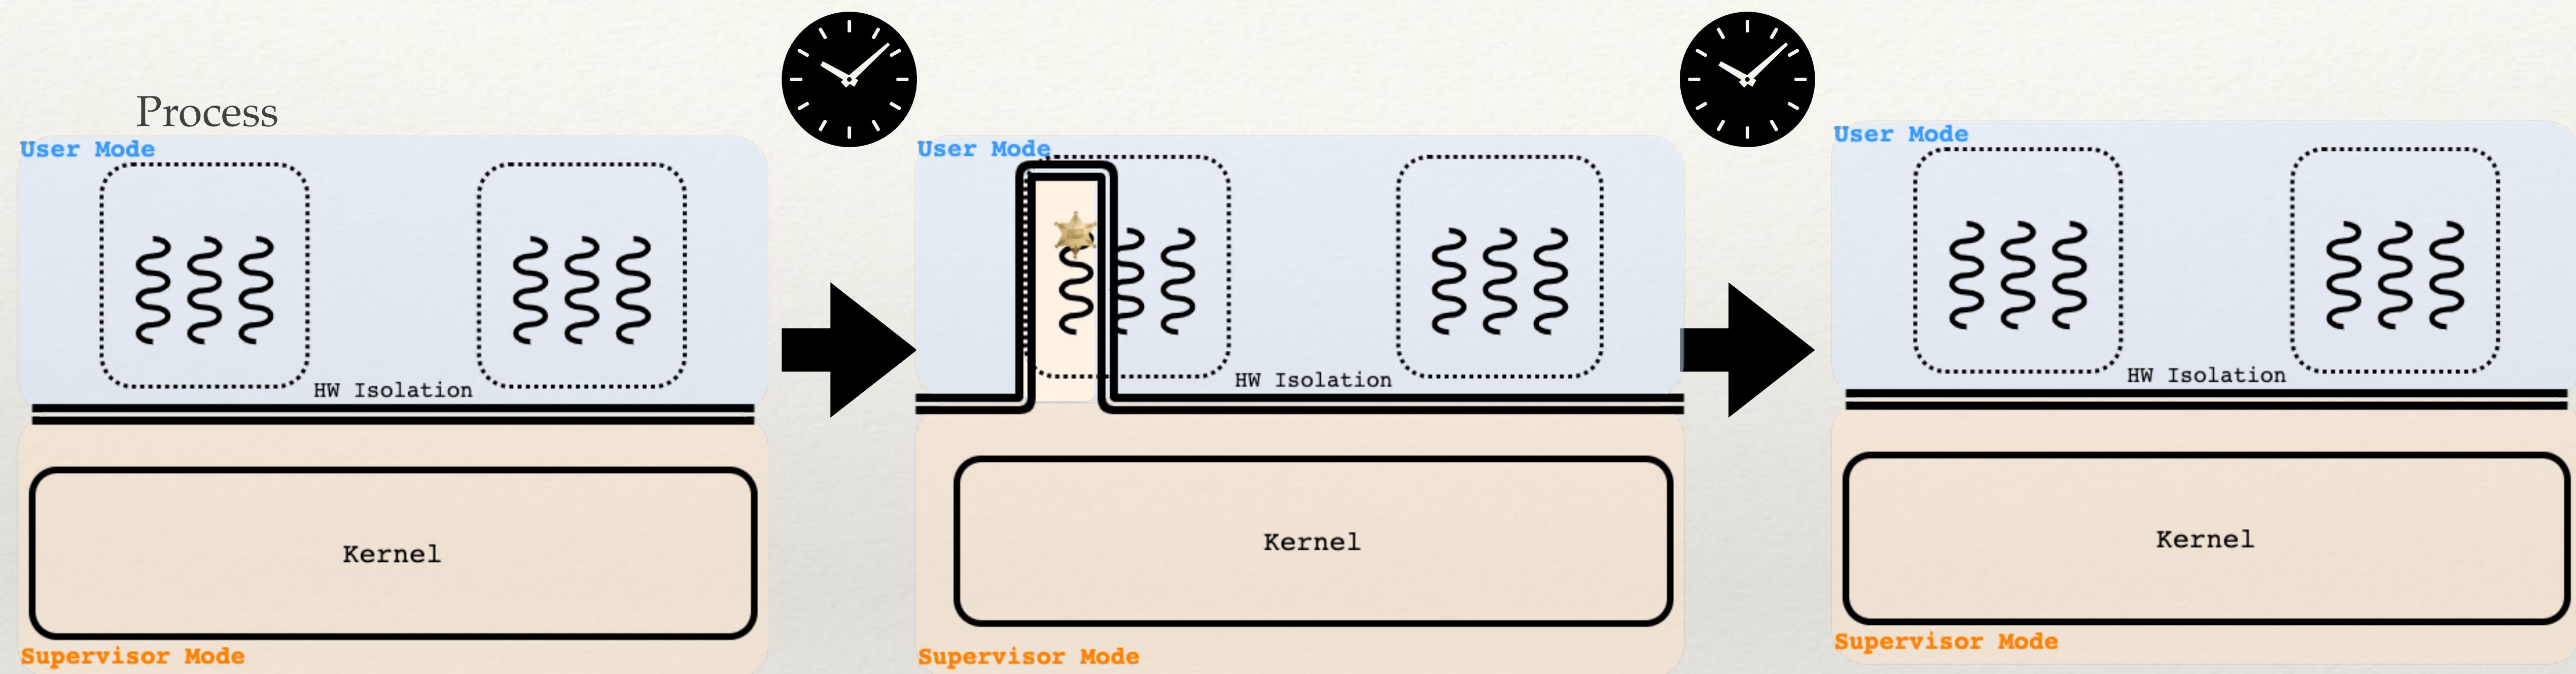
Kernel



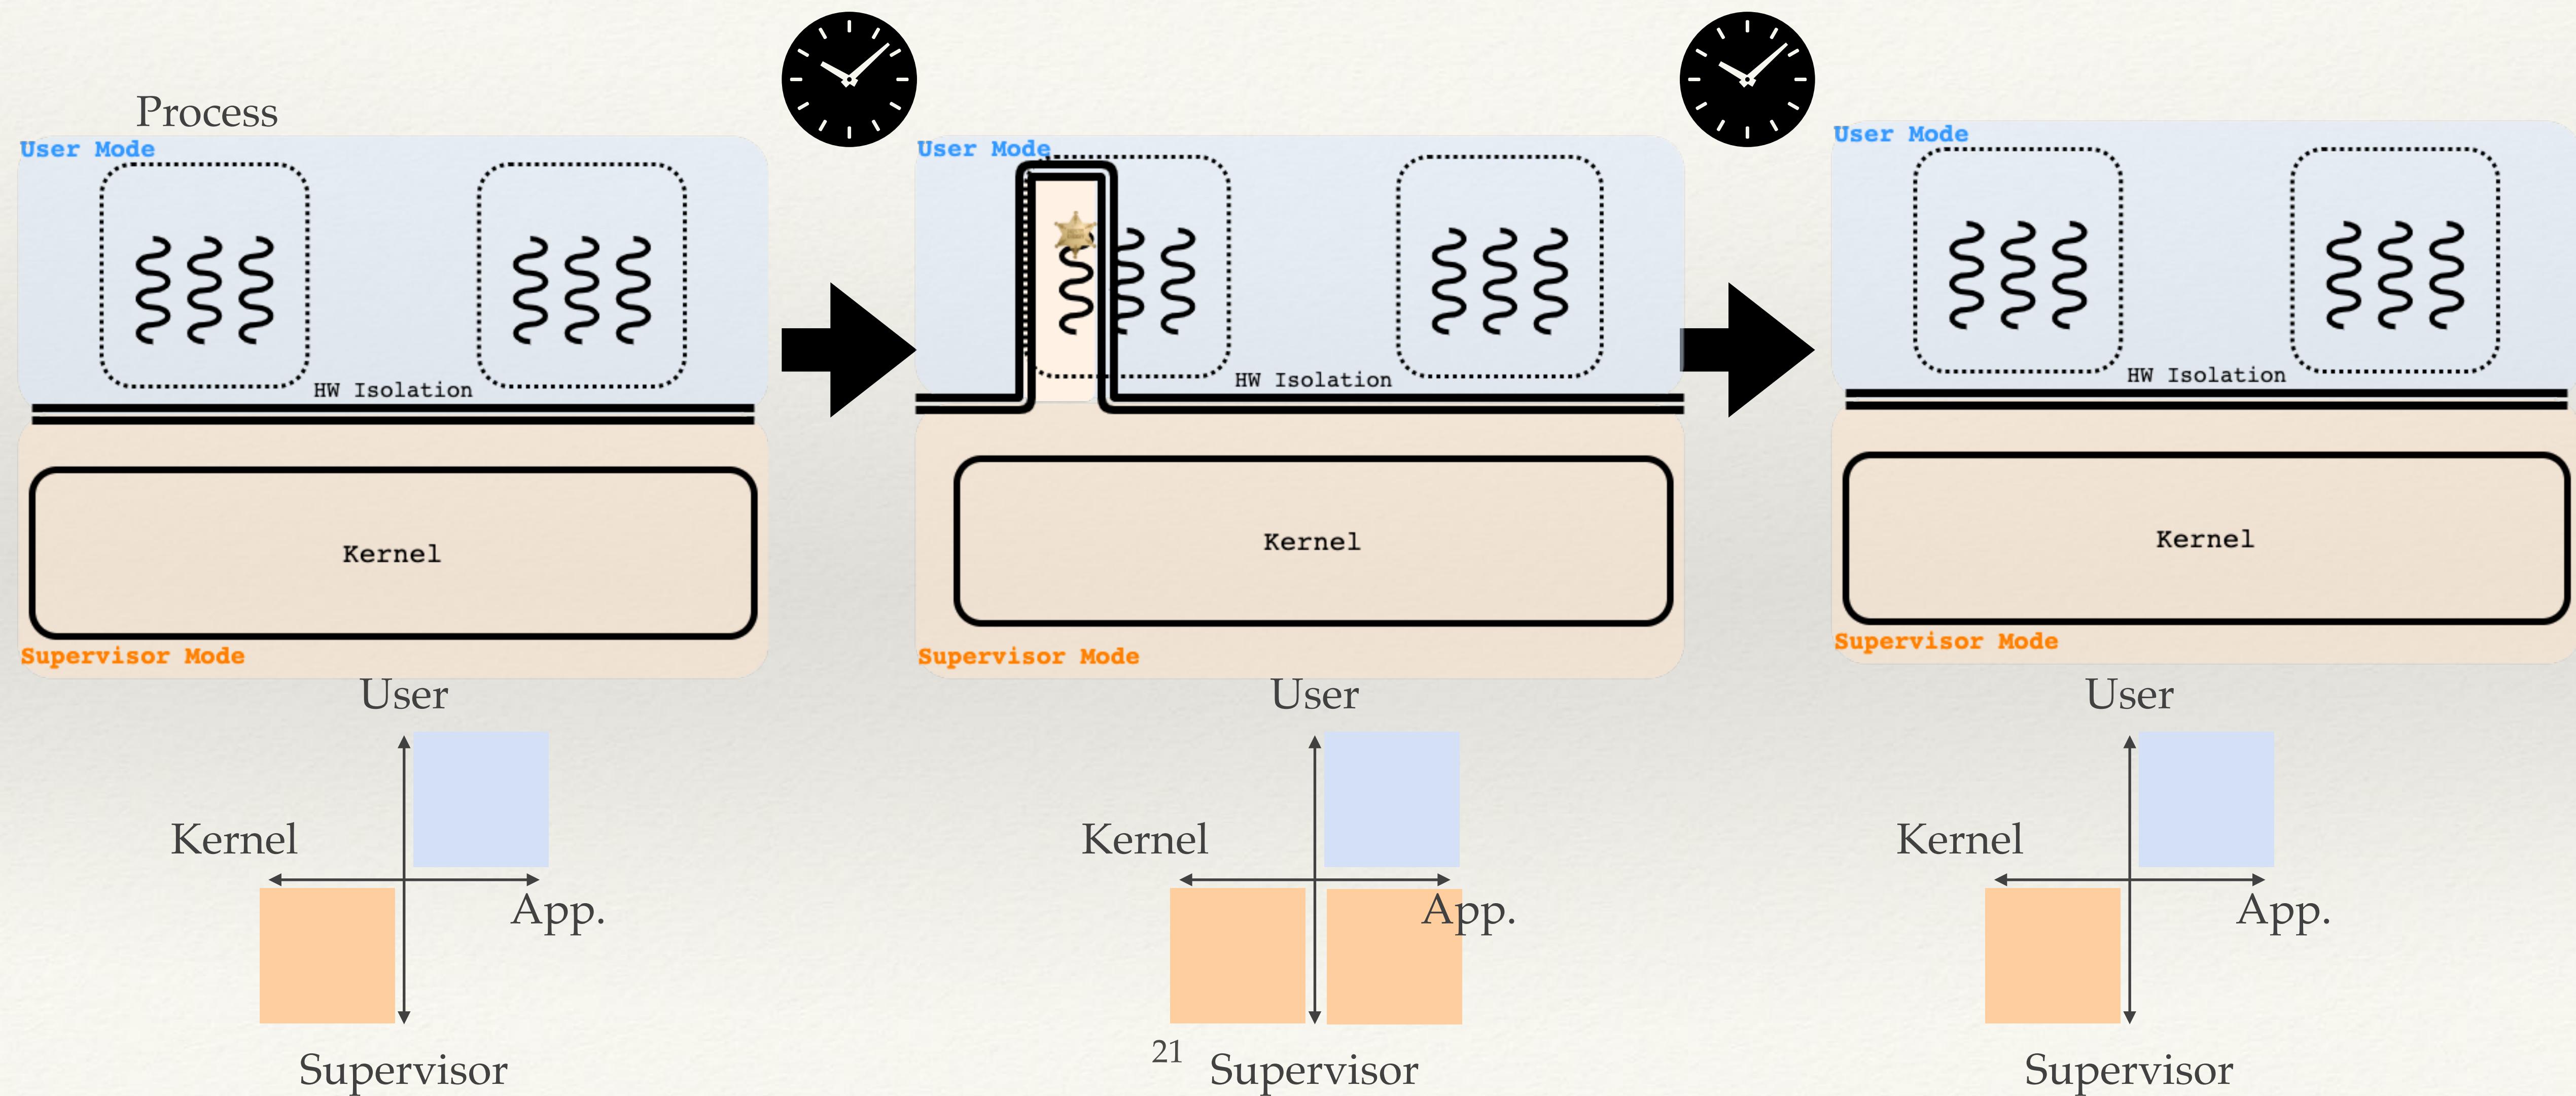
# Dynamic Privilege:



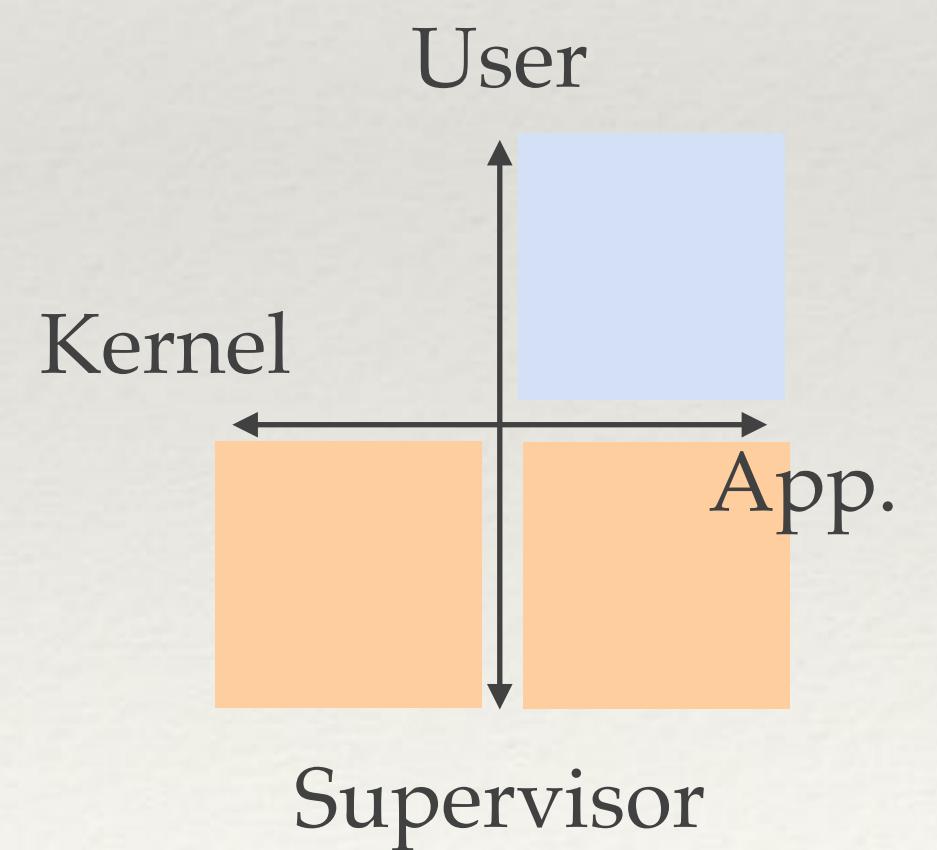
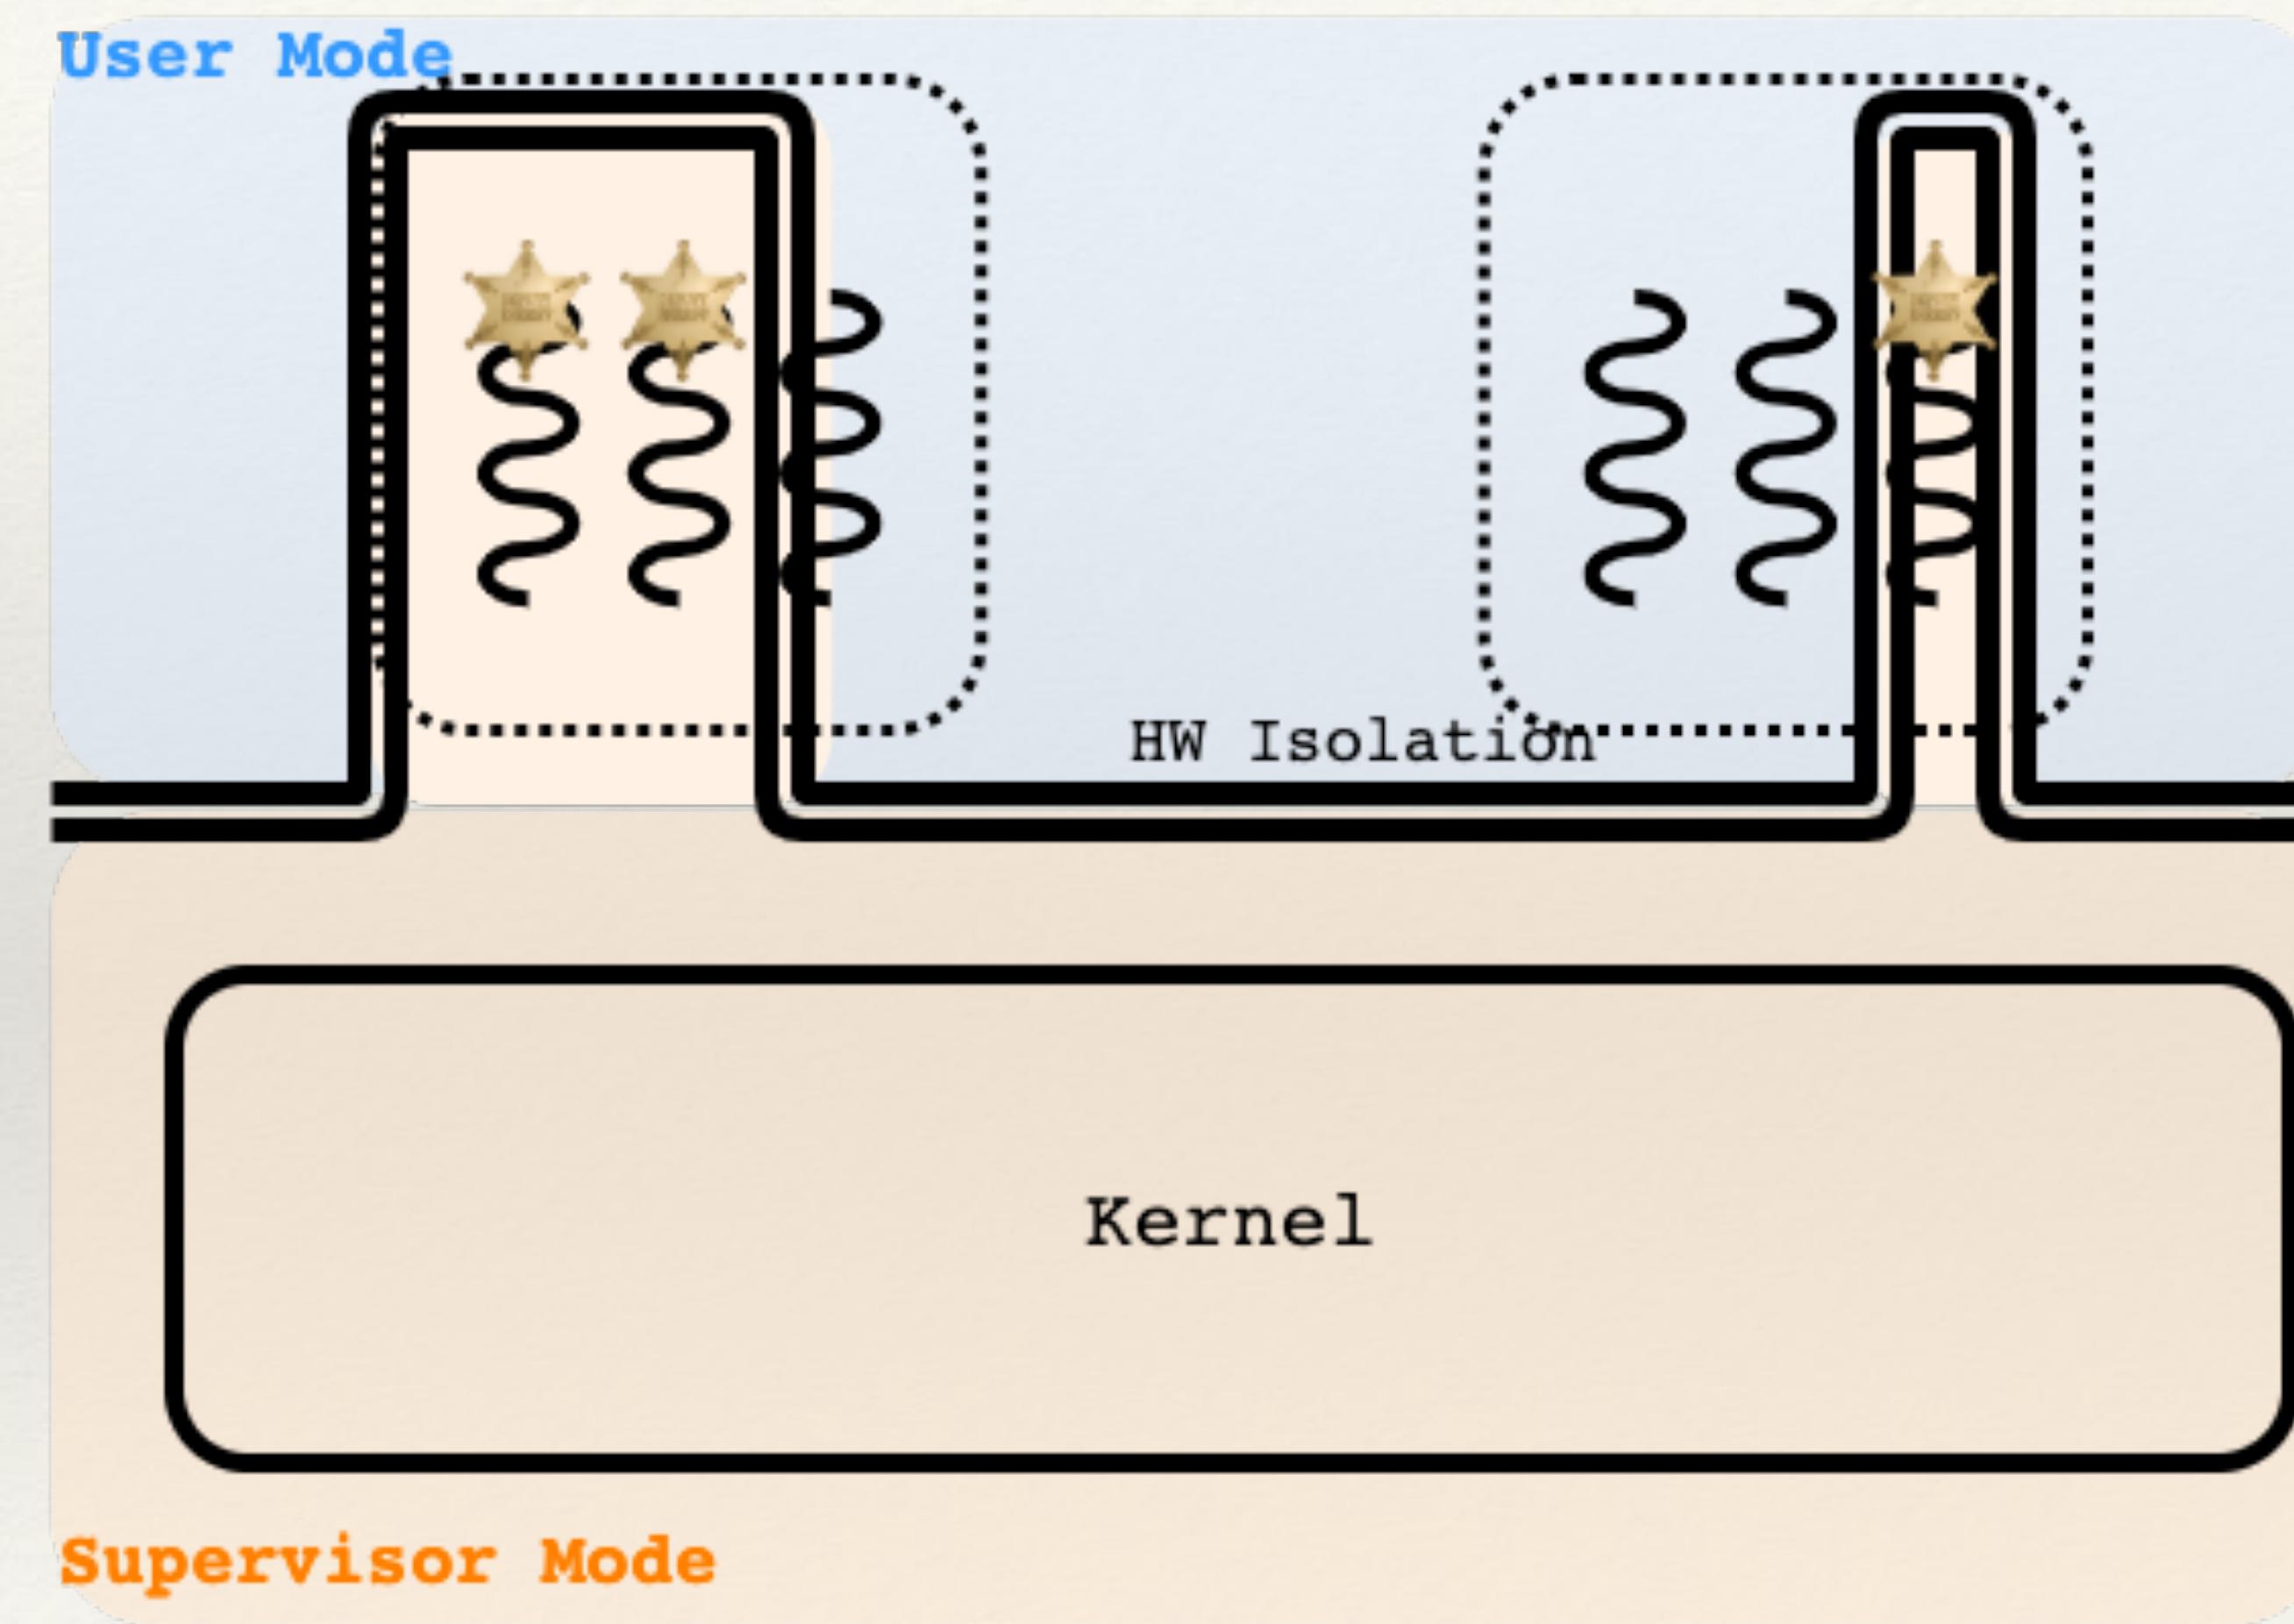
# Dynamic Privilege



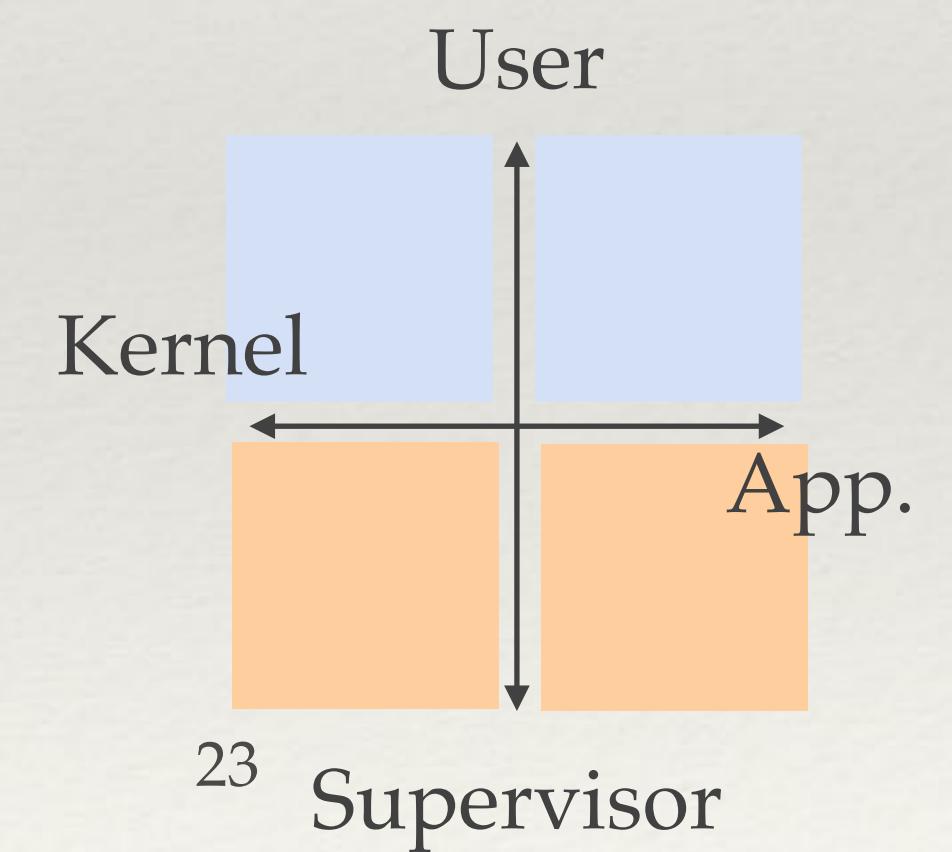
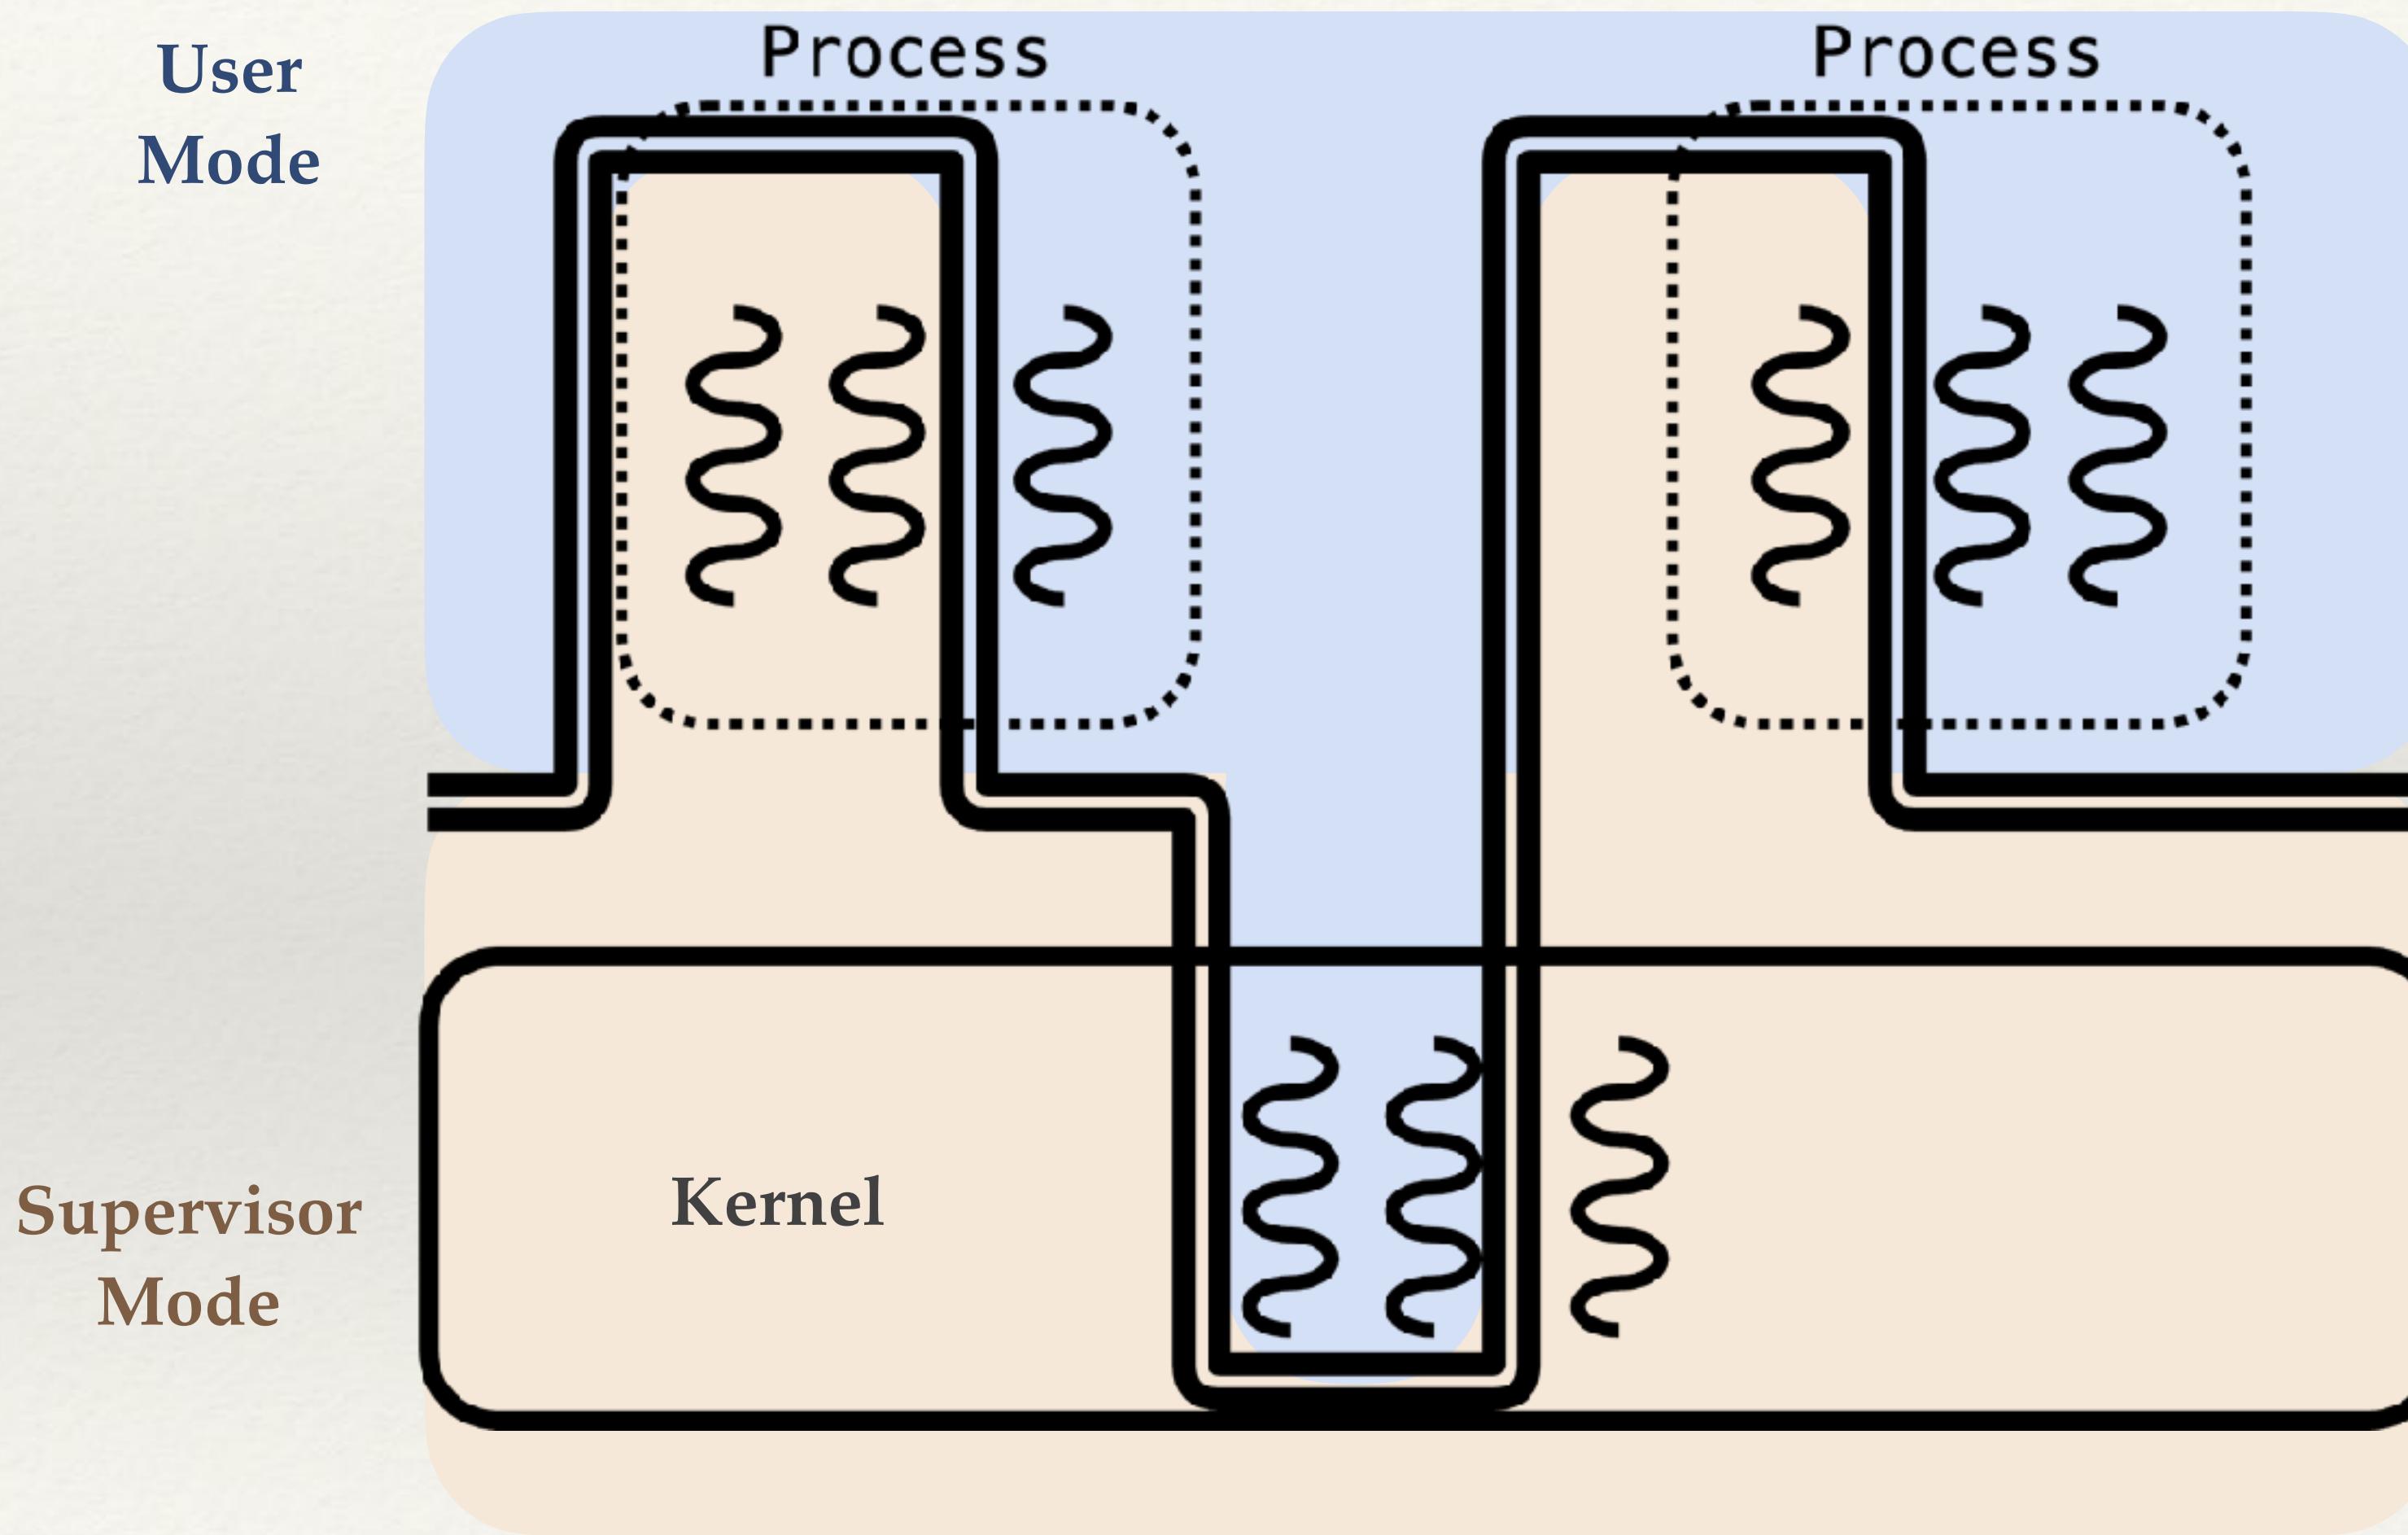
# Dynamic Privilege



# Dynamic Privilege

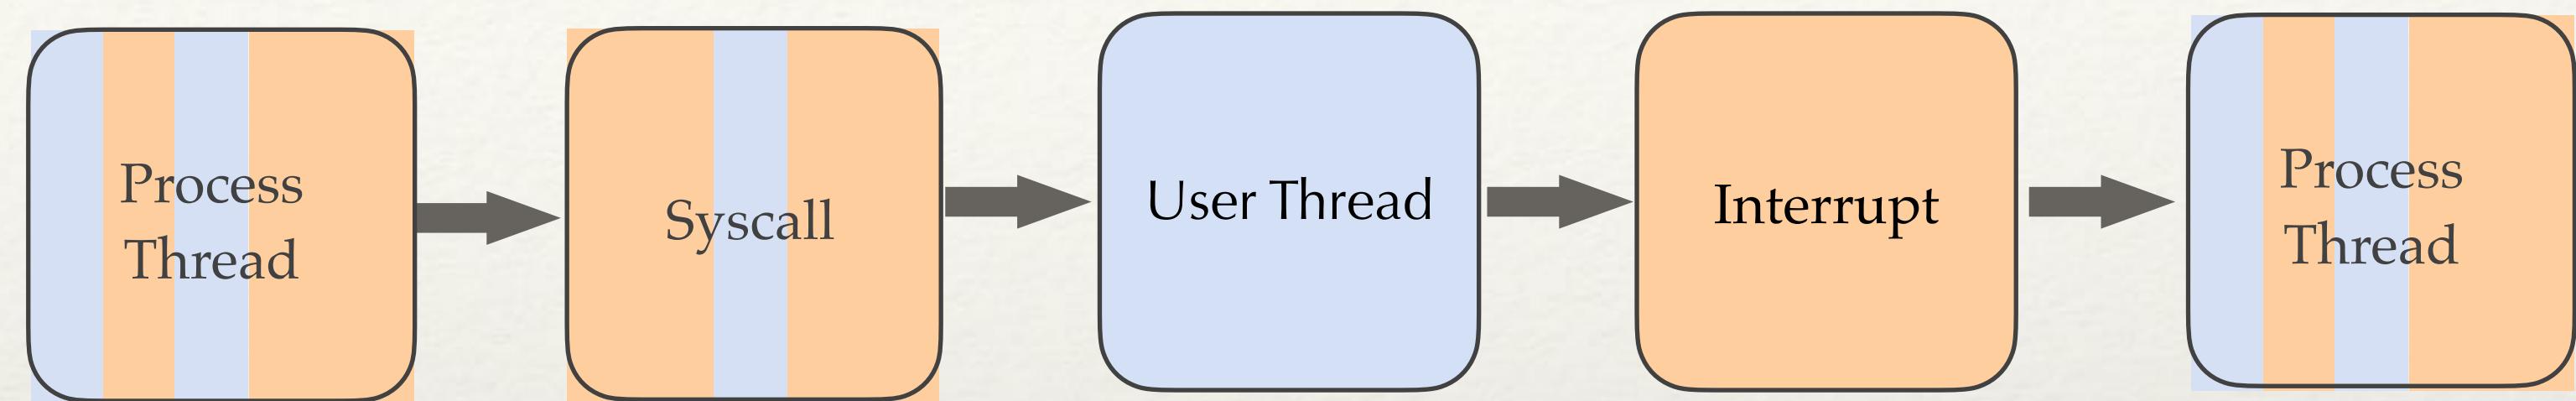


# Dynamic Privilege

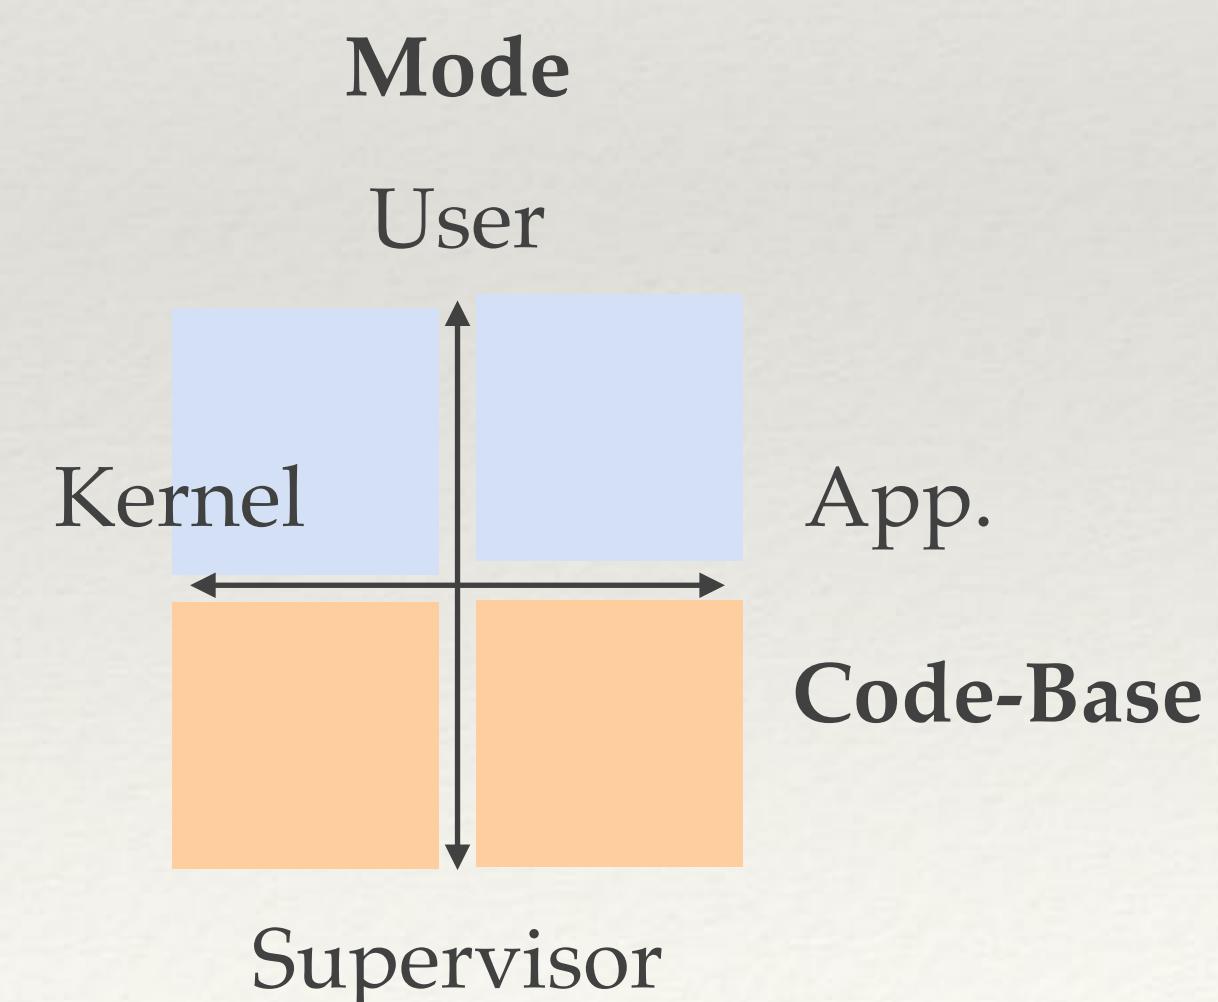


23 Supervisor

# Decoupling Mode and Code-Base

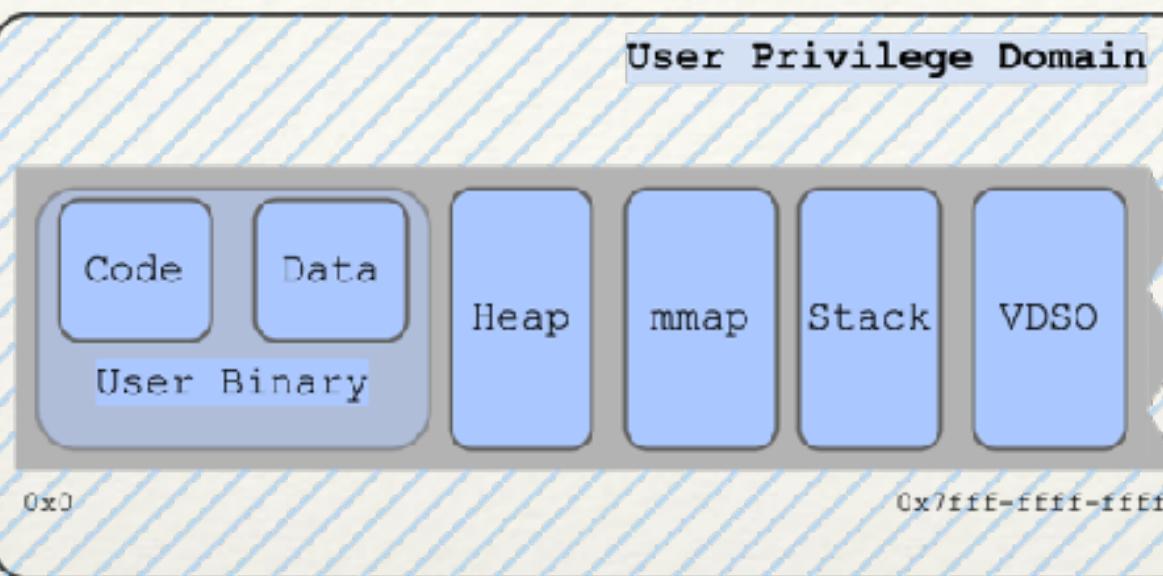


- ❖ Decoupled transitions
  - ❖ Codebase Change  $\neq$  Privilege Level Change

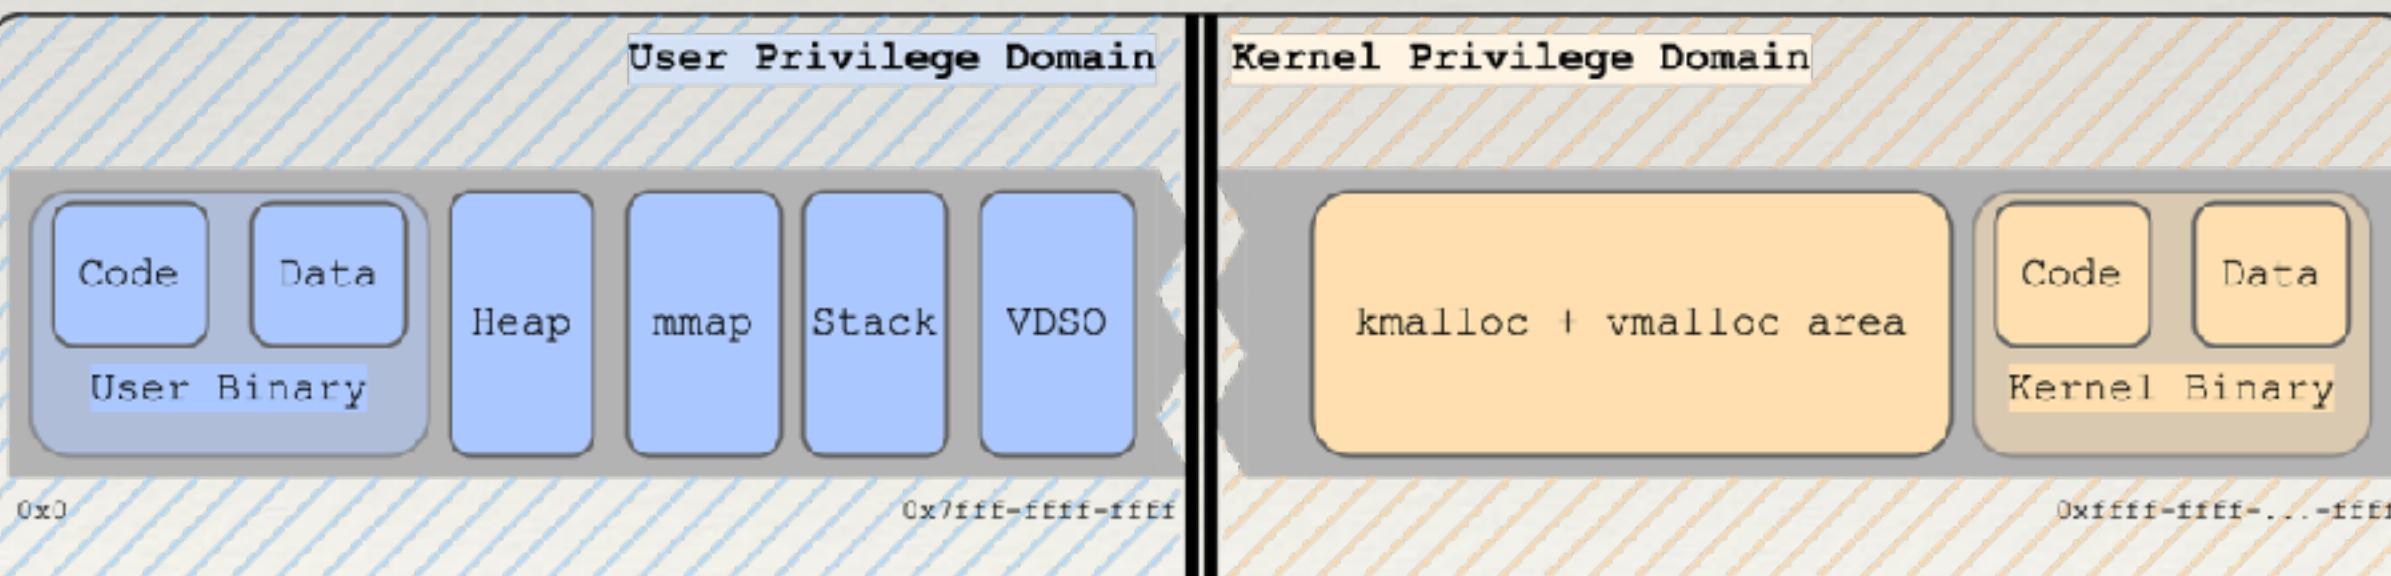


# Virtual Address Space (In Memory)

User No Privilege

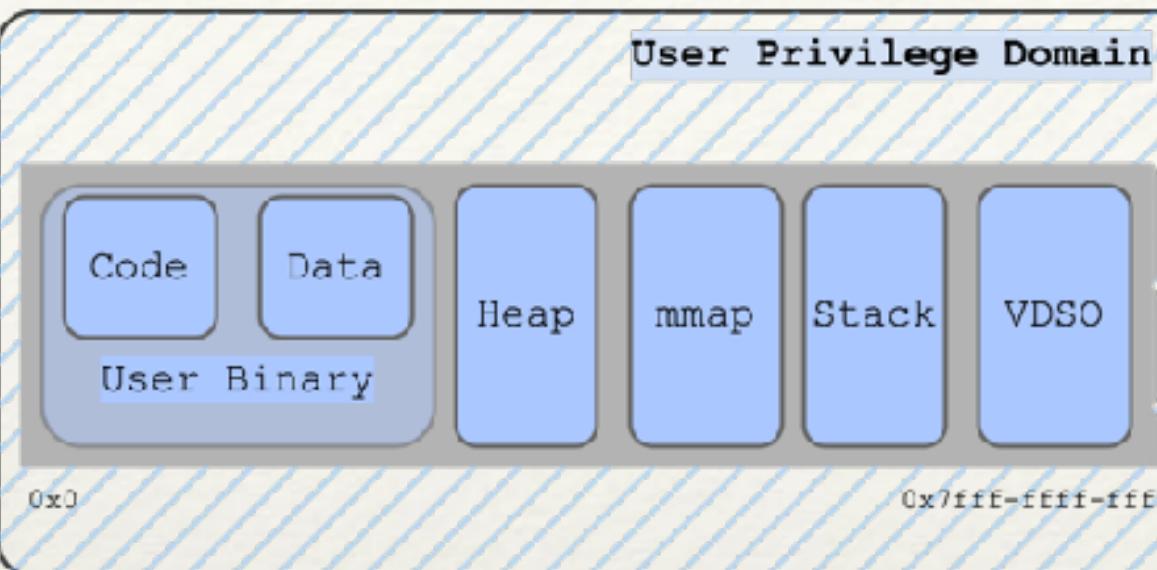


Kernel



# Virtual Address Space (In Memory)

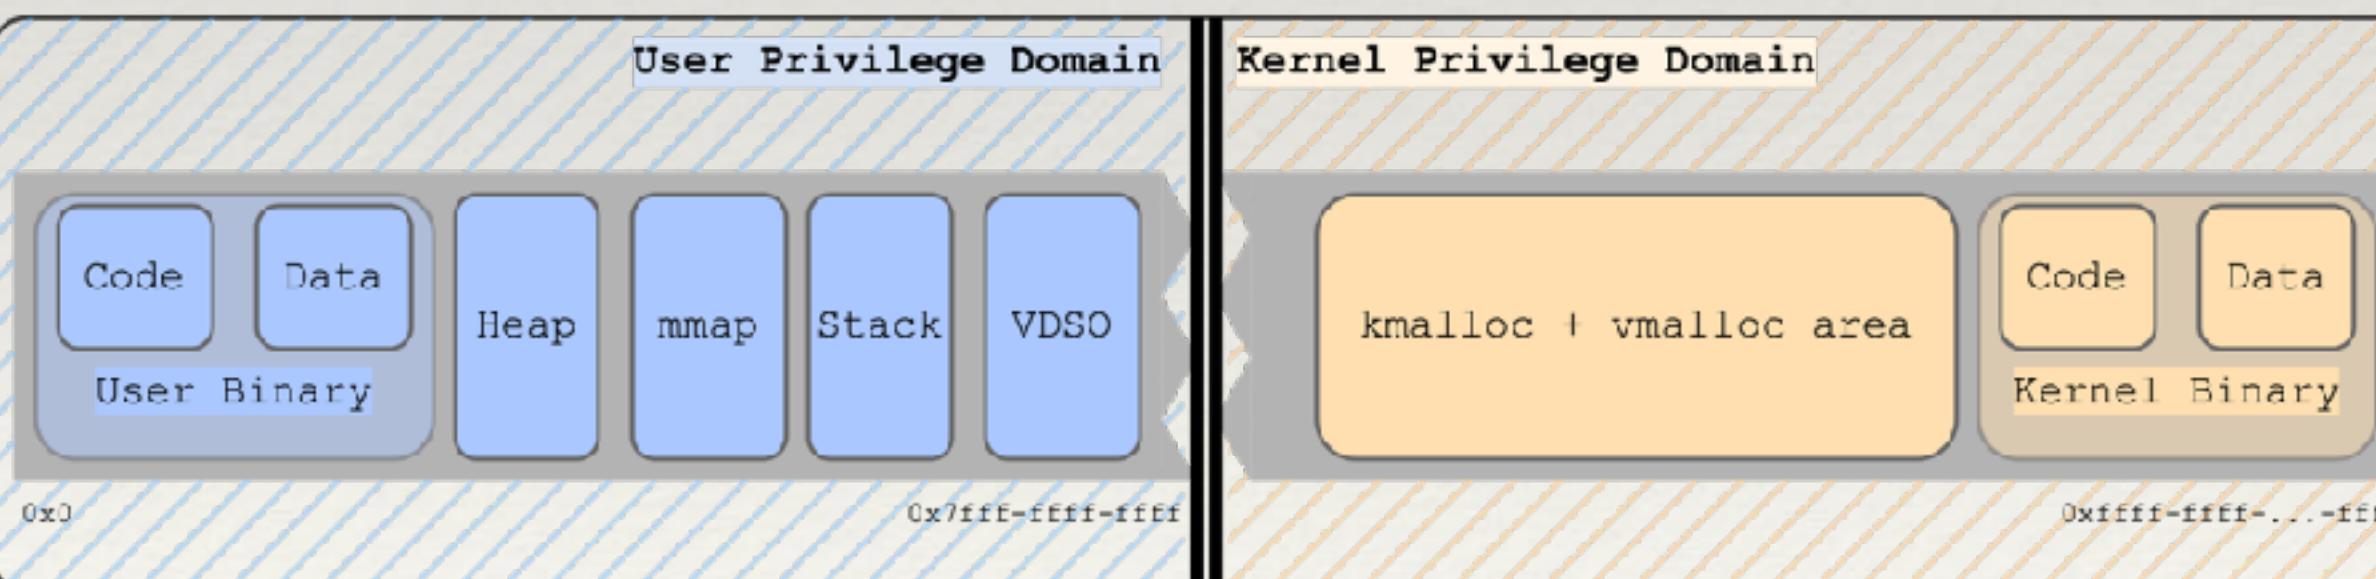
User No Privilege



User With Privilege



Kernel



---

# A Definition

---

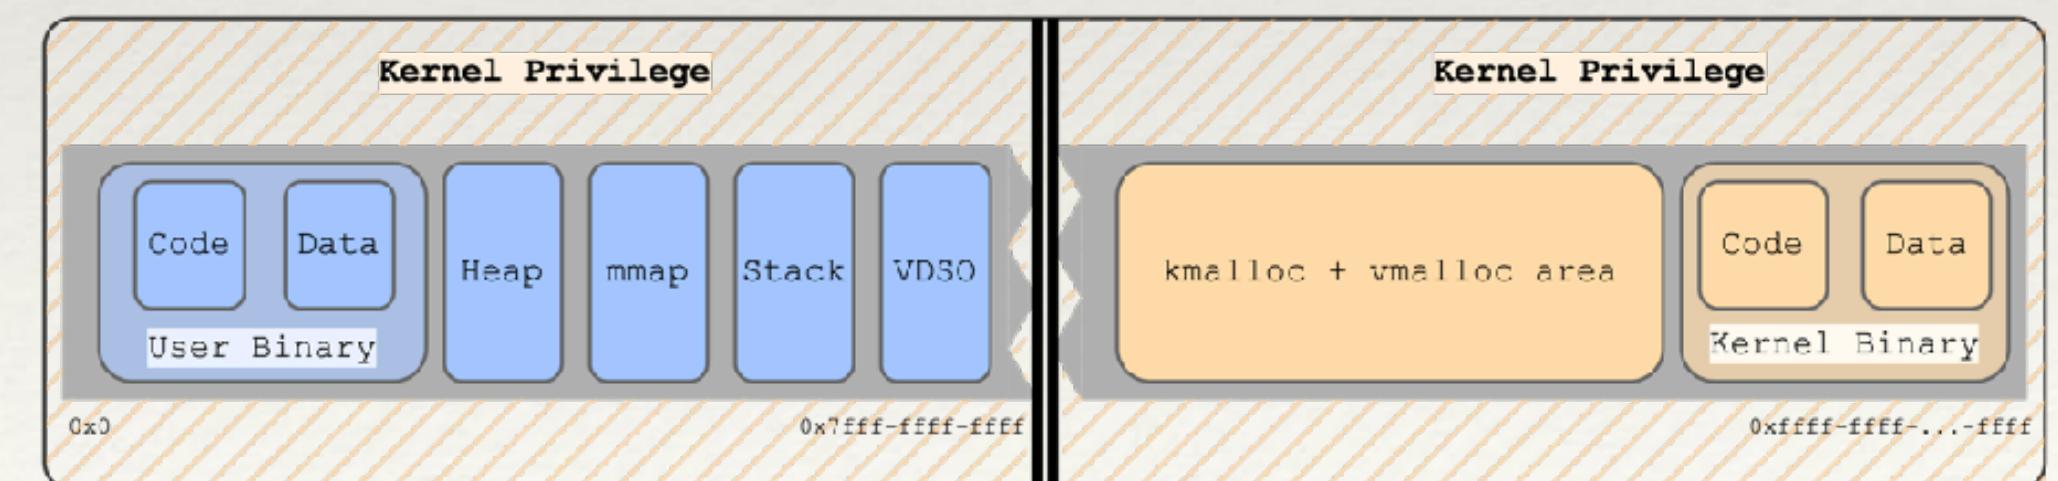
## Dynamic Privilege:

The ability for a unit of execution (e.g. a thread) to independently transition between hardware modes of execution.

# A Definition

## Dynamic Privilege Attributes:

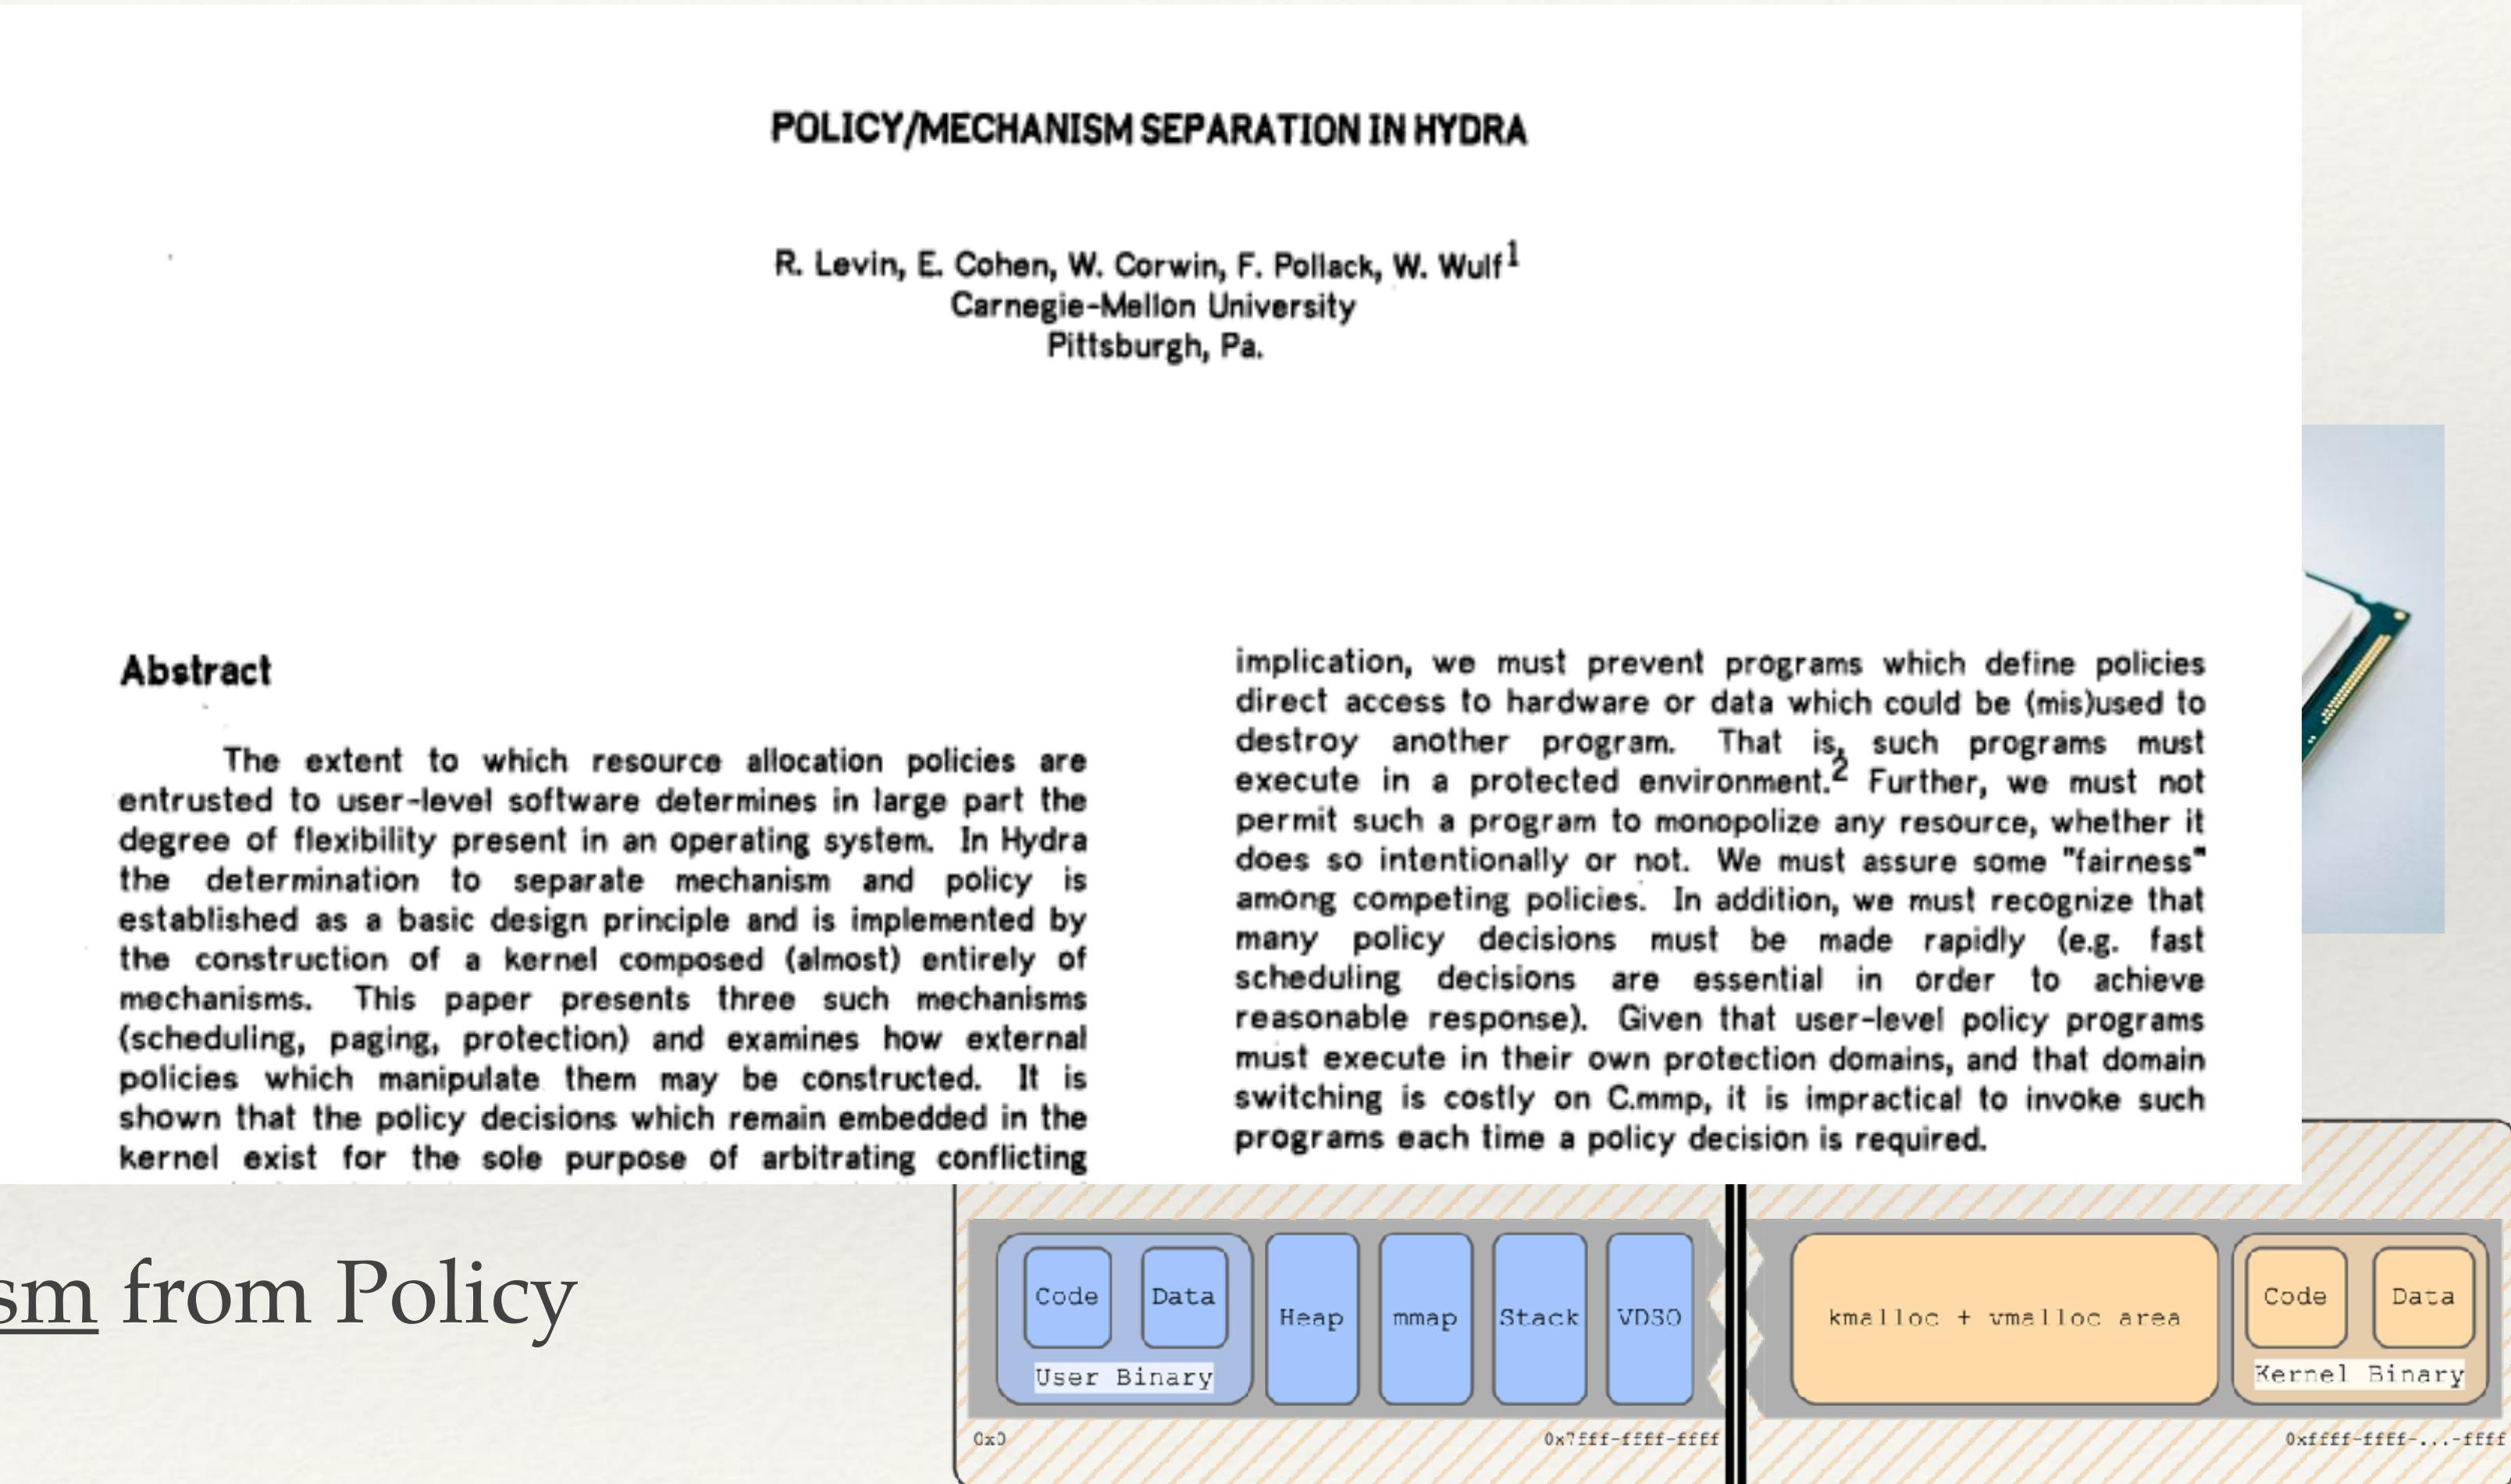
- ❖ Hardware-centric, OS-agnostic
- ❖ Address Space Neutral
- ❖ Separates Mechanism from Policy



# A Definition

## Dynamic Privacy

- ❖ Hardware-centric, Cmmpl
- ❖ Address Space Neutral
- ❖ Separates Mechanism from Policy



# Policy/Mechanism Separation

## Policy / Mechanism Separation in Hydra

The kernel cannot possibly support all conceivable user-defined policies, since some will violate fairness guarantees or protection requirements. At best it can provide a mechanism adequate to implement a large class of desirable resource allocation policies. The decision to exclude certain policies is

# Policy/Mechanism Separation

- ❖ Dynamic Privilege could be used to implement
  - ❖ Static Privilege
  - ❖ Just-in-Time Privilege
  - ❖ On-Demand Privilege
  - ❖ Lazy Privilege

## Policy / Mechanism Separation in Hydra

The kernel cannot possibly support all conceivable user-defined policies, since some will violate fairness guarantees or protection requirements. At best it can provide a mechanism adequate to implement a large class of desirable resource allocation policies. The decision to exclude certain policies is

---

# Implementing Dynamic Privilege

---

- ❖ Many options
  - ❖ Add a new instruction to hardware
  - ❖ Implemented a new OS
  - ❖ Enabling this in an existing OS

# Intersecting Research

## Kernel architecture

- Monolith
- Microkernel
- Exokernel
- Anykernel
- Unikernel
- SAOSs

## Existing OS Mechanisms

- Kernel Modules
- FUSE
- VMMs
- e-BPF

## Capabilities

- Mungi
- Grasshopper
- Protected Shared Libraries

## OS-Level Access

- EbbRT
- UKL
- Dynamic Privilege
- Unix Process

## Control Flow



## Application Personality

## Extension

- K42
- Vino
- Code Downloading
- Software Fault Isolation
- Hardware Transactions
- OO OSs

## Playing with privilege

- Kernel Mode Linux
- Dune
- Privbox
- Lupine
- UKL

---

# Thesis Statement

---

We introduce Dynamic Privilege, an OS model enabling threads to adjust their hardware privilege levels on the fly, challenging the traditional separation between kernel and application.

Dynamic Privilege provides granular control and requires only minor modifications to support it in an existing OS.

Further, Dynamic Privilege enables new system evolution and optimization while maintaining compatibility with existing software ecosystems.

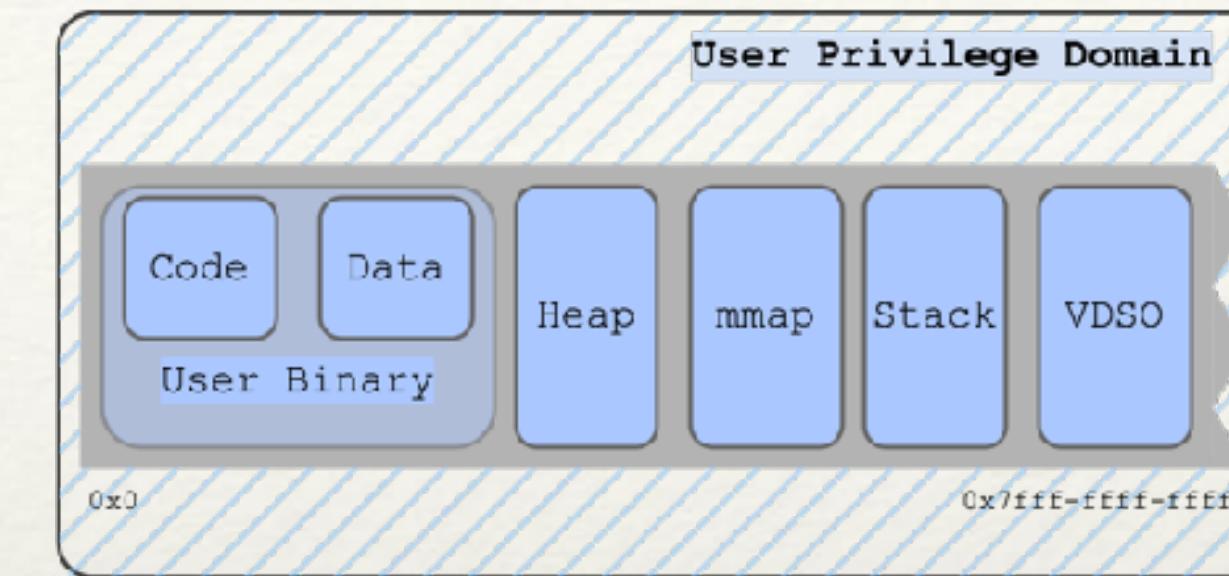
# Ch 3: Demonstrating Low-Level and OS-Level Access

---

- ❖ Implementing Dynamic Privilege as system call, kElevate
- ❖ Demonstrating Access
  - ❖ Low-Level (Hardware)
  - ❖ Runtime OS (Symbol Table)
  - ❖ Compile Time OS (Macros / Types / Inline-Functions)

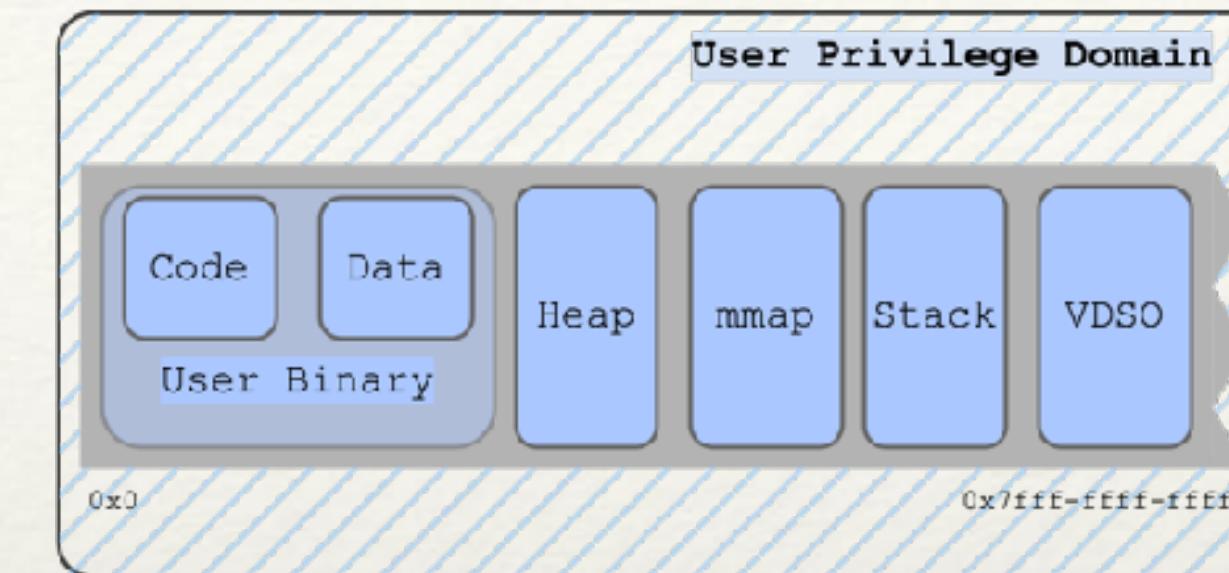
# Ch 3: Demonstrating Low-Level and OS-Level Access

- ❖ Implementing Dynamic Privilege as system call, kElevate
- ❖ Demonstrating Access
  - ❖ Low-Level (Hardware)
  - ❖ Runtime OS (Symbol Table)
  - ❖ Compile Time OS (Macros / Types / Inline-Functions)



# Ch 3: Demonstrating Low-Level and OS-Level Access

- ❖ Implementing Dynamic Privilege as system call, kElevate



- ❖ Demonstrating Access

- ❖ Low-Level (Hardware)



- ❖ Runtime OS (Symbol Table)

- ❖ Compile Time OS (Macros / Types / Inline-Functions)



# Quick Implementation

- ❖ While this could be implemented in hardware...
- ❖ Until then, we'll emulate it with a system call
- ❖ kElevate() is the name of our syscall, the **mechanism** we use to implement Dynamic Privilege
- ❖ It's about as fast as our lowest latency syscall  
getppid() ~100ns

```
static long do_kElevate(int target_mode){  
    return syscall(NR_ELEVATE_SYSCALL, target_mode);  
}  
  
long k_elevate(){  
    return do_kElevate(SYM_ELEVATE_FLAG);  
}  
  
long k_lower(){  
    return do_kElevate(SYM_LOWER_FLAG);  
}  
  
long sym_check_elevate(){  
    return do_kElevate(SYM_QUERY_FLAG);  
}
```

Using kElevate to access and relinquish hardware privilege

# Quick Implementation

- ❖ The entire kElevate patch for Linux (x86\_64) is 173 lines of new kernel code.
  - ❖ What does this really enable?
  - ❖ Adaptors
- ❖ Developed for the x86\_64 architecture
- ❖ ARM64 prototype



```
static long do_kElevate(int target_mode){  
    return syscall(NR_ELEVATE_SYSCALL, target_mode);  
}  
  
long k_elevate(){  
    return do_kElevate(SYM_ELEVATE_FLAG);  
}  
  
long k_lower(){  
    return do_kElevate(SYM_LOWER_FLAG);  
}  
  
long sym_check_elevate(){  
    return do_kElevate(SYM_QUERY_FLAG);  
}
```

Using kElevate to access and relinquish hardware privilege

# Similar to Kernel Modules Security Model

- ❖ The only thing protecting system calls from unprivileged users is a check w/ the capability system
  - ❖ systemctl, mount, insmod
- ❖ We introduce a use a very similar check to determine if the executable has the corresponding capability

```
if (!capable(CAP_SYS_ADMIN))  
    return -EPERM;
```

```
1095     struct SymbiReg sreg;  
1096     sreg.raw = flags;  
1097  
1098     if (!capable(CAP_SYMBI_ELEV))|  
1099         return -EPERM;  
1100
```



# kElevate, the Gory Details

- ❖ Kernel Changes:

- ❖ Add system call kElevate()
- ❖ Add kernel config KELEVATE
- ❖ Add conditional system call return path
  - ❖ Modified IRET frame to return with privilege
    - ❖ Alternative RET implementation
  - ❖ Add 2 bits to the (per thread) task\_struct
    - ❖ Elevated / Migrated
  - ❖ Migration path
    - ❖ Update GS register to new CPU
- ❖ Boot Arguments:
  - ❖ No SMEP / No SMAP / No KASLR

```
static long do_kElevate(int target_mode){  
    return syscall(NR_ELEVATE_SYSCALL, target_mode);  
}  
  
long k_elevate(){  
    return do_kElevate(SYM_ELEVATE_FLAG);  
}  
  
long k_lower(){  
    return do_kElevate(SYM_LOWER_FLAG);  
}  
  
long sym_check_elevate(){  
    return do_kElevate(SYM_QUERY_FLAG);  
}
```

Using kElevate to access and relinquish hardware privilege



---

# Low-Level Access

---

- ❖ With Dynamic Privilege, all we get is hardware access
  - ❖ Bracketing a thread's privilege access in time
  - ❖ Accessing / Executing:
    - ❖ instructions
    - ❖ registers
    - ❖ memory

# Low-Level Access

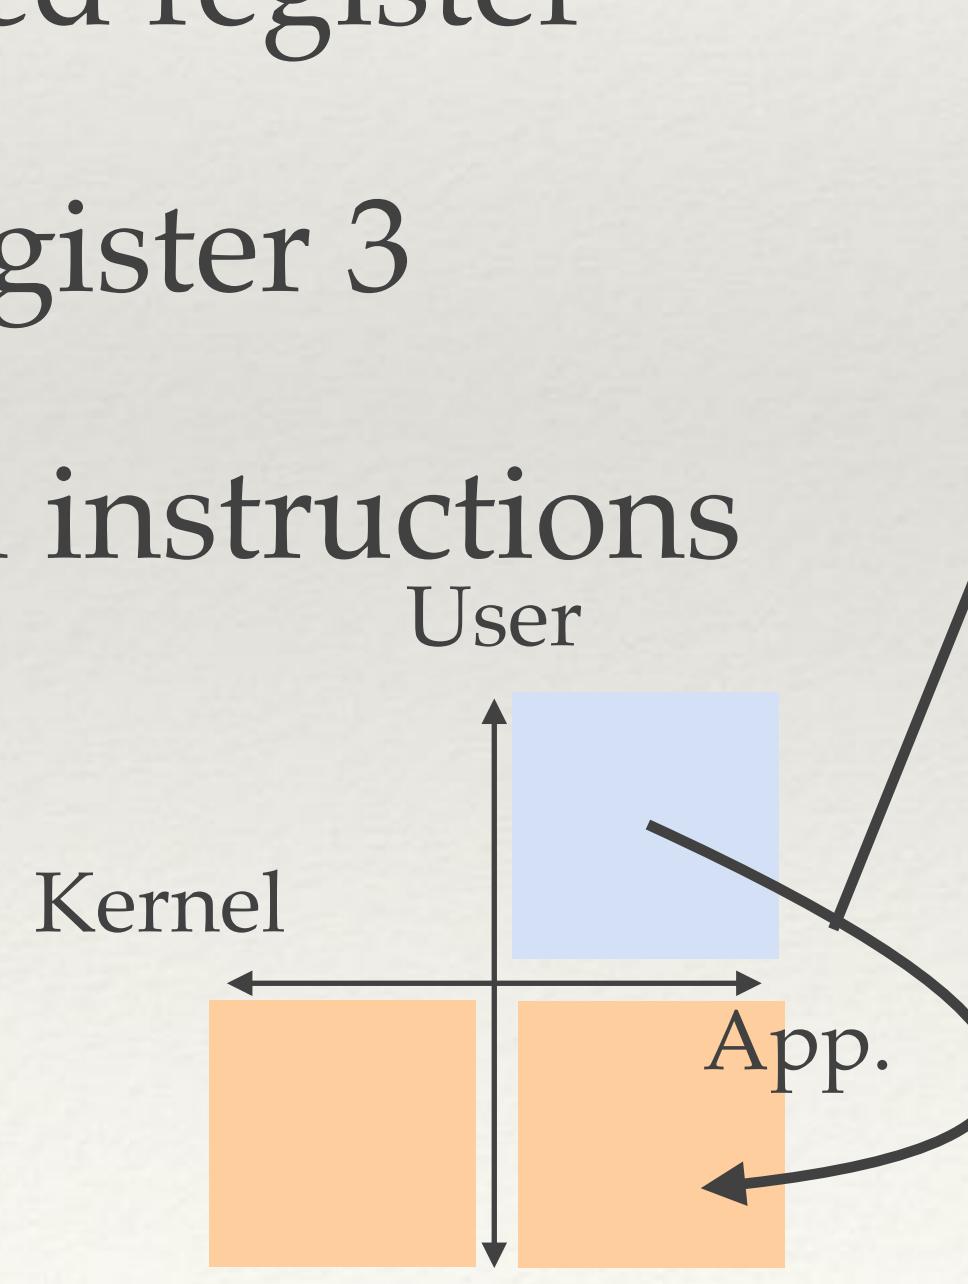
- ❖ Bracketing privilege access in time
- ❖ Accessing a privileged register
  - ❖ x86\_64 Control Register 3
- ❖ Executing privileged instructions
  - ❖ `mov cr3, %rax`

```
1 priv = loadMod(priv_lib_path)
2 kele = loadMod(kele_lib_path)
3 kele.kElevate()
4 cr3 = priv.getcr3()
5 kele.kLower()
6 print(f"CR3 = 0x{cr3:x}")
```

```
1 CR3 = 0x2e868002
```

# Low-Level Access

- ❖ Bracketing privilege access in time
- ❖ Accessing a privileged register
  - ❖ x86\_64 Control Register 3
- ❖ Executing privileged instructions
  - ❖ `mov cr3, %rax`

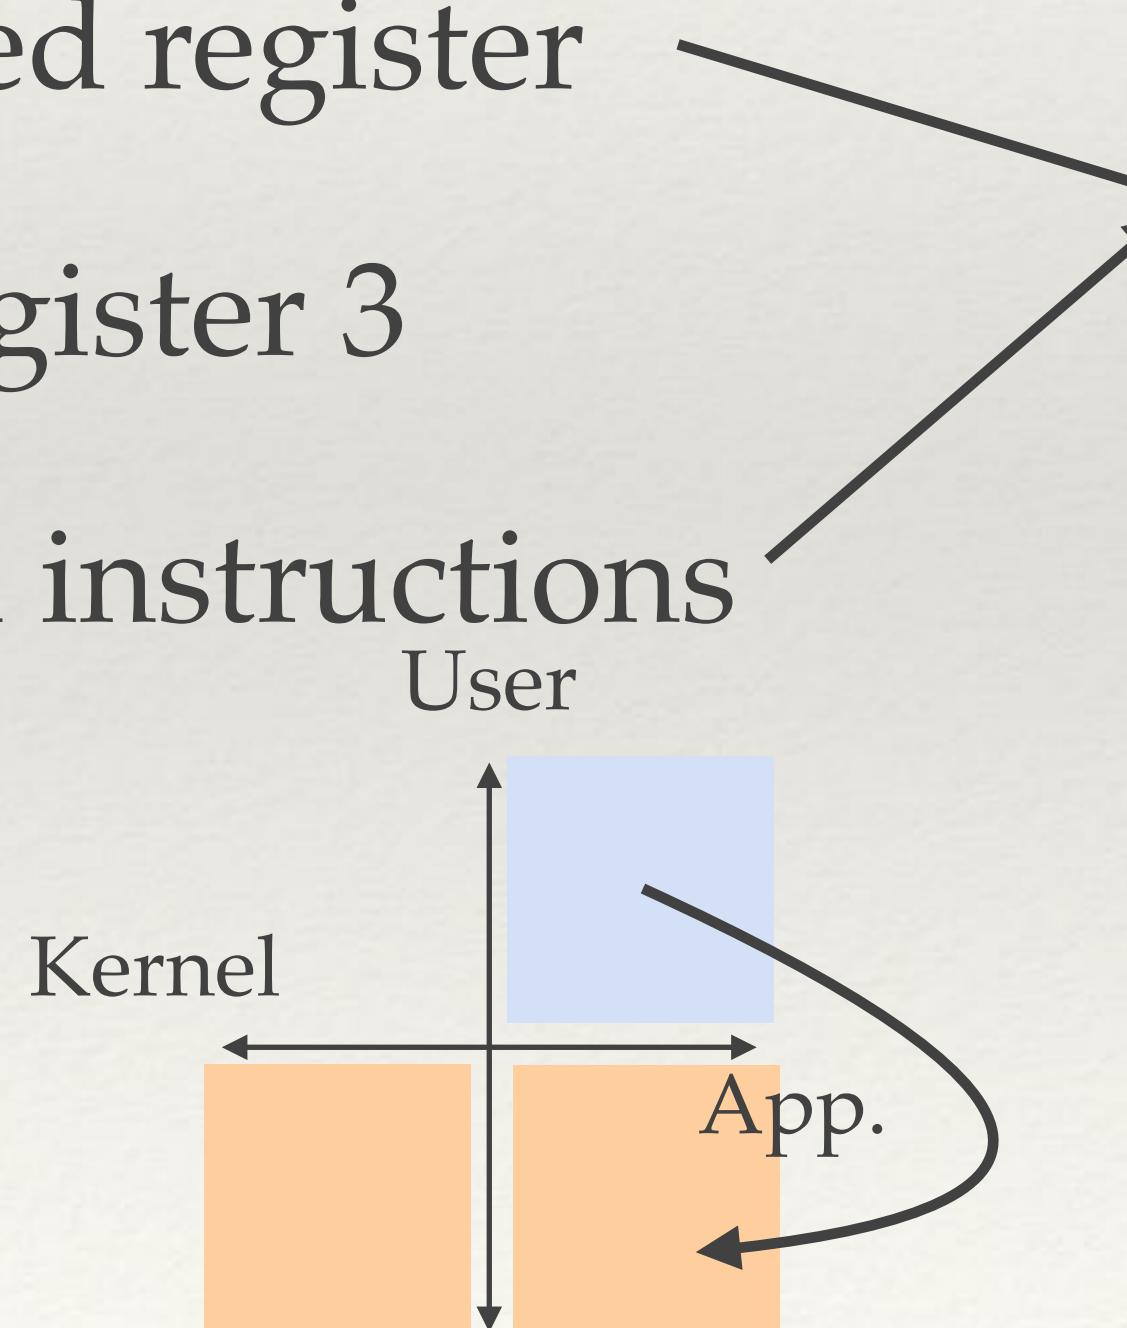


```
1 priv = loadMod(priv_lib_path)
2 kele = loadMod(kele_lib_path)
3 kele.kElevate()
4 cr3 = priv.getcr3()
5 kele.kLower()
6 print(f"CR3 = 0x{cr3:x}")
```

```
1 CR3 = 0x2e868002
```

# Low-Level Access

- ❖ Bracketing privilege access in time
- ❖ Accessing a privileged register
  - ❖ x86\_64 Control Register 3
- ❖ Executing privileged instructions
  - ❖ `mov cr3, %rax`

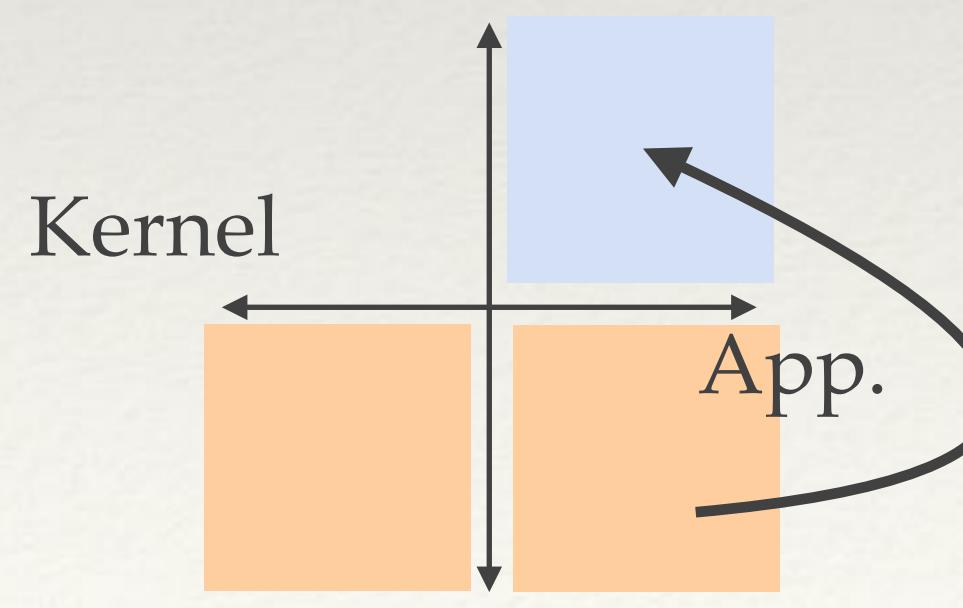


```
1 priv = loadMod(priv_lib_path)
2 kele = loadMod(kele_lib_path)
3 kele.kElevate()
4 cr3 = priv.getcr3()
5 kele.kLower()
6 print(f"CR3 = 0x{cr3:x}")
```

```
1 CR3 = 0x2e868002
```

# Low-Level Access

- ❖ Bracketing privilege access in time
- ❖ Accessing a privileged register
  - ❖ x86\_64 Control Register 3
- ❖ Executing privileged instructions
  - ❖ `mov cr3, %rax`



```
1 priv = loadMod(priv_lib_path)
2 kele = loadMod(kele_lib_path)
3 kele.kElevate()
4 cr3 = priv.getcr3()
5 kele.kLower()
6 print(f"CR3 = 0x{cr3:x}")
```

```
1 CR3 = 0x2e868002
```

# Dynamic Privilege from Anywhere

- ❖ We've mainly prototyped in C, C++, and Python, but have written prototypes to demonstrate kElevate can be used in others
  - ❖ Java
  - ❖ Node
  - ❖ Rust
  - ❖ Go
  - ❖ Assembly

```
In [10]: ret = com.run_cmd("taskset -c 0 idt_tool -g")
addr_old_idt = ret.stdout.splitlines()[0]
# Prepend 0x if not there.
addr_old_idt = hex(int(addr_old_idt, 16))
# Print the old IDT is located at this address

print("Old IDT is located at: ", addr_old_idt)

Old IDT is located at: 0xfffffc90000880f000

In [11]: # Allocate a kernel page, copy the old idt onto it, and return the address of this page
ret = com.run_cmd("taskset -c 0 idt_tool -c")
addr_new_idt = ret.stdout.splitlines()[0]
addr_new_idt = hex(int(addr_new_idt, 16))

print("New IDT is located at: ", addr_new_idt)

New IDT is located at: 0xfffffc90003f57000

In [13]: # Install the new IDT
com.run_cmd("taskset -c 0 idt_tool -i -a " + addr_new_idt)

# Get the currently loaded IDT ptr
ret = com.run_cmd("taskset -c 0 idt_tool -g")
addr_current_idt = ret.stdout.splitlines()[0]
addr_current_idt = hex(int(addr_current_idt, 16))

print("Current IDT is located at: ", addr_current_idt)

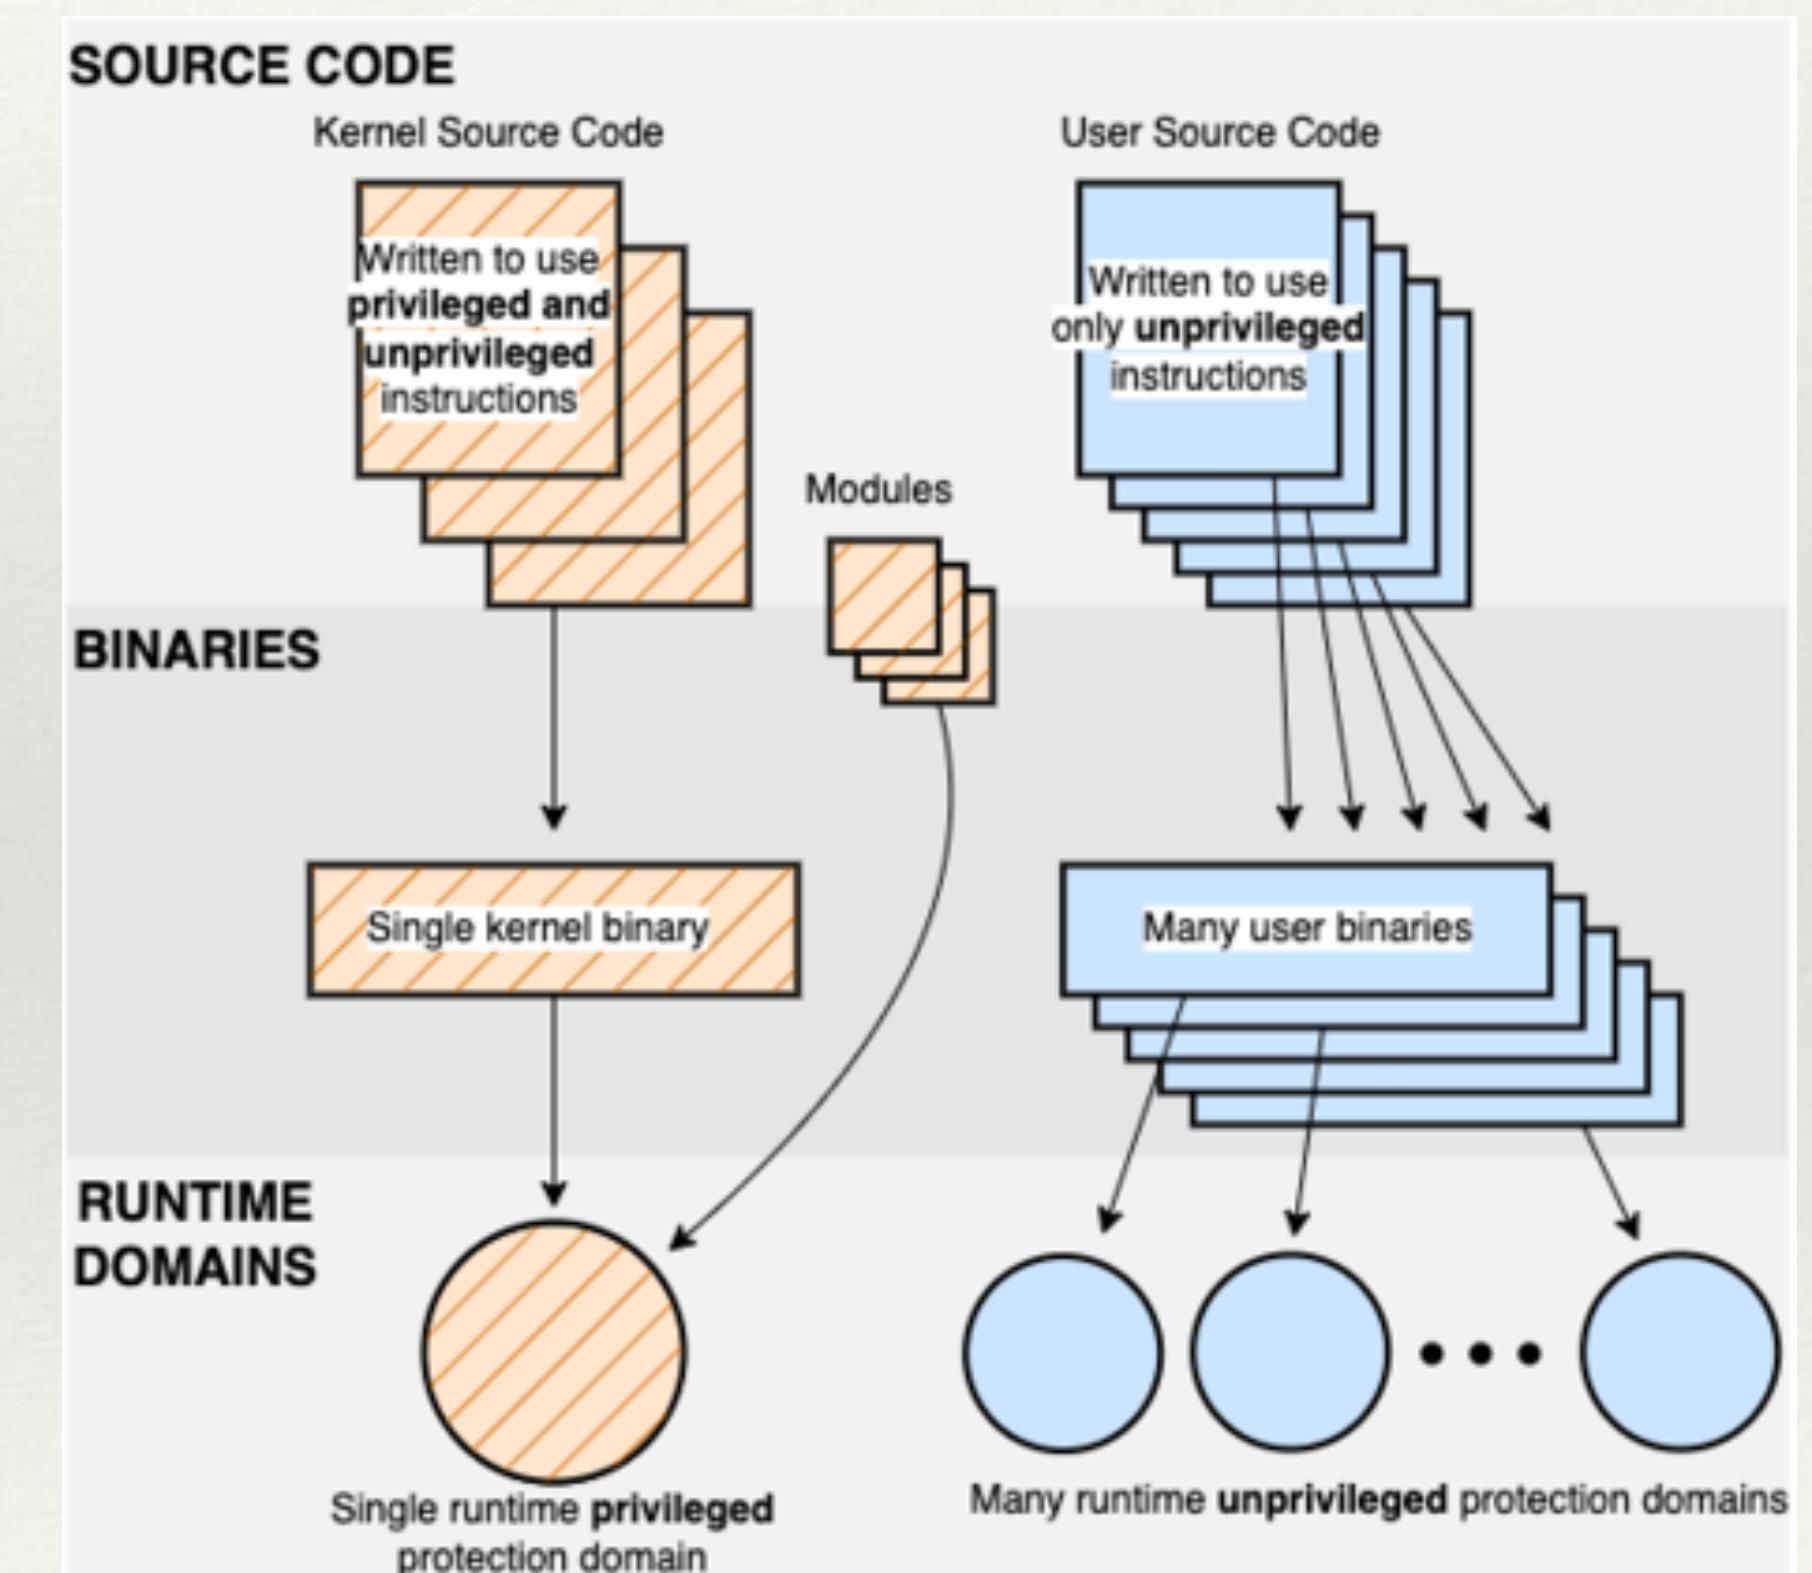
Current IDT is located at: 0xfffffc90003f57000
```

Who says you can't  
hack the kernel from a  
Jupyter Notebook



# OS-Level Access

- ❖ What kernel developers need to get their work done
- ❖ Runtime
  - ❖ Executing standalone kernel functions
  - ❖ Accessing kernel data



# Runtime OS Access

- ❖ Intermixing user and kernel codepaths!

```
→ with open("output.txt", "w") as f:  
    # Standard write.  
    f.write("Hello ")  
  
    kele.kElevate()  
    # Invoking kernel's ksys_write directly.  
    kele.ksys_write(f.fileno(), "World!", 7)  
    kele.kLower()
```

```
:  
    # Open the file and read the contents  
    with open("output.txt", "r") as f:  
        print(f.read())
```

Hello world!

# Runtime OS Access

- ❖ Intermixing user and kernel codepaths!

```
with open("output.txt", "w") as f:  
    # Standard write.  
    f.write("Hello ")  
  
    kele.kElevate()  
    # Invoking kernel's ksys_write directly.  
    kele.ksys_write(f.fileno(), "World!", 7)  
    kele.kLower()
```

```
:  
    # Open the file and read the contents  
    with open("output.txt", "r") as f:  
        print(f.read())
```

Hello world!

# Runtime OS Access

- ❖ Intermixing user and kernel codepaths!

```
with open("output.txt", "w") as f:  
    # Standard write.  
    f.write("Hello ")  
  
    → kele.kElevate()  
    # Invoking kernel's ksys_write directly.  
    kele.ksys_write(f.fileno(), "World!", 7)  
    kele.kLower()
```

```
:  
    # Open the file and read the contents  
    with open("output.txt", "r") as f:  
        print(f.read())
```

Hello world!

# Runtime OS Access

- ❖ Intermixing user and kernel codepaths!
- ❖ Location of ksys\_write() known by kallsyms subsystem (or the symbol table)

```
with open("output.txt", "w") as f:  
    # Standard write.  
    f.write("Hello ")  
  
    kele.kElevate()  
    # Invoking kernel's ksys_write directly.  
    kele.ksys_write(f.fileno(), "World!", 7)  
    kele.kLower()
```

```
:  
    # Open the file and read the contents  
    with open("output.txt", "r") as f:  
        print(f.read())
```

Hello world!

# Runtime OS Access

- ❖ Intermixing user and kernel codepaths!

```
with open("output.txt", "w") as f:  
    # Standard write.  
    f.write("Hello ")  
  
    kele.kElevate()  
    # Invoking kernel's ksys_write directly.  
    kele.ksys_write(f.fileno(), "World!", 7)  
    kele.kLower()
```

```
:  
    # Open the file and read the contents  
    with open("output.txt", "r") as f:  
        print(f.read())
```

Hello world!

# Runtime OS Access

- ❖ Intermixing user and kernel codepaths!

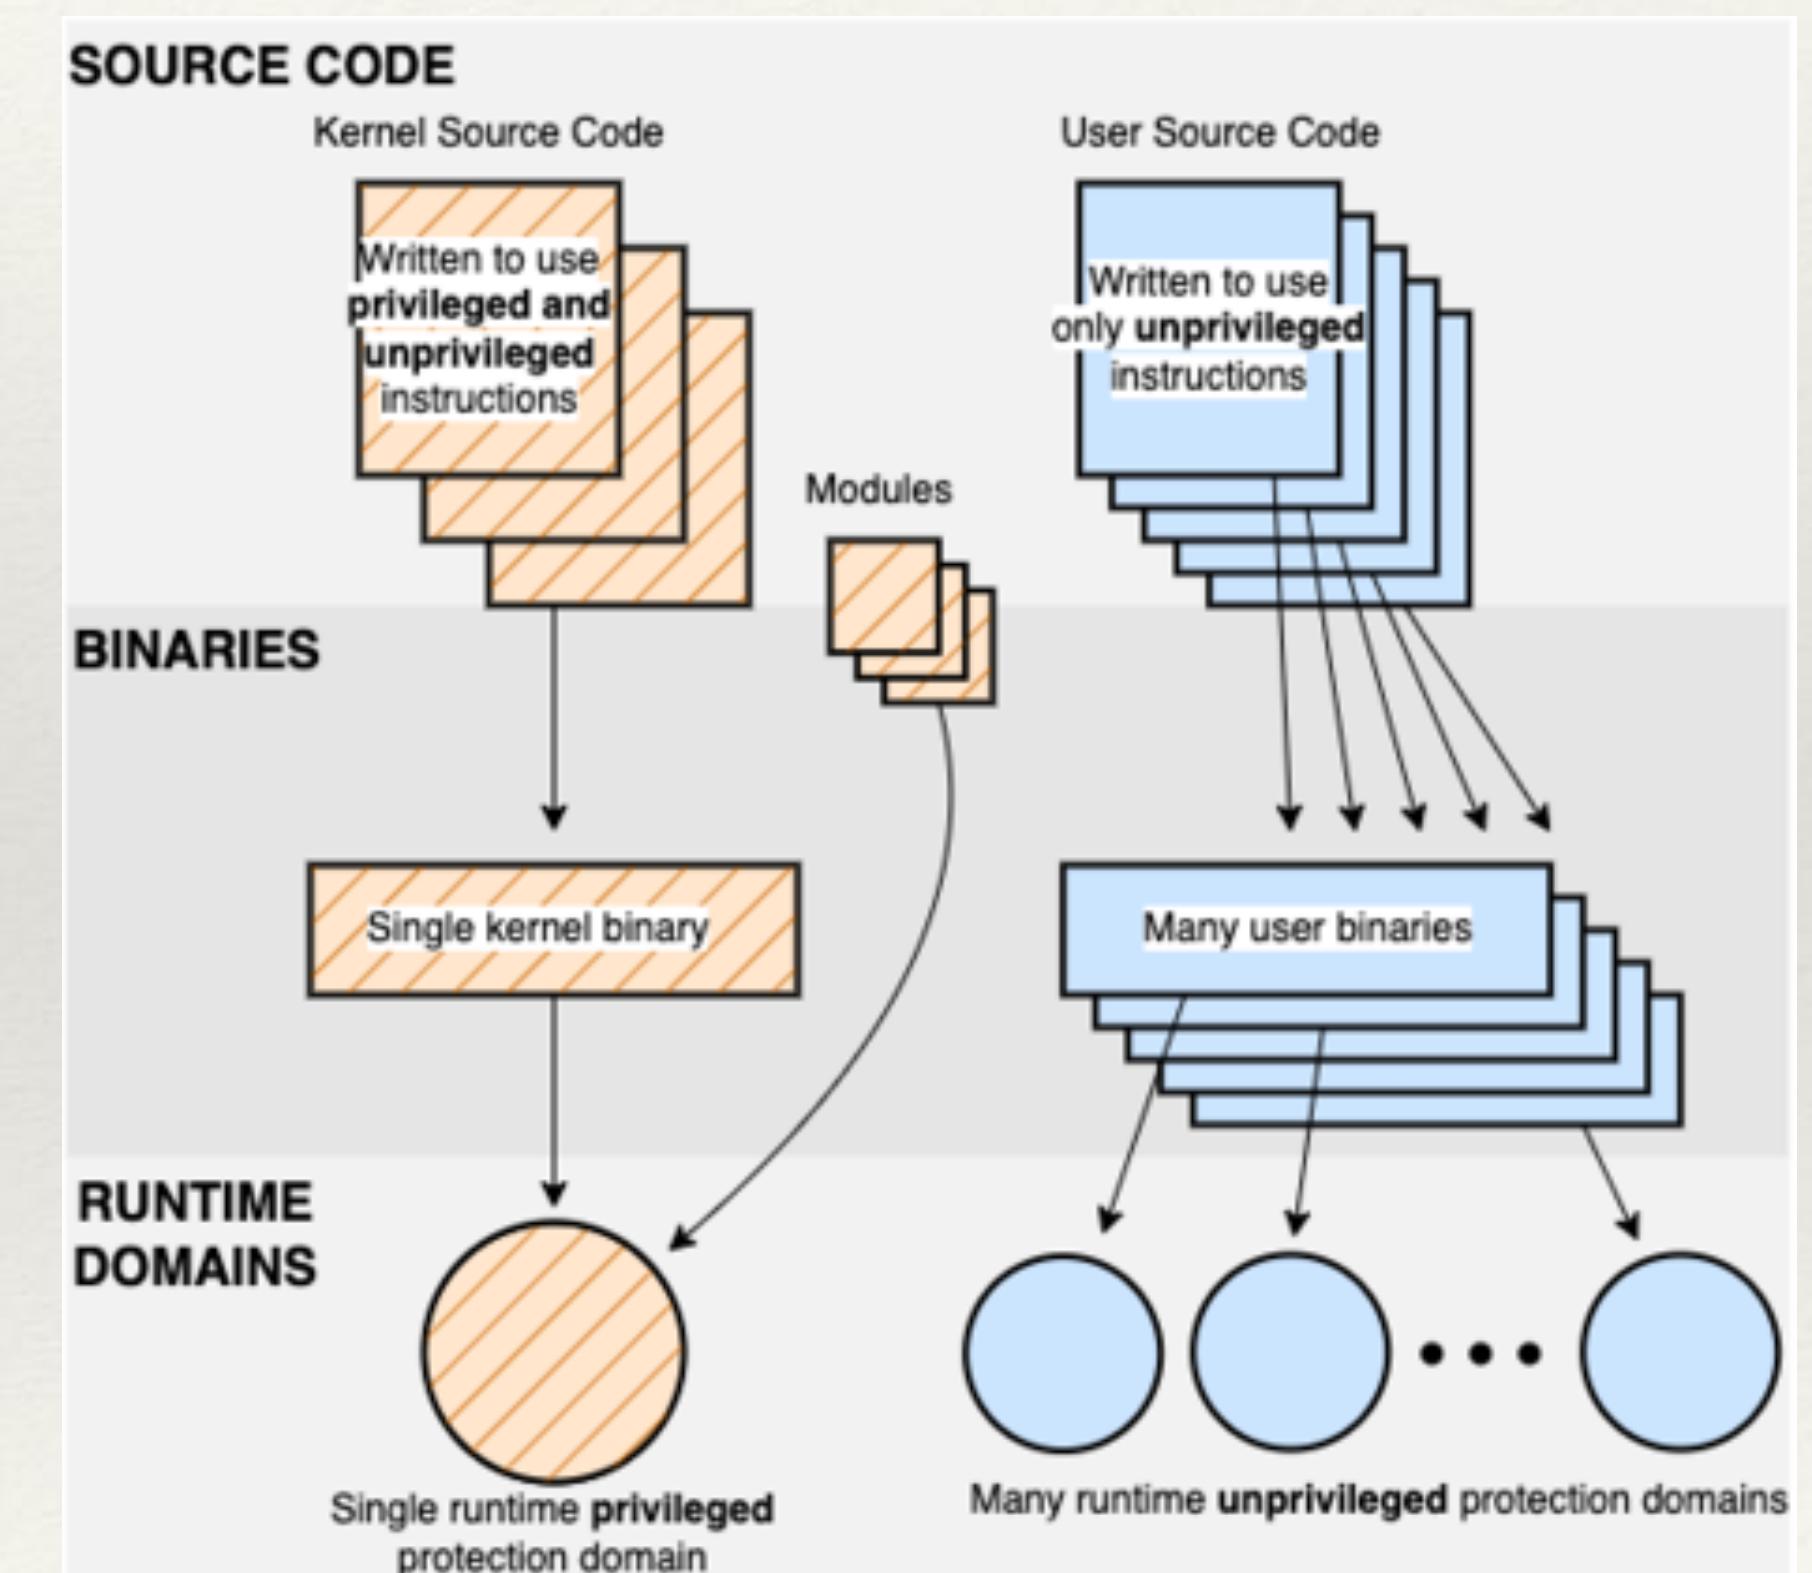
```
with open("output.txt", "w") as f:  
    # Standard write.  
    f.write("Hello ")  
  
    kele.kElevate()  
    # Invoking kernel's ksys_write directly.  
    kele.ksys_write(f.fileno(), "World!", 7)  
    kele.kLower()
```

```
:  
# Open the file and read the contents  
with open("output.txt", "r") as f:  
    print(f.read())
```

Hello world!

# OS-Level Access

- ❖ What kernel developers need to get their work done
  - ❖ Runtime
    - ❖ Executing standalone kernel functions
    - ❖ Accessing kernel data
  - ❖ Compile time
    - ❖ Macros
    - ❖ Types
    - ❖ Inline Functions



# Compile Time OS-Level Access

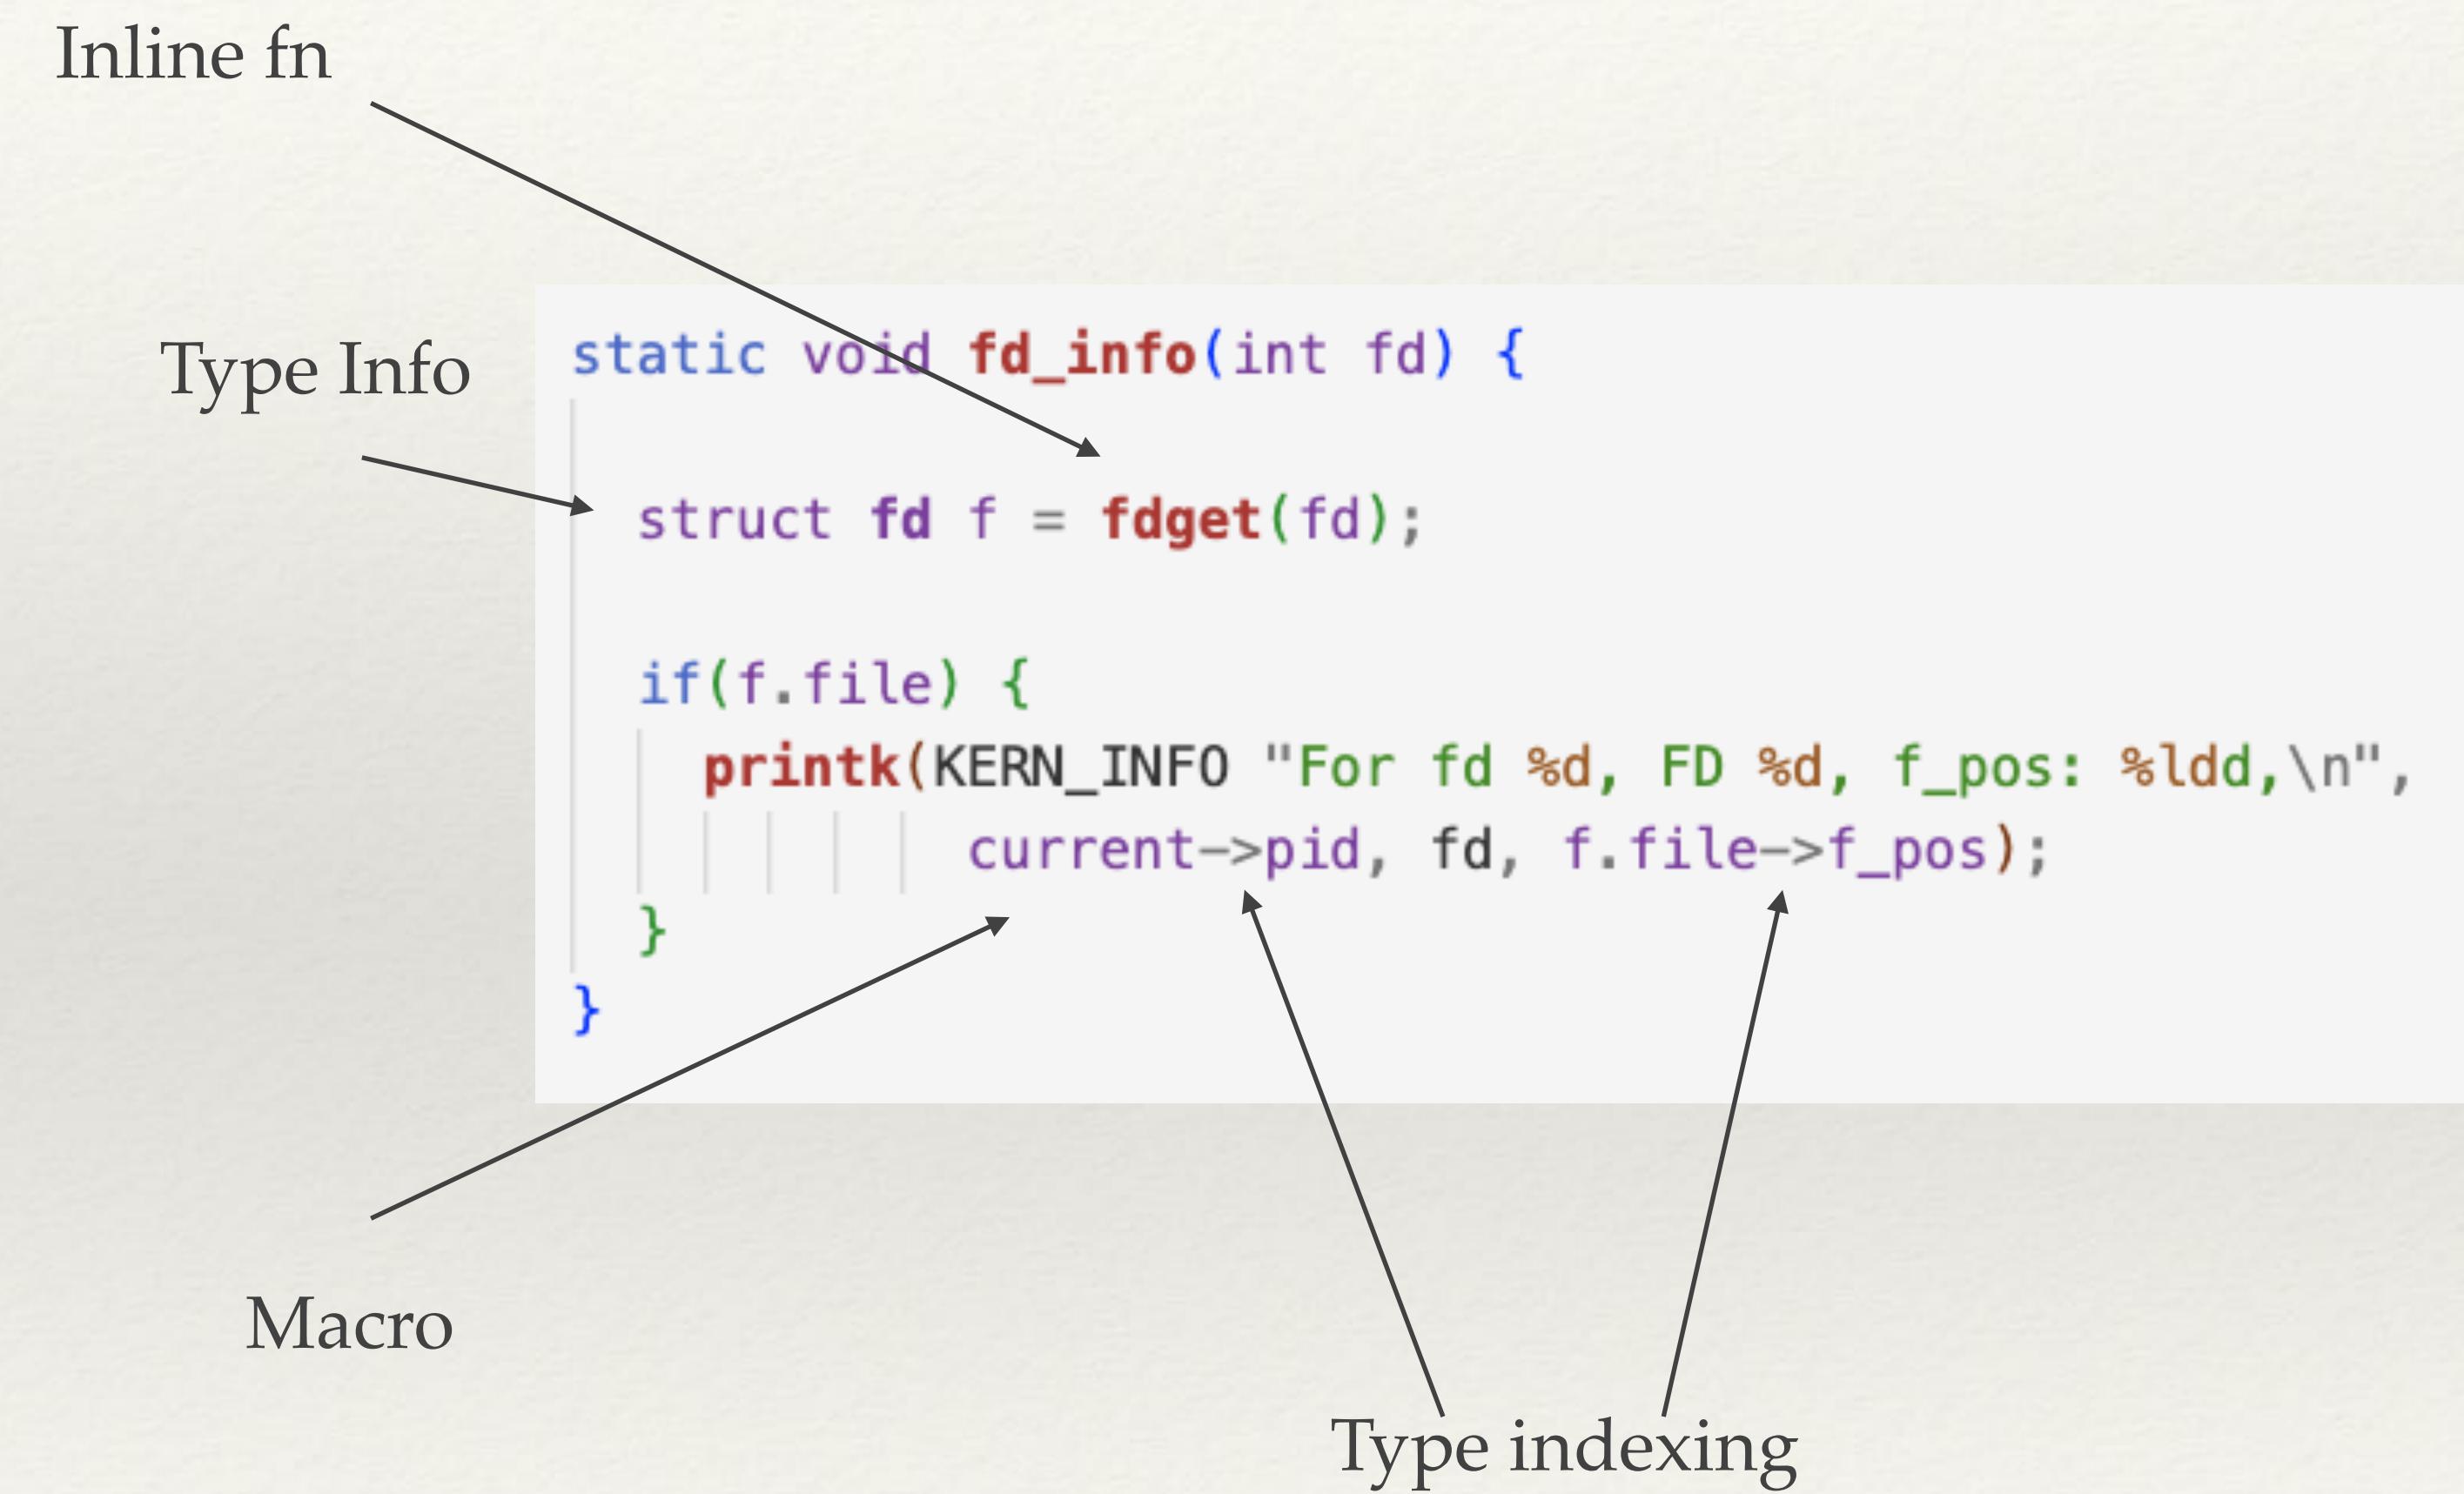
```
static void fd_info(int fd) {  
    struct fd f = fdget(fd);  
  
    if(f.file) {  
        printk(KERN_INFO "For fd %d, FD %d, f_pos: %lld,\n",  
              current->pid, fd, f.file->f_pos);  
    }  
}
```

Inline fn

Type Info

Macro

Type indexing



# Compile Time OS-Level Access

- ❖ How: Utilizing the kernel's “out-of-tree” kernel module build system
- ❖ Building this first would have knocked a year off my phd

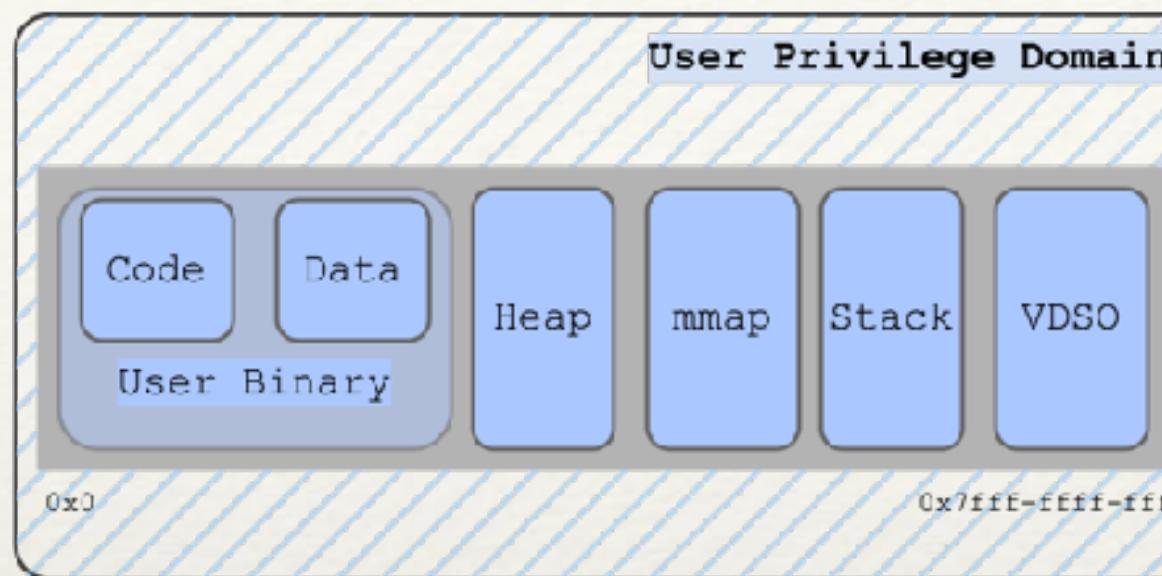
The diagram shows a snippet of kernel code with several annotations:

- Inline fn**: An arrow points from this label to the function definition `static void fd_info(int fd) {`.
- Type Info**: An arrow points from this label to the variable declaration `struct fd f = fdget(fd);`.
- Macro**: An arrow points from this label to the macro call `current->pid` in the `printf` statement.
- Type indexing**: Two arrows point from this label to the type index `f.file` and the type index `f_pos` in the `printf` statement.

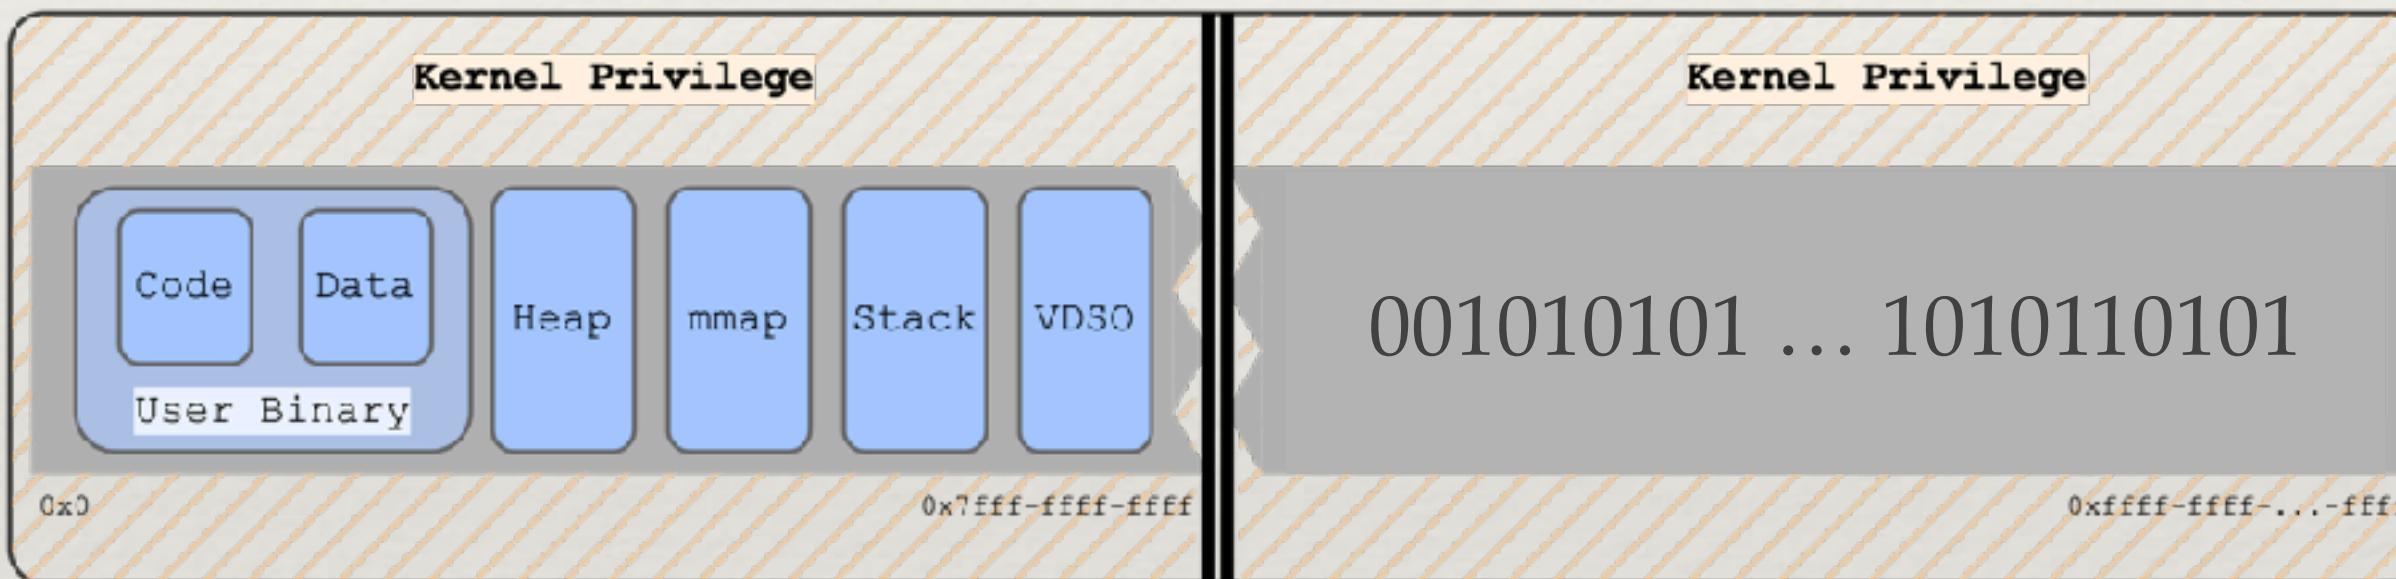
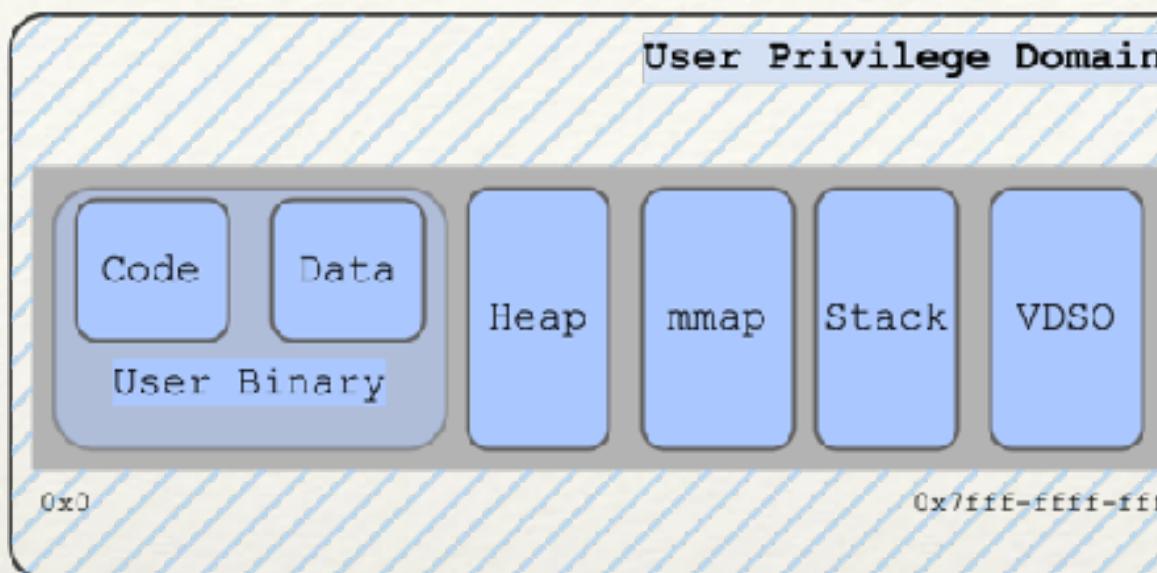
```
static void fd_info(int fd) {
    struct fd f = fdget(fd);

    if(f.file) {
        printk(KERN_INFO "For fd %d, FD %d, f_pos: %ld,\n",
               current->pid, fd, f.file->f_pos);
    }
}
```

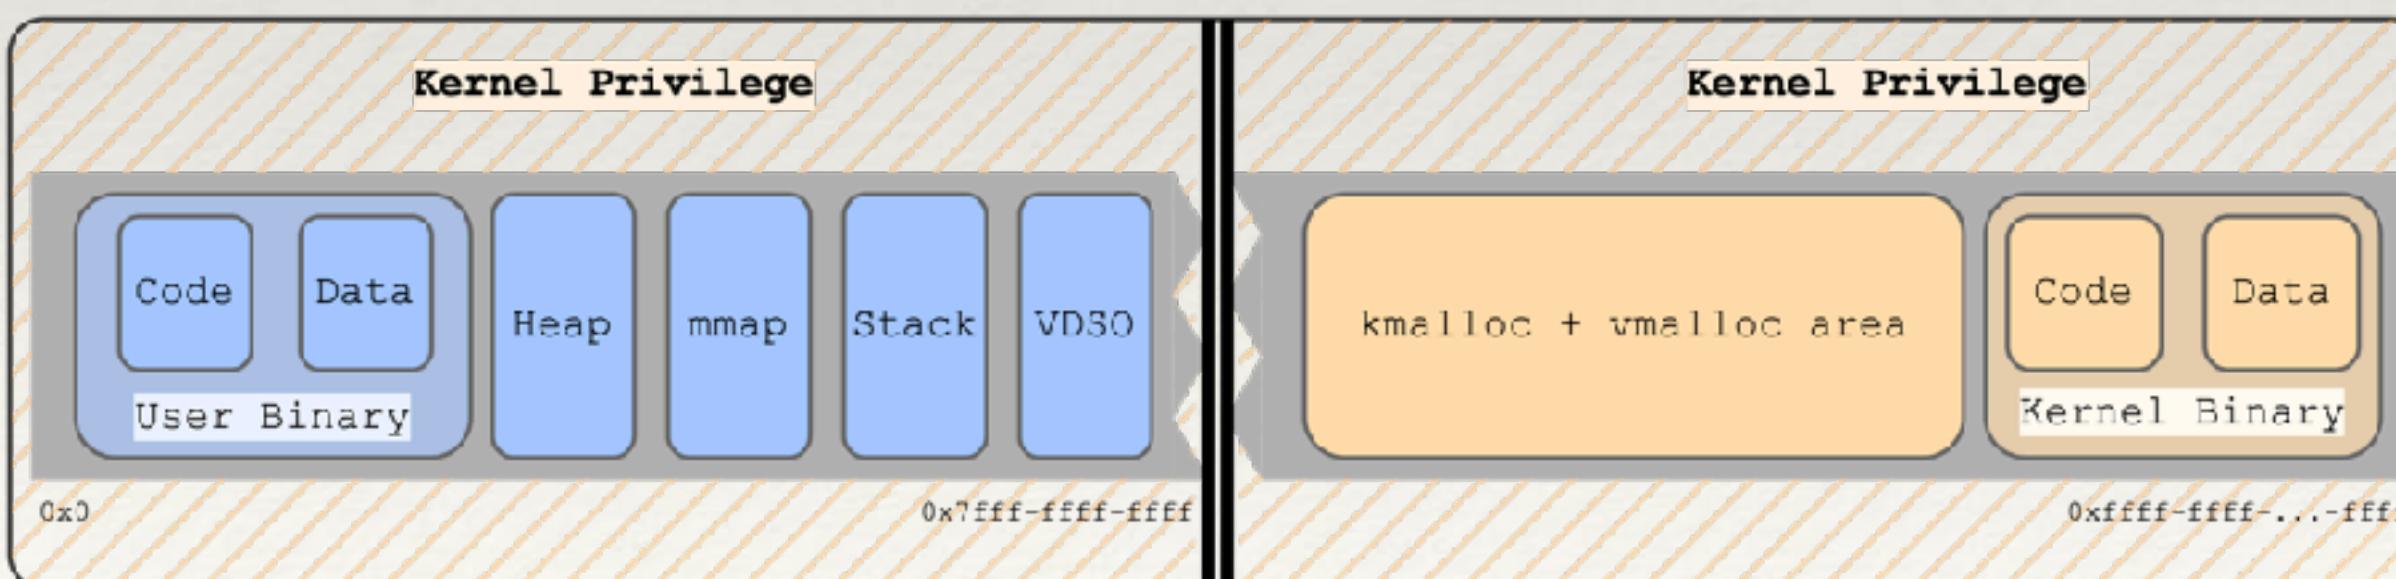
# Low-Level to OS-Level Access



# Low-Level to OS-Level Access



Looks a bit like  
dynamic loading /  
linking...



---

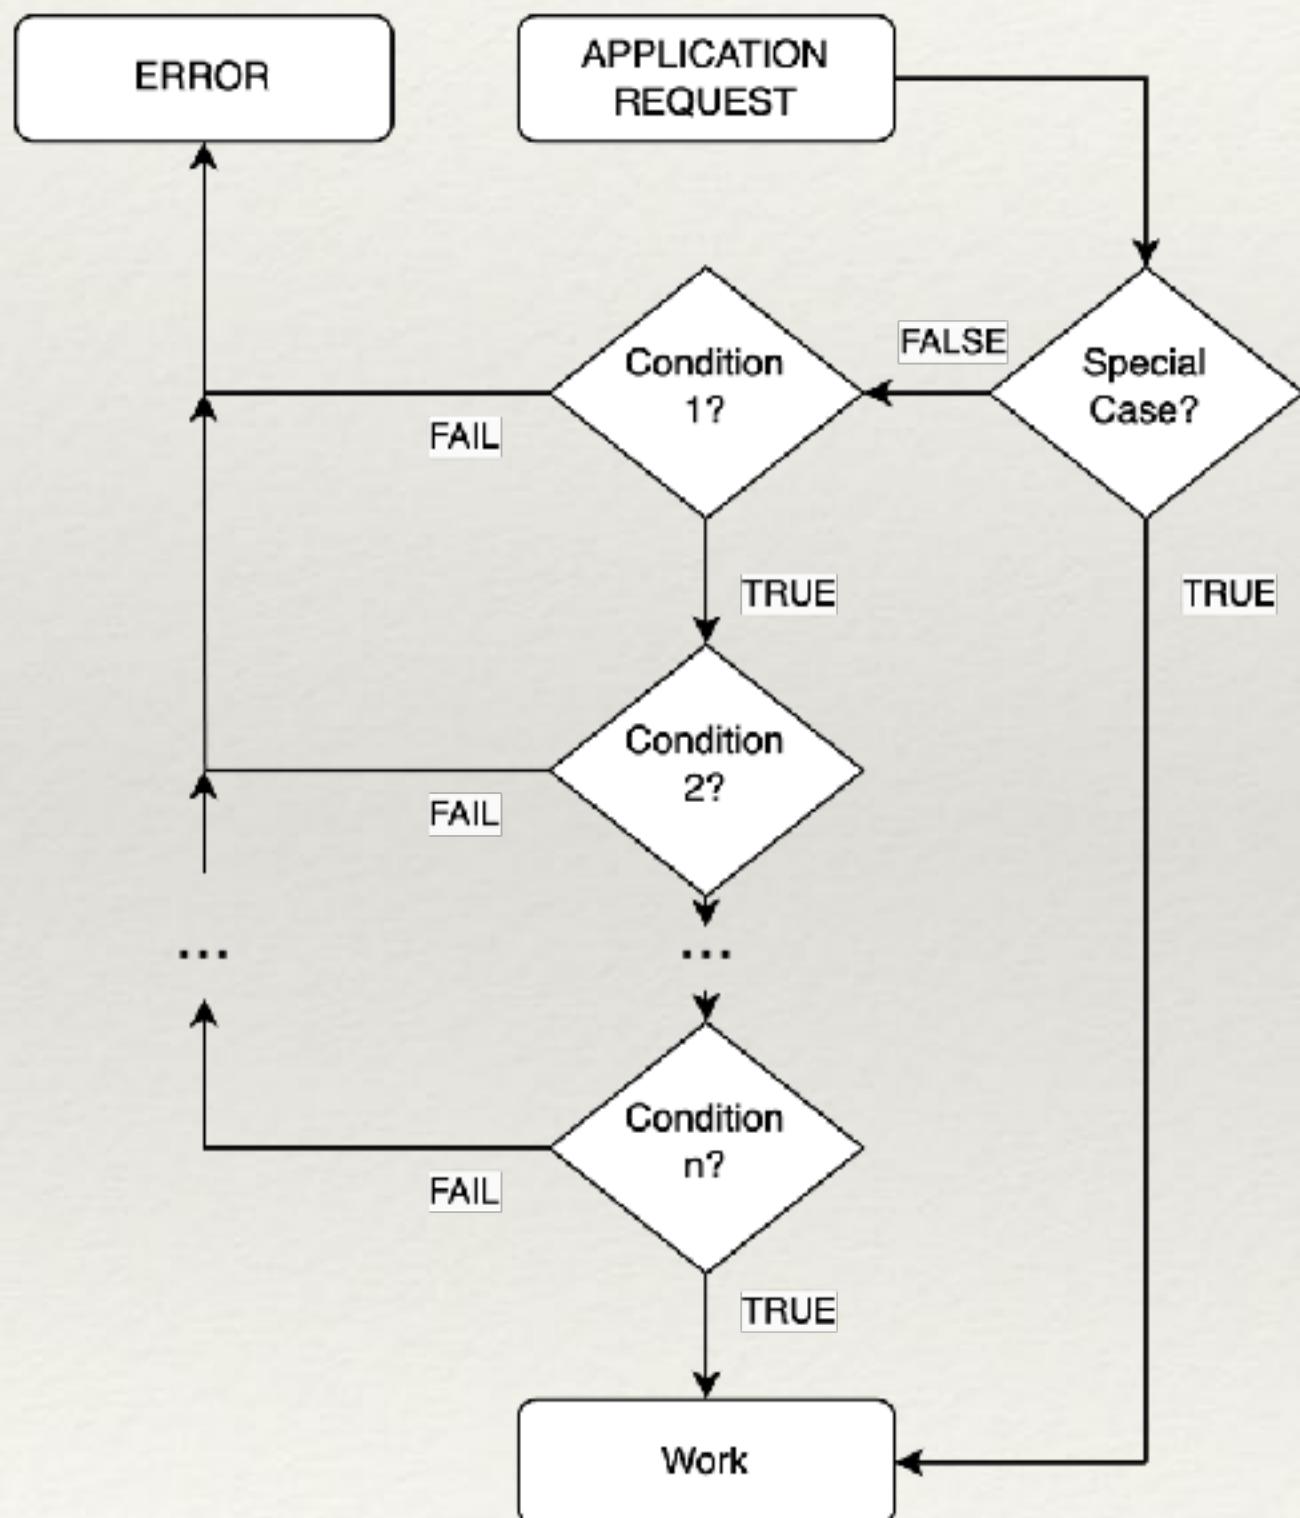
# Ch 4: Case Studies

---

- ❖ \*Application performance optimization
  - ❖ System call shortcircuiting
- ❖ Kernel Code Reuse
  - ❖ Observing page table structure

# Case Study: Shortcutting

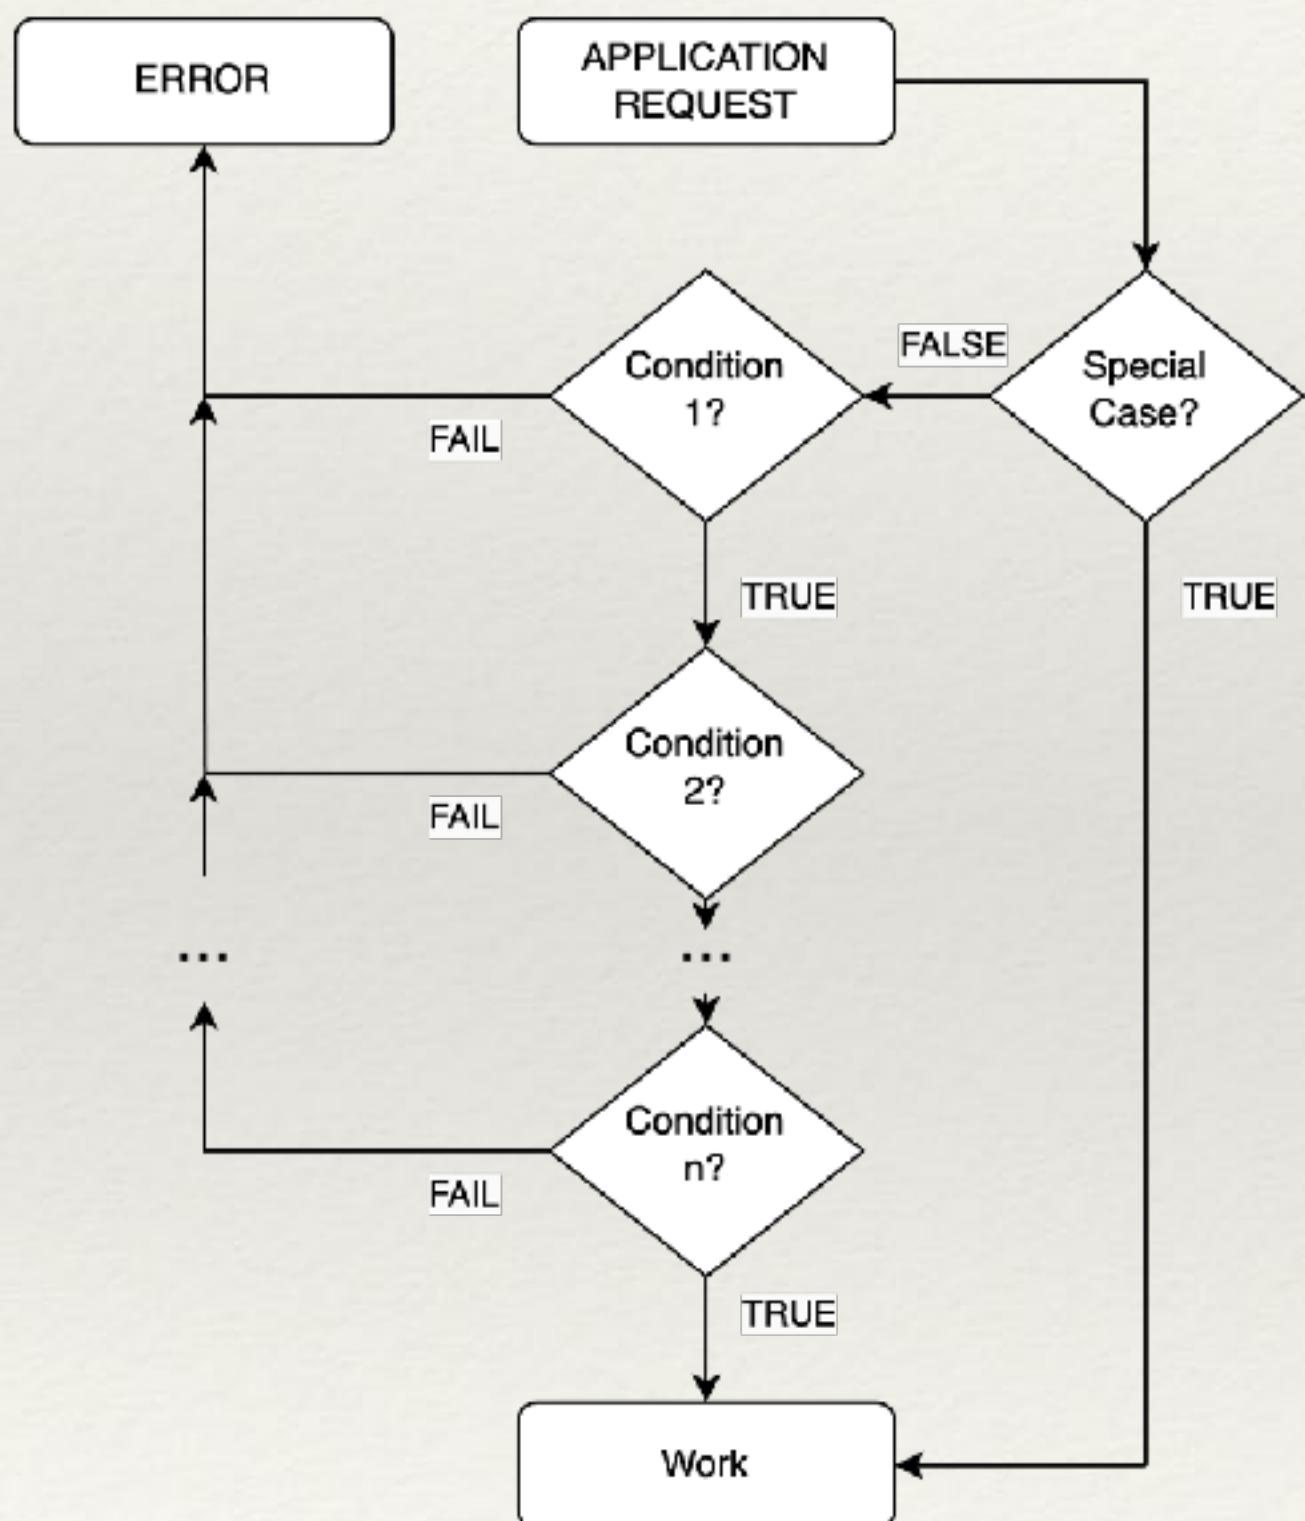
Can we just build this?



Opportunity

# Shortcutting

Can we just build this?



Opportunity

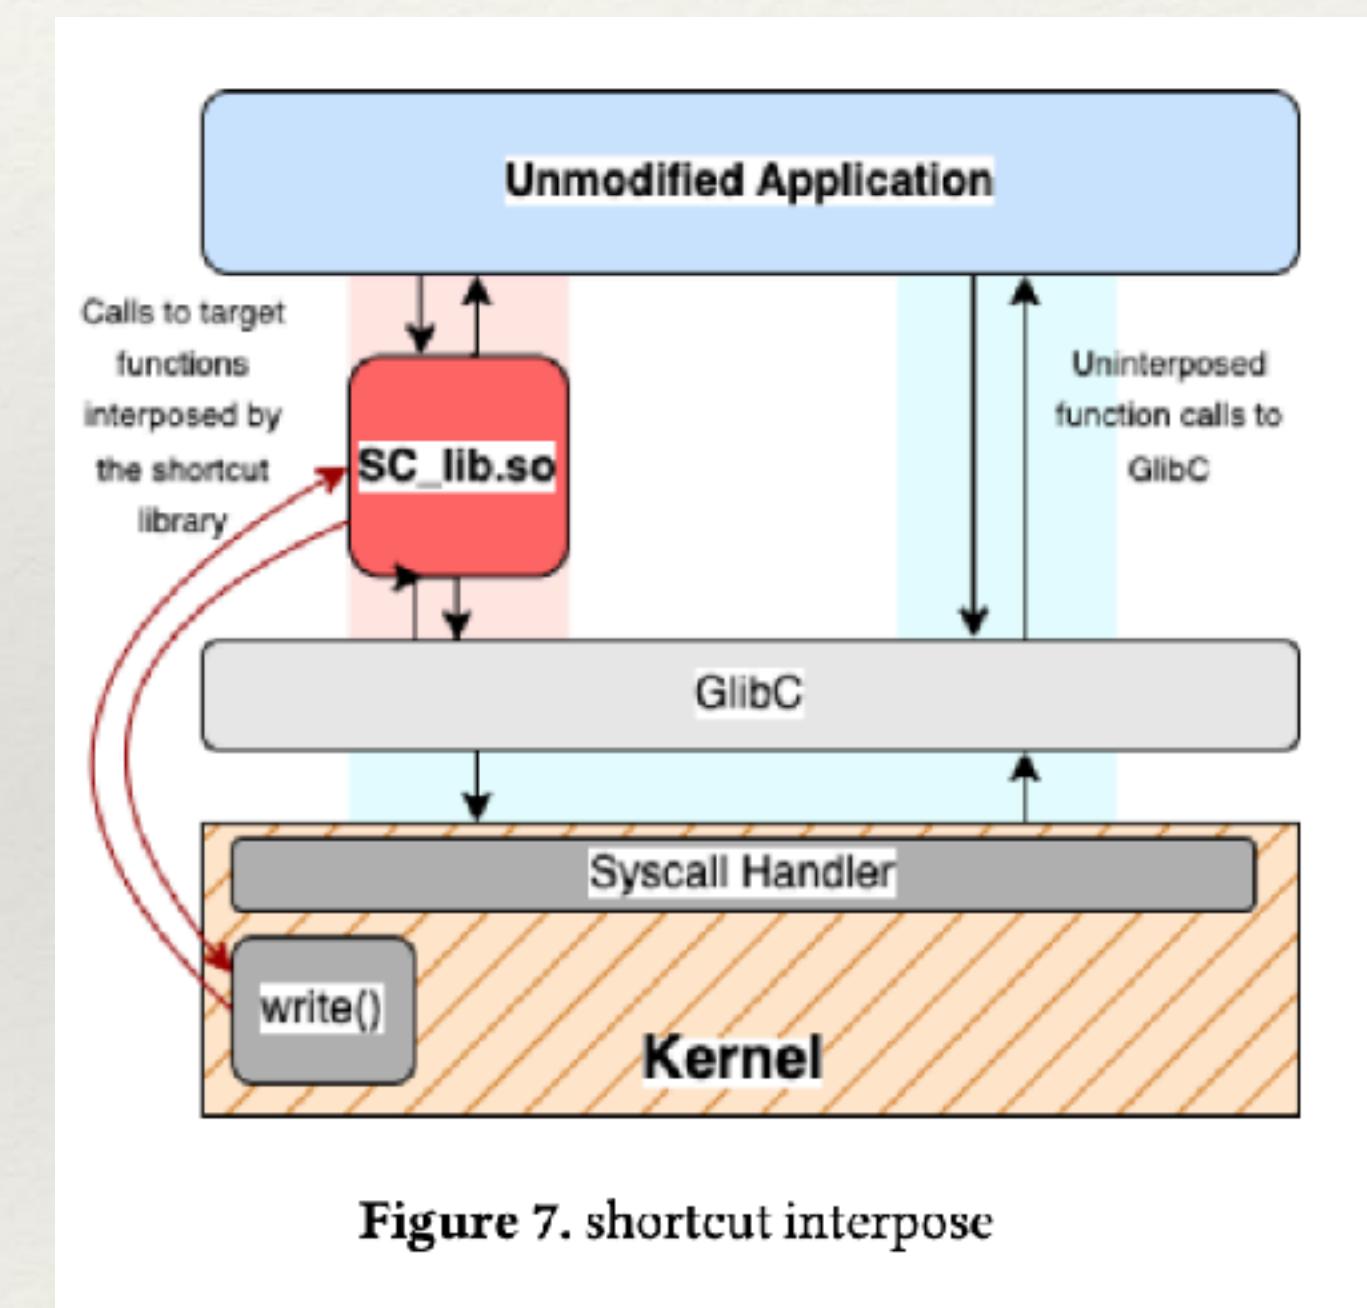
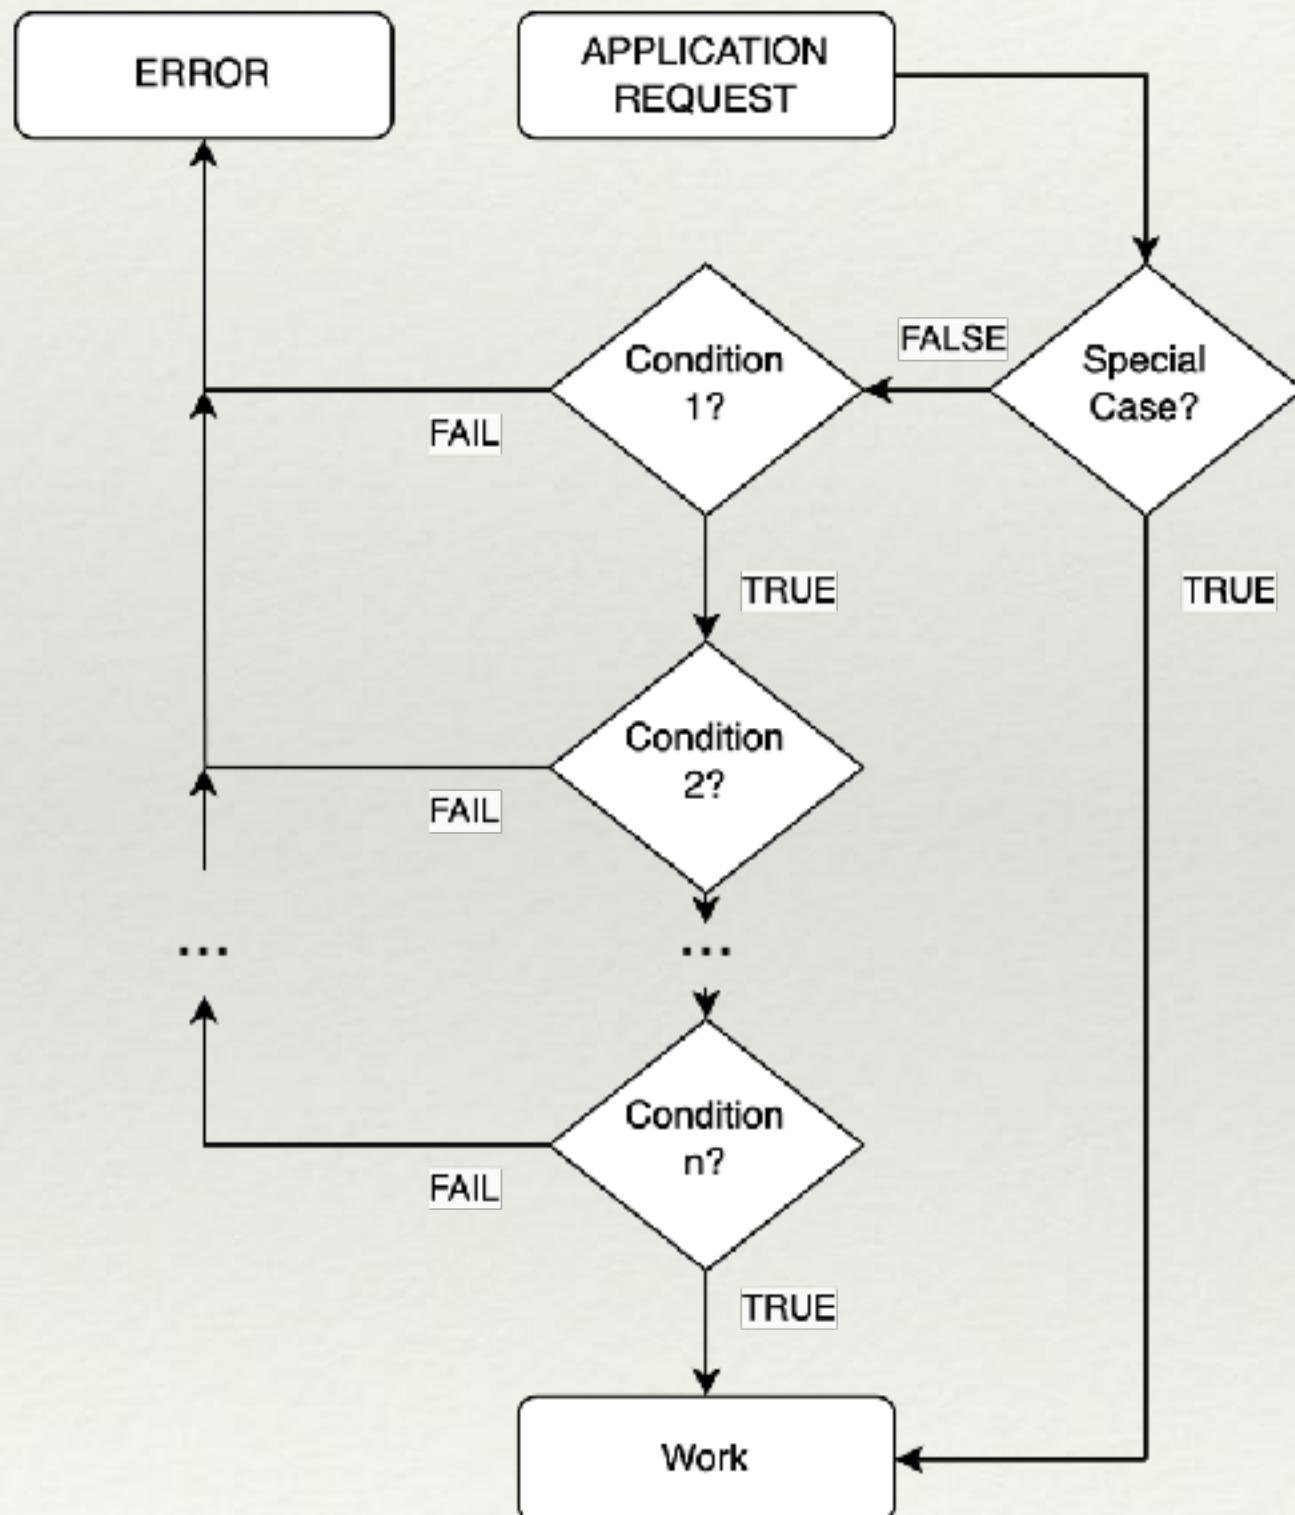


Figure 7. shortcut interpose

Architecture

# Shortcutting

Can we just build this?



```
root: [~/Symbi-OS/Tools/bin/shortcut - [master]]  
$ ./shortcut.sh -h  
#####  
shortcut.sh [args] --- <app> [app args]  
Script for launching apps with options for elevation and shortcutting.  
Flags:  
-p: passthrough mode, run application without interposer library.  
Ignores all other args,  
The following options may interact with kElevate  
-be / -bl: begin with executable elevated / lowered  
These ones can be repeated to specify multiple functions:  
-e <fns>: Elevate around these functions  
-l <fns>: Lower around these functions  
-s <fn1->fn2>: Shortcut between these functions: eg: 'write->ksys_write'  
Don't forget to use quotes around the function names!  
Debugging options  
-v: verbose mode for debugging  
-d: dry run, do everything EXCEPT running the application  
#####
```

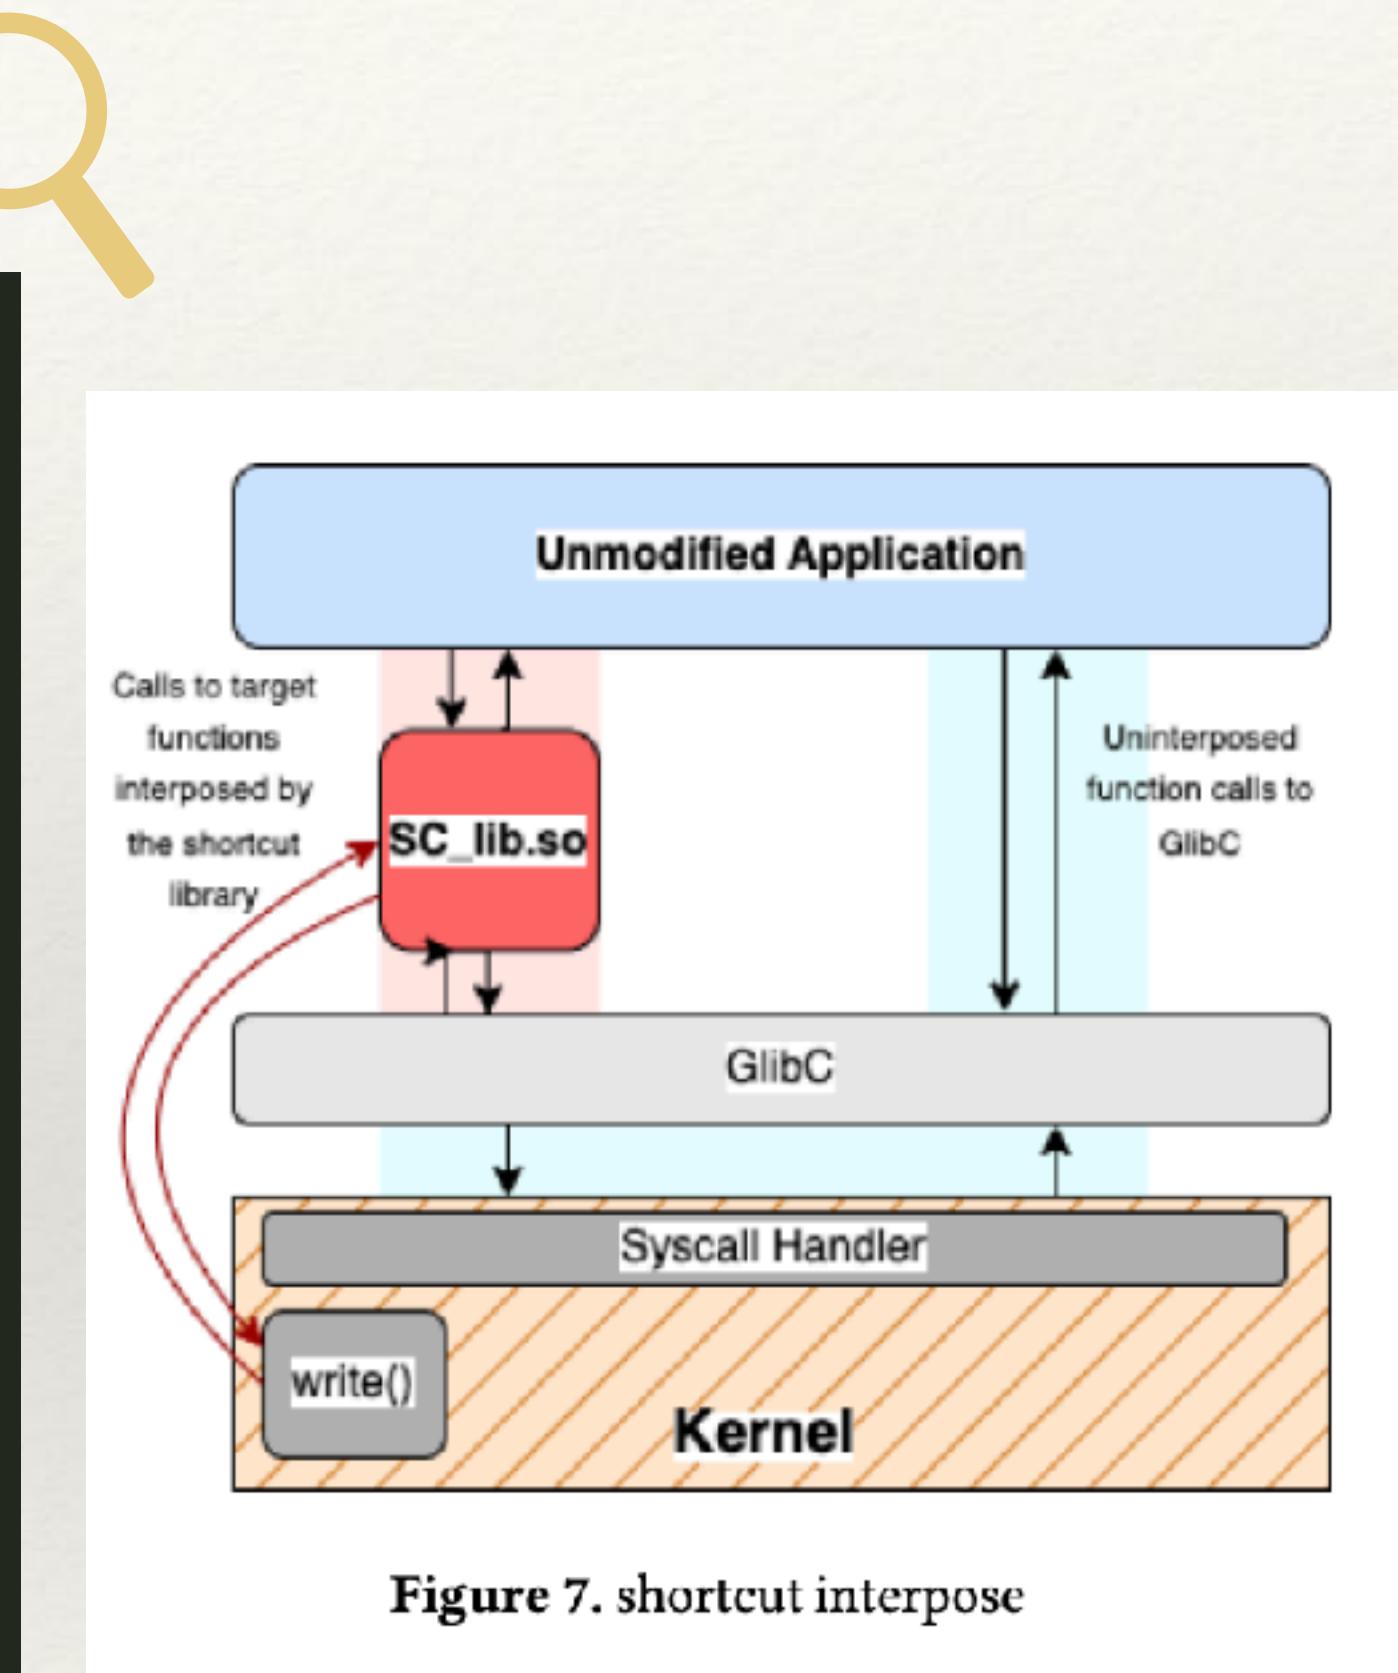


Figure 7. shortcut interpose

Opportunity

Interface

Architecture

# Workload

- A C program, *write\_loop*, makes write function calls
- Built with a standard toolchain
- With writes destined for a Ramdisk FS

```
void write_loop(int fd, char *buf, int size, int iter){  
  
    for (int i = 0; i < iter; i++) {  
        int rc = write(fd, buf, size);  
        if (rc == -1) {  
            printf("write failed\n");  
            exit(1);  
        }  
    }  
}
```

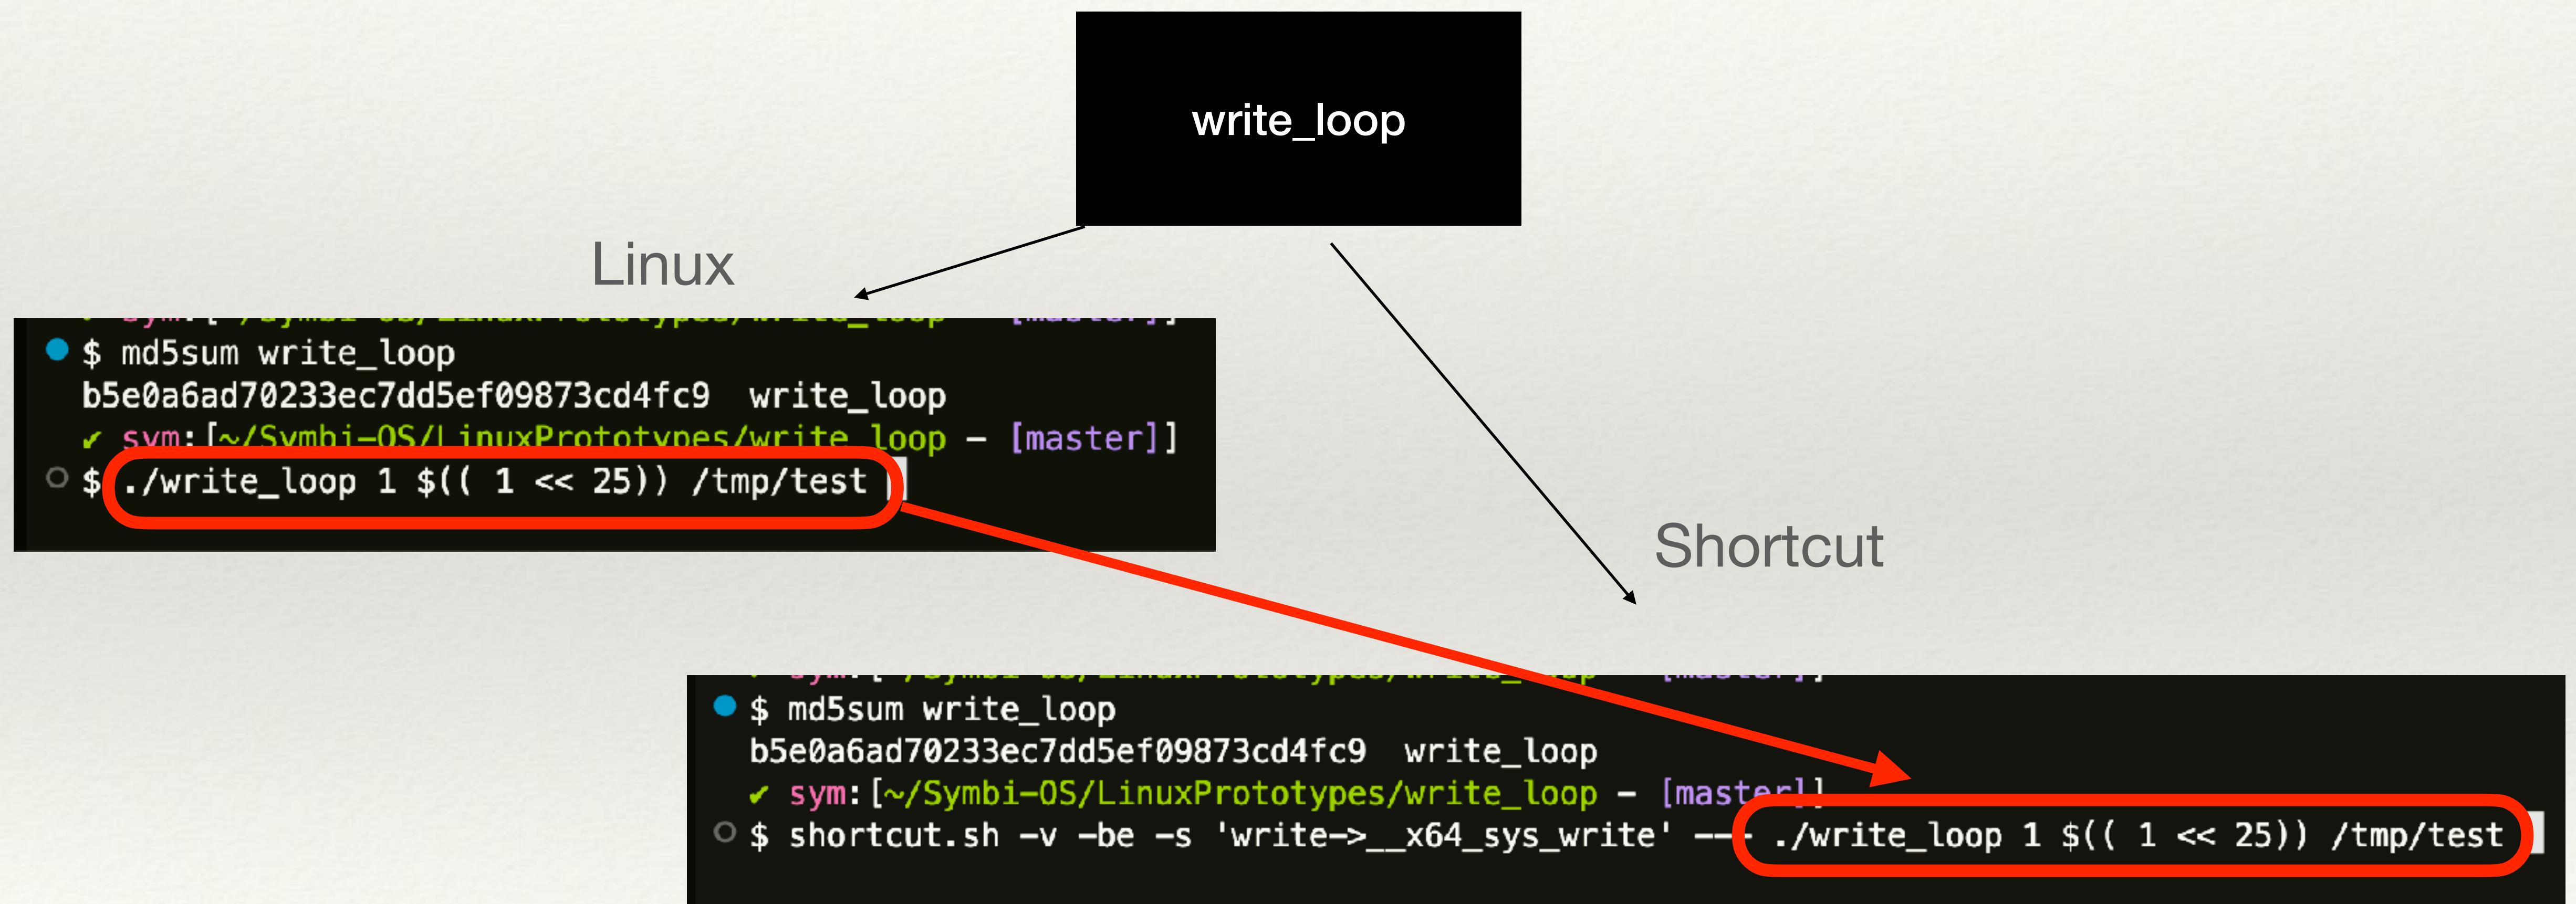
# A Comparison



# A Comparison



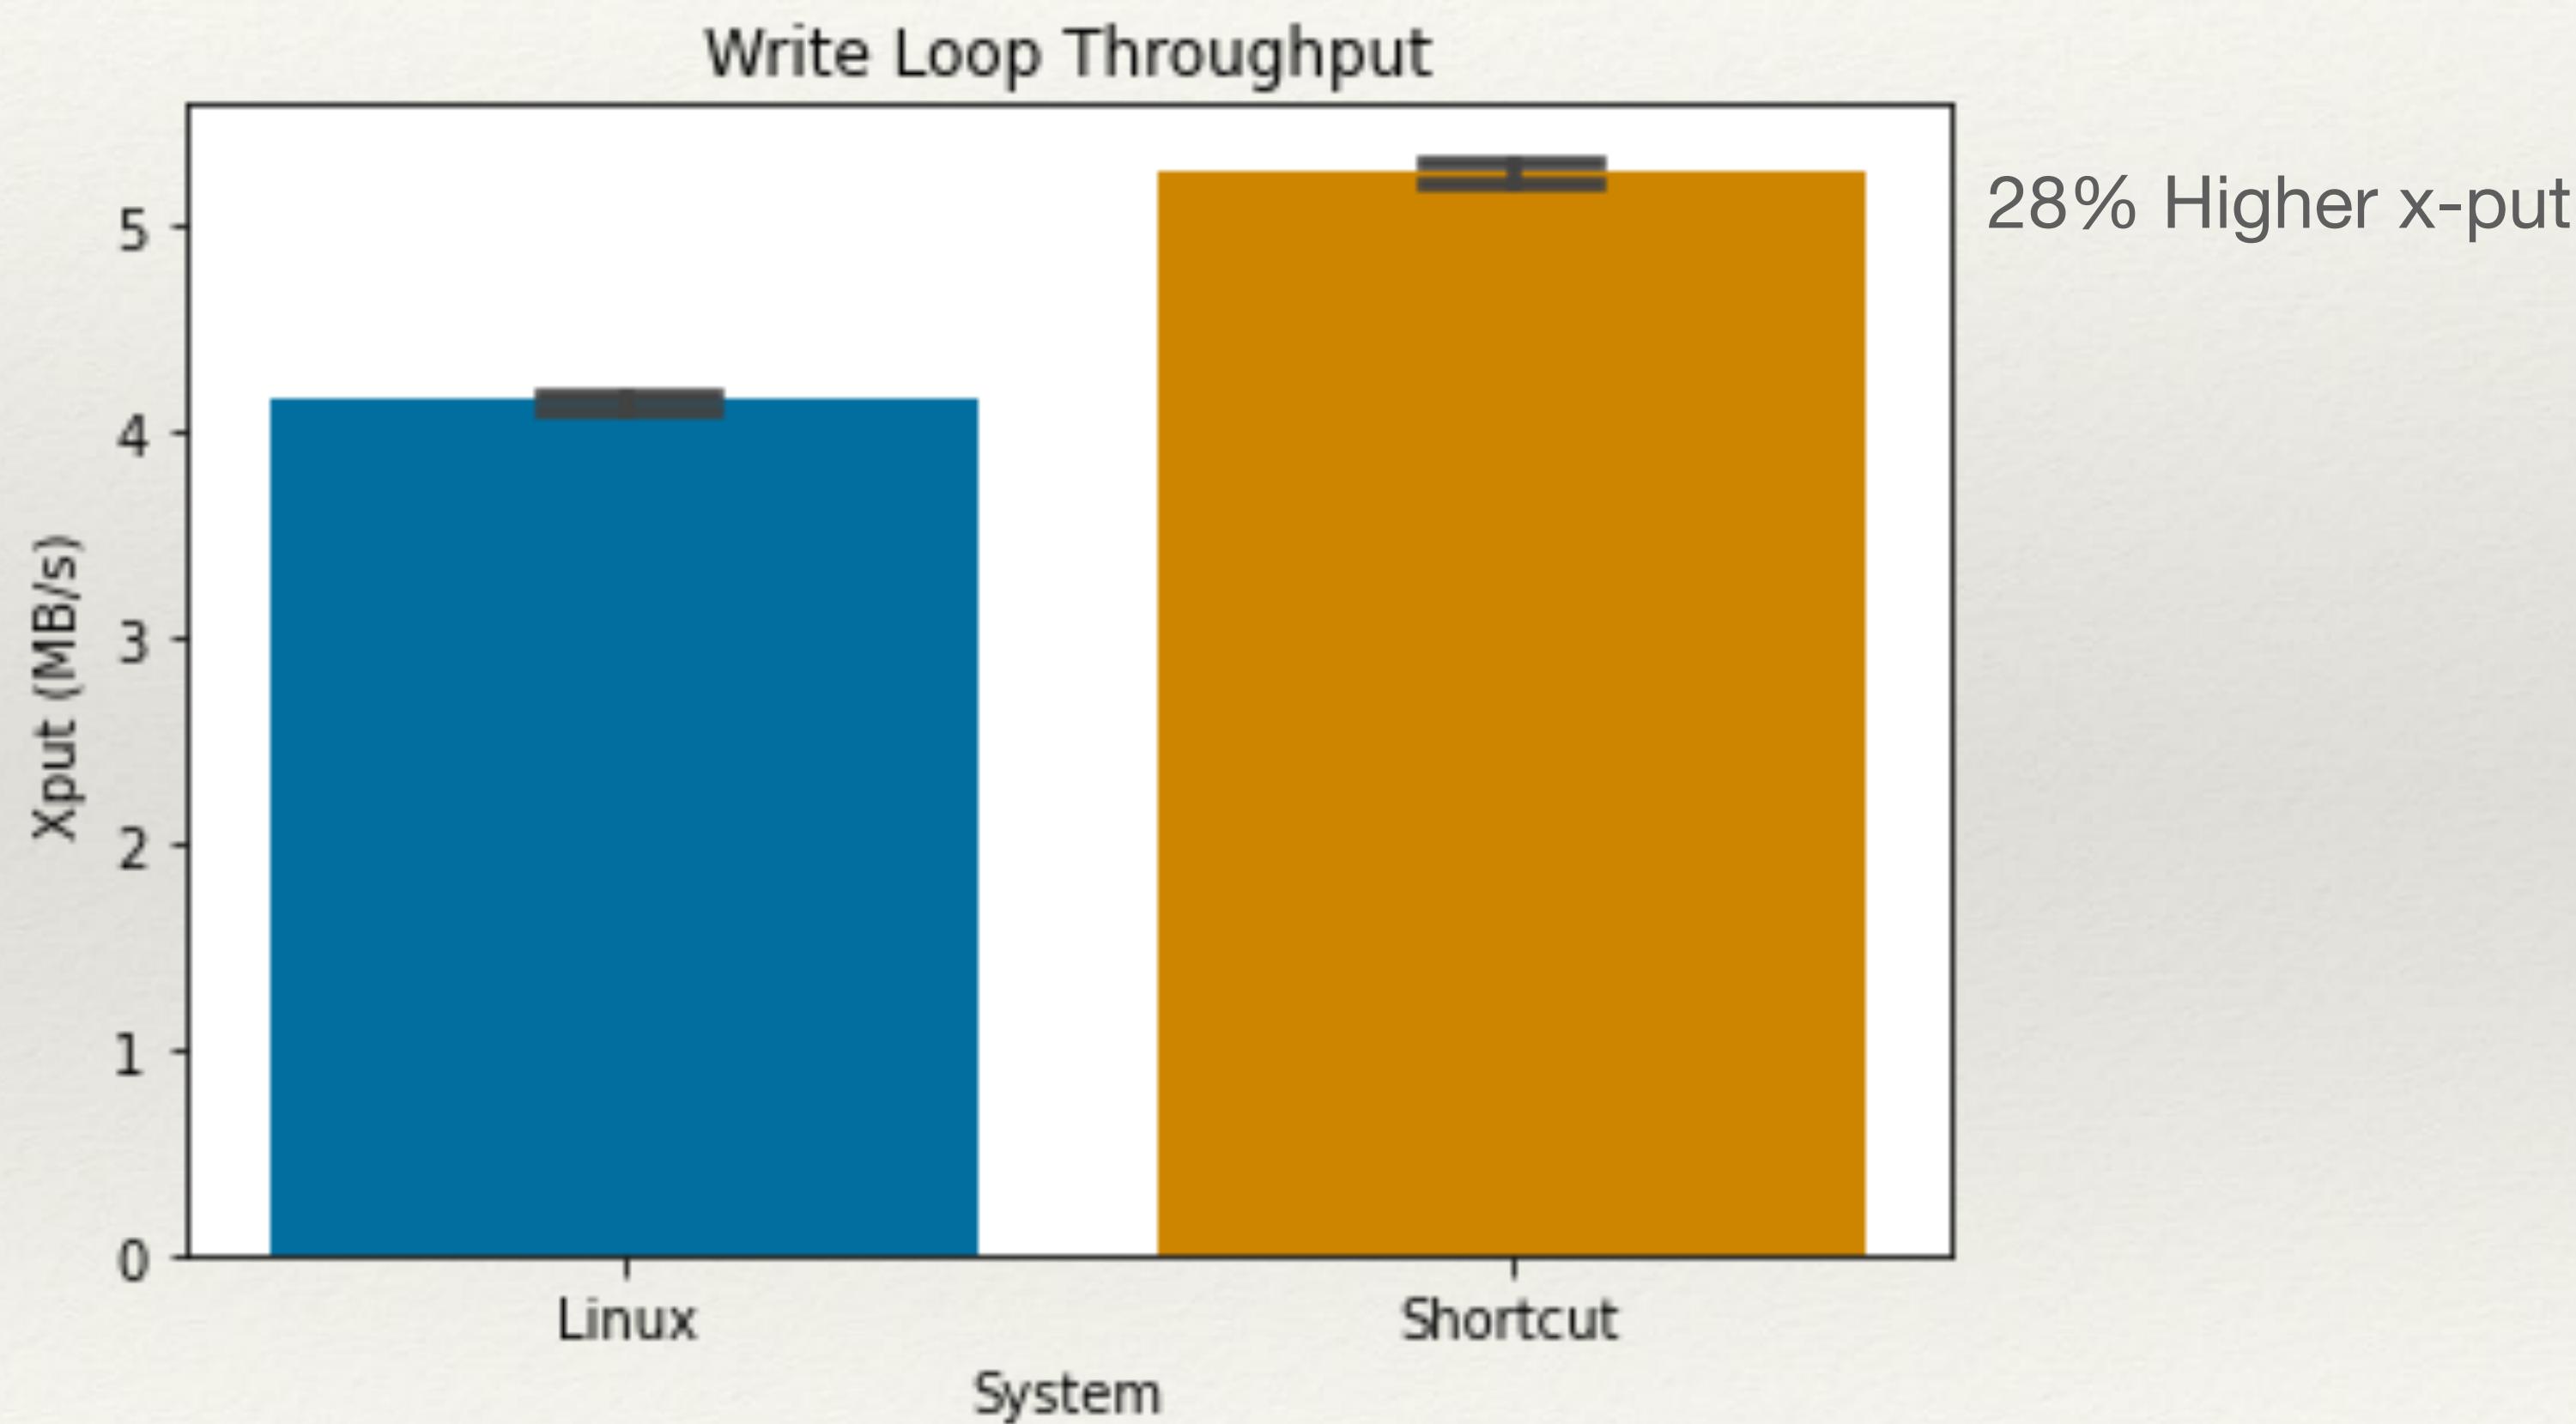
# A Comparison



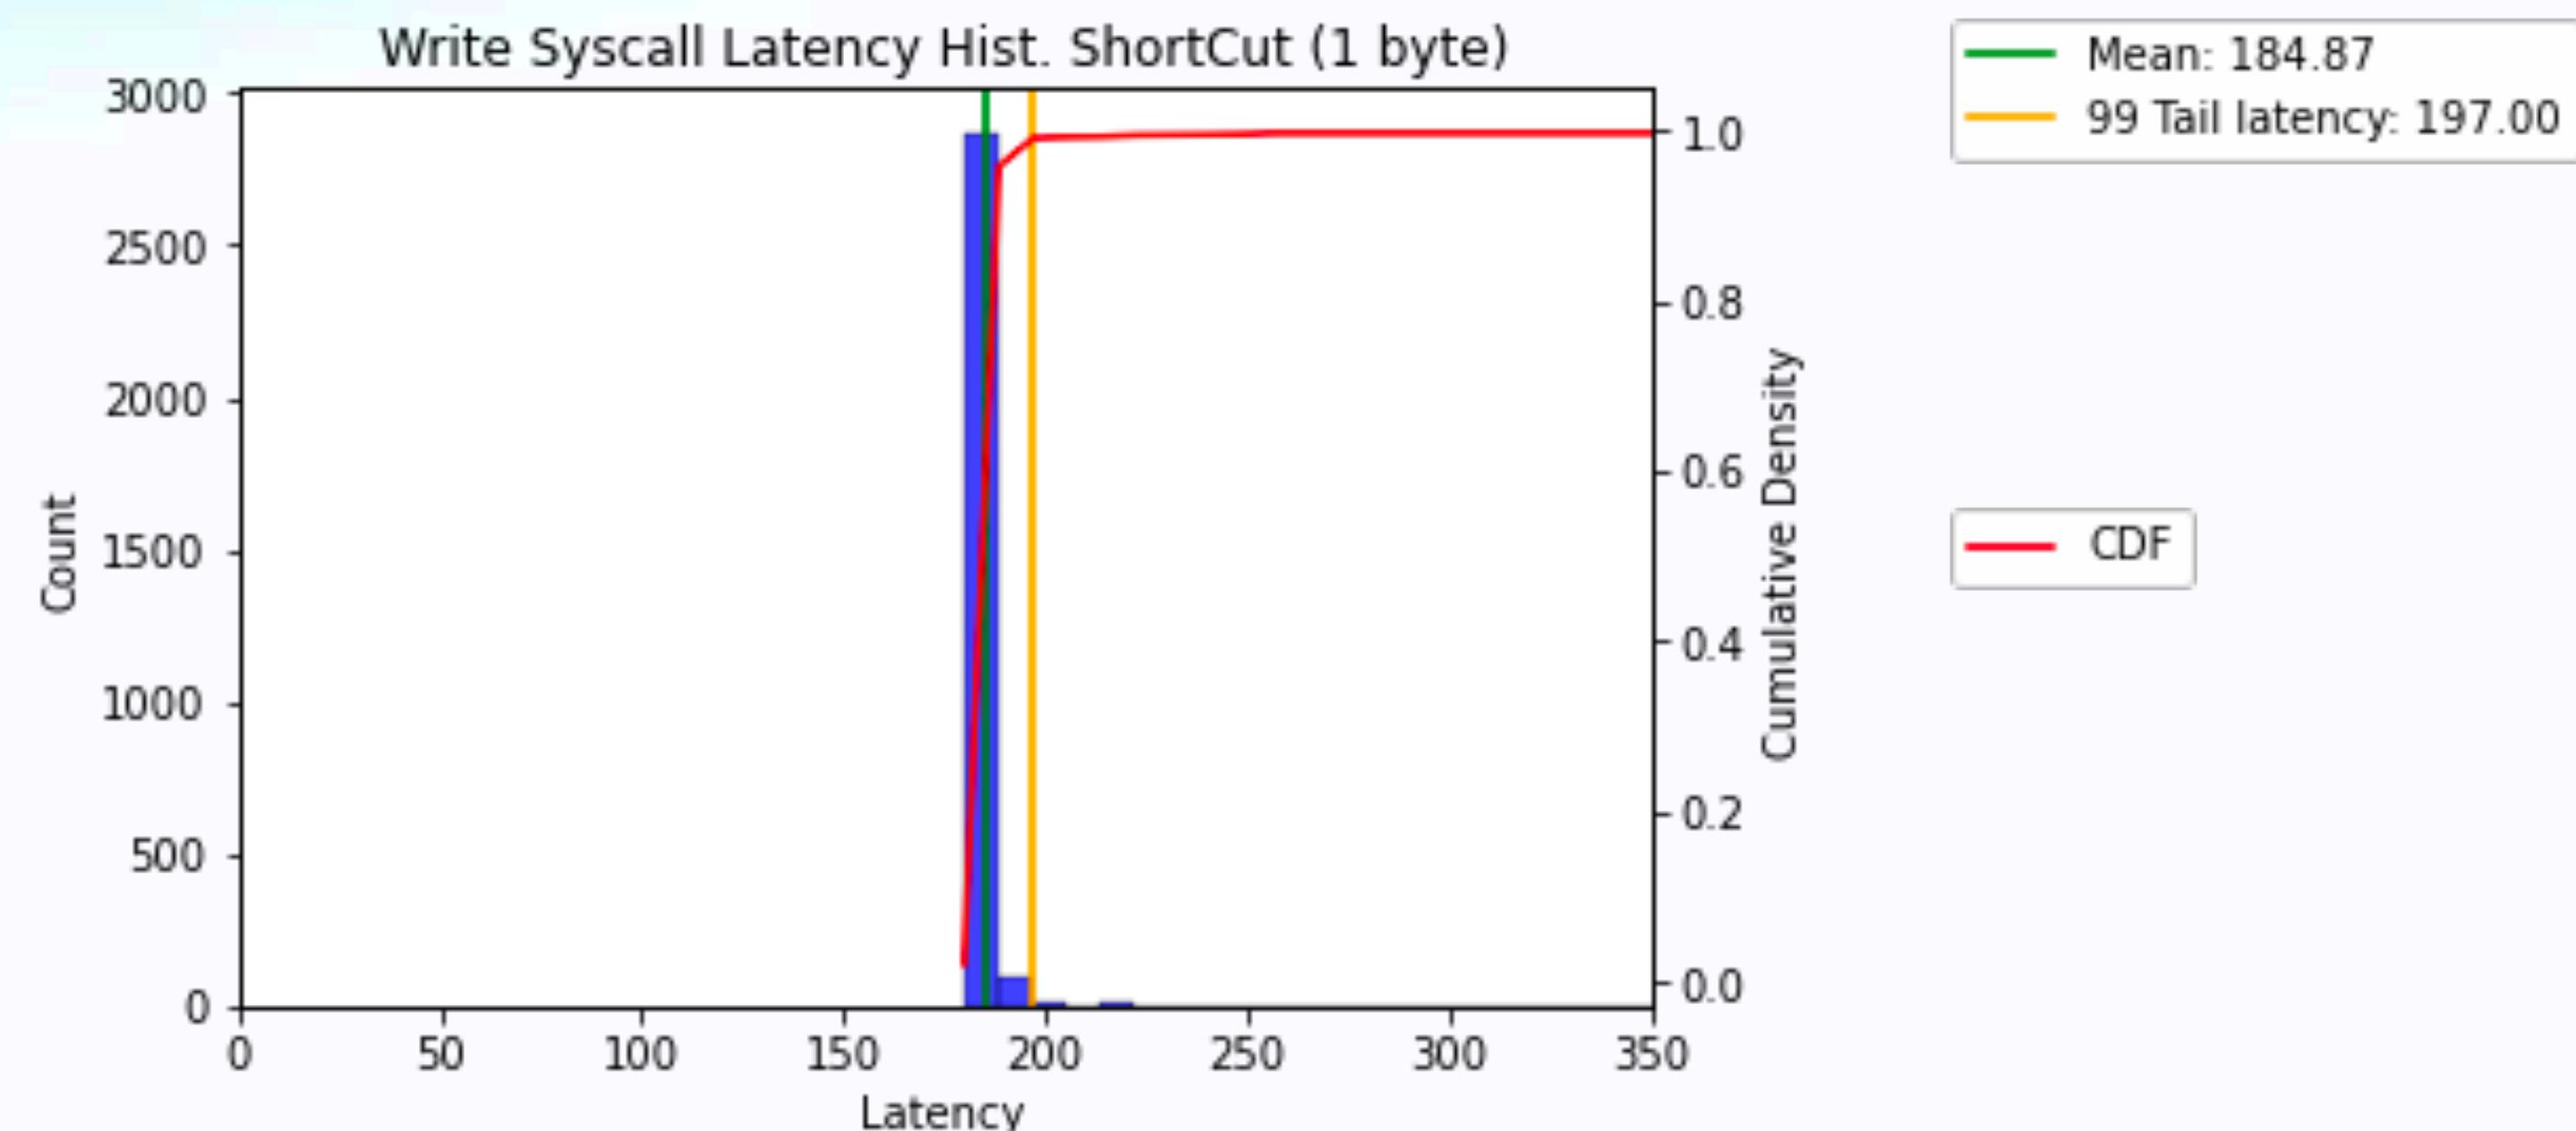
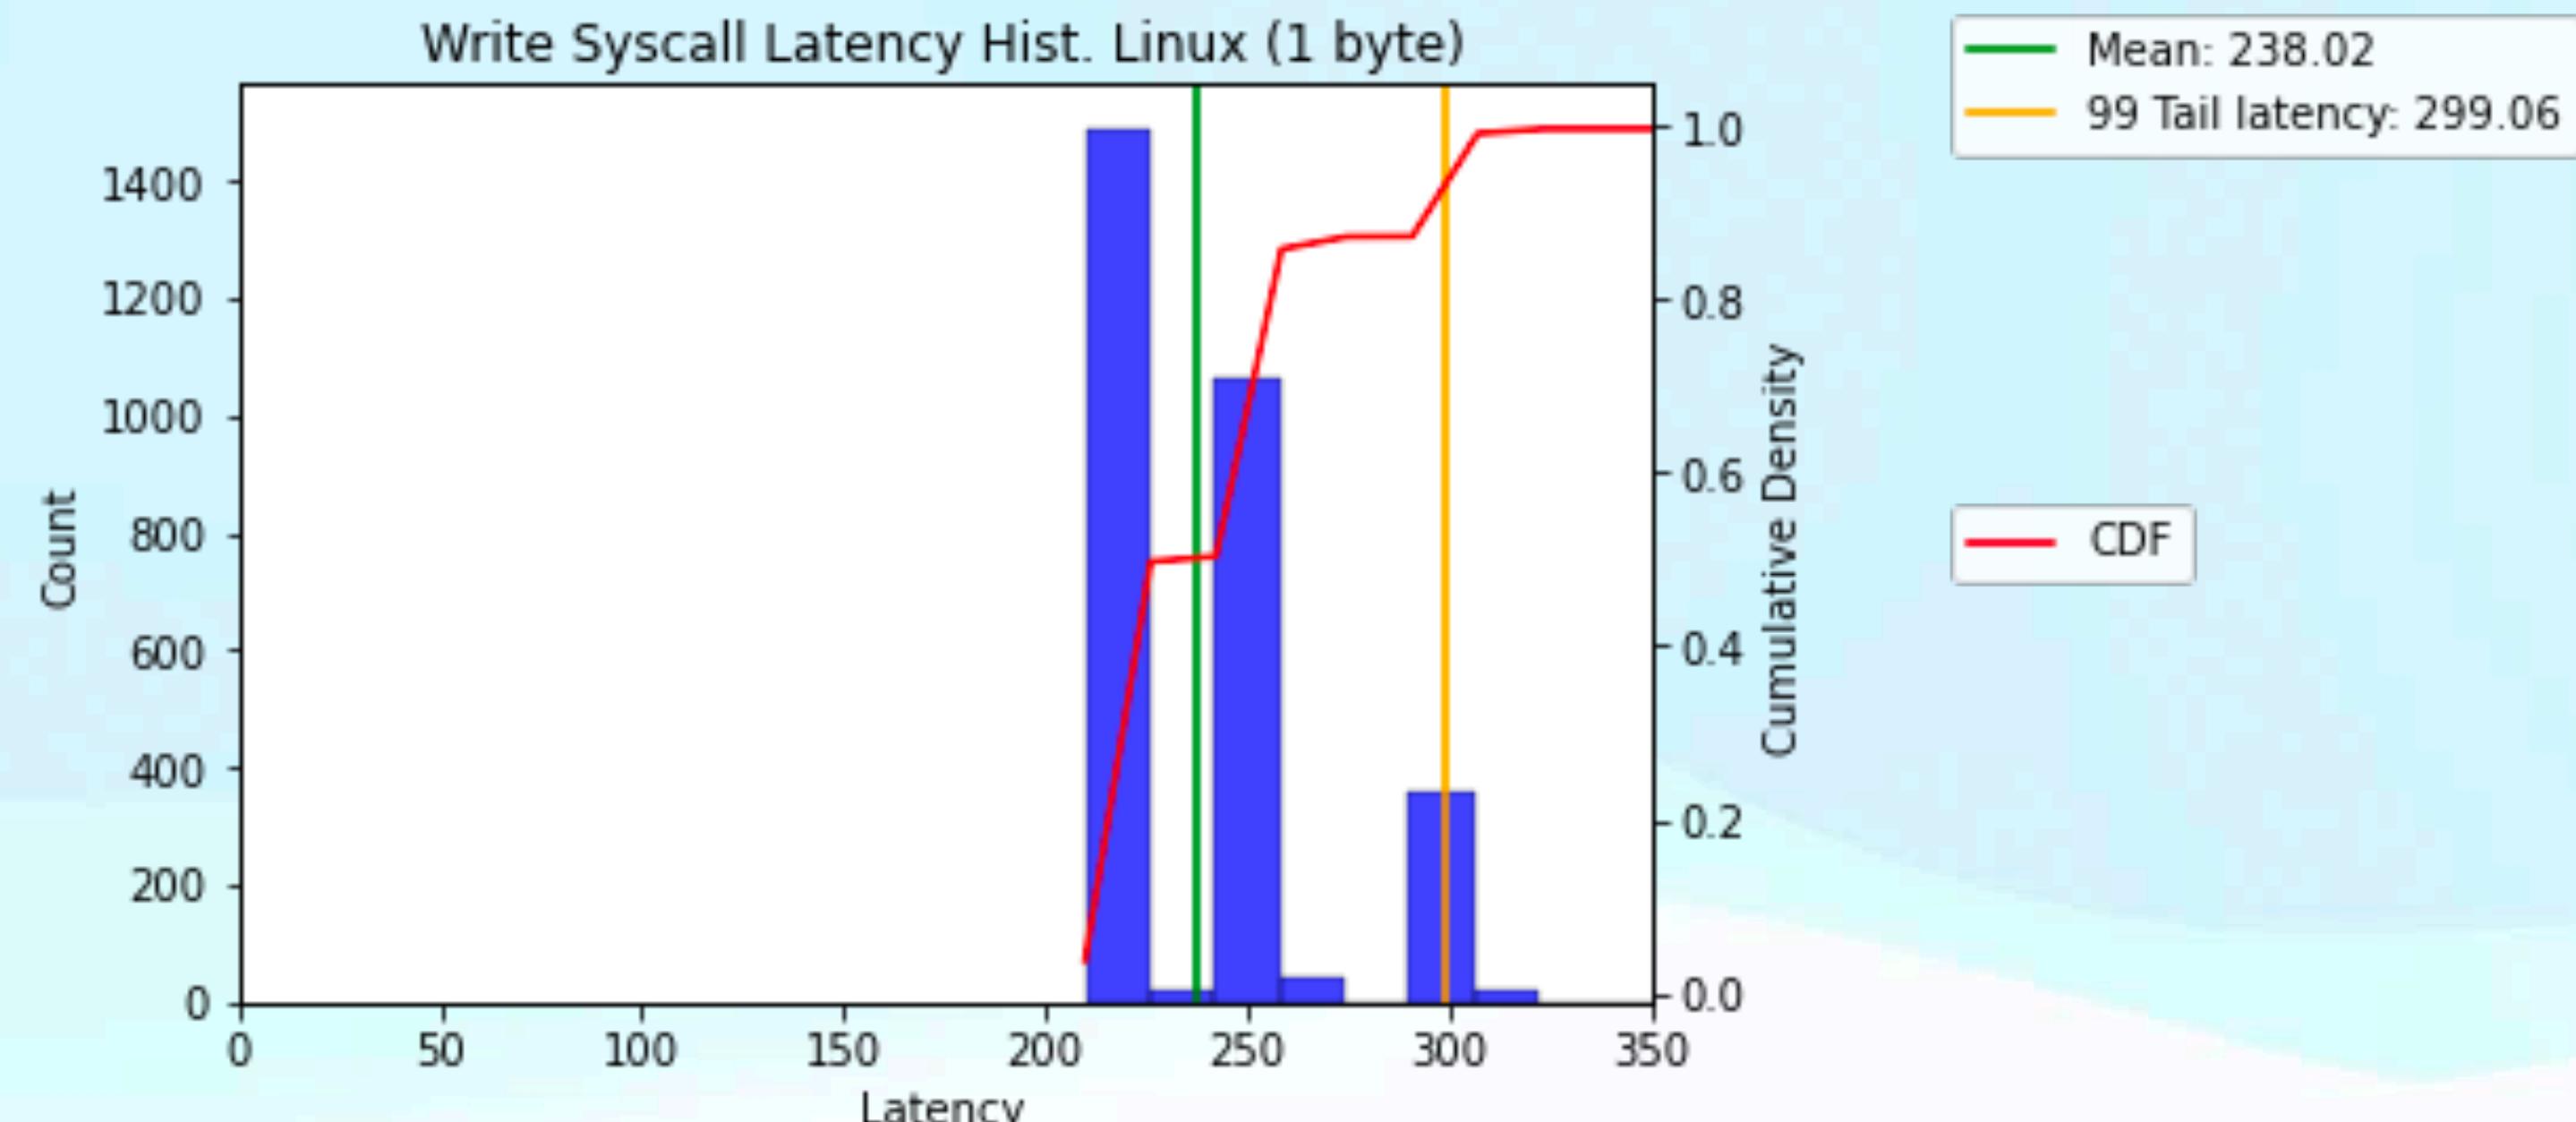
# A Comparison



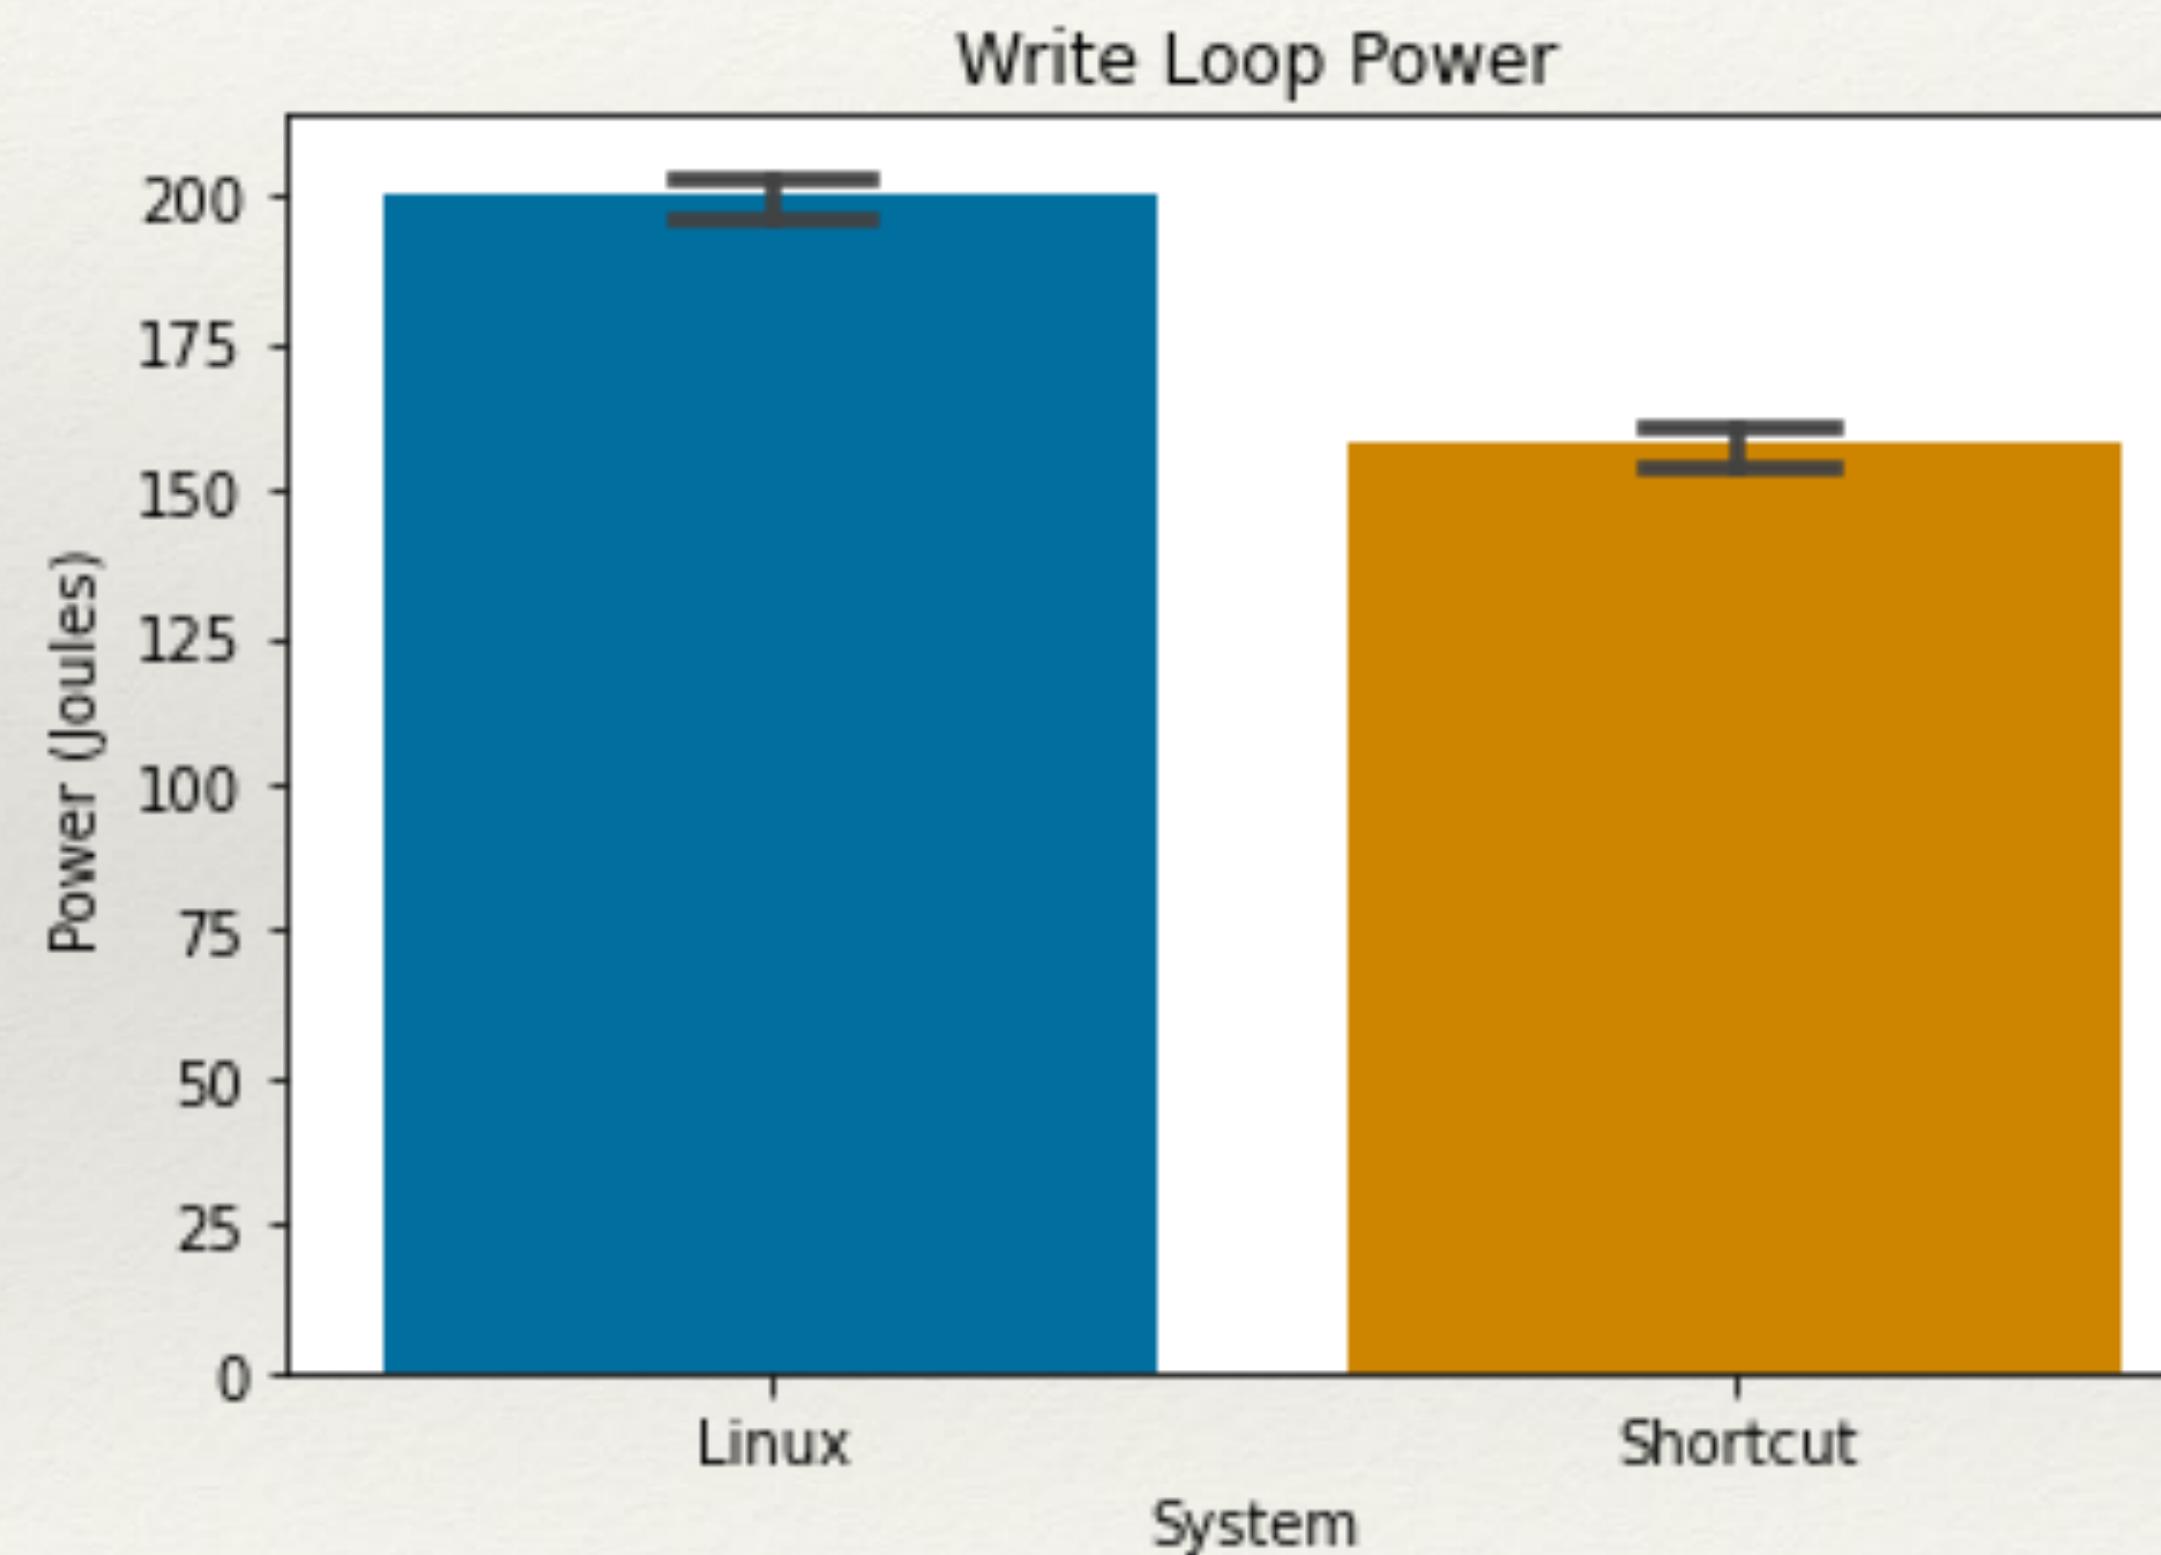
# Throughput Comparison



# Syscall Latencies



# Energy Comparison



24% Reduction in Energy

---

# Shortcutting: Macrobenchmarks

---

- ❖ Refer to dissertation
  - ❖ Redis:
    - ❖ Shallow shortcuts: 9-11%  
xput improvement
    - ❖ Deep shortcuts: 20-22%  
xput improvement
  - ❖ Memcached
    - ❖ Shallow shortcuts: 9-11%  
xput improvement

# Shortcutting: Macrobenchmarks

- ❖ Refer to dissertation

- ❖ Redis:

- ❖ Shallow shortcuts: 9-11%  
xput improvement

- ❖ Deep shortcuts: 20-22%  
xput improvement

- ❖ Memcached

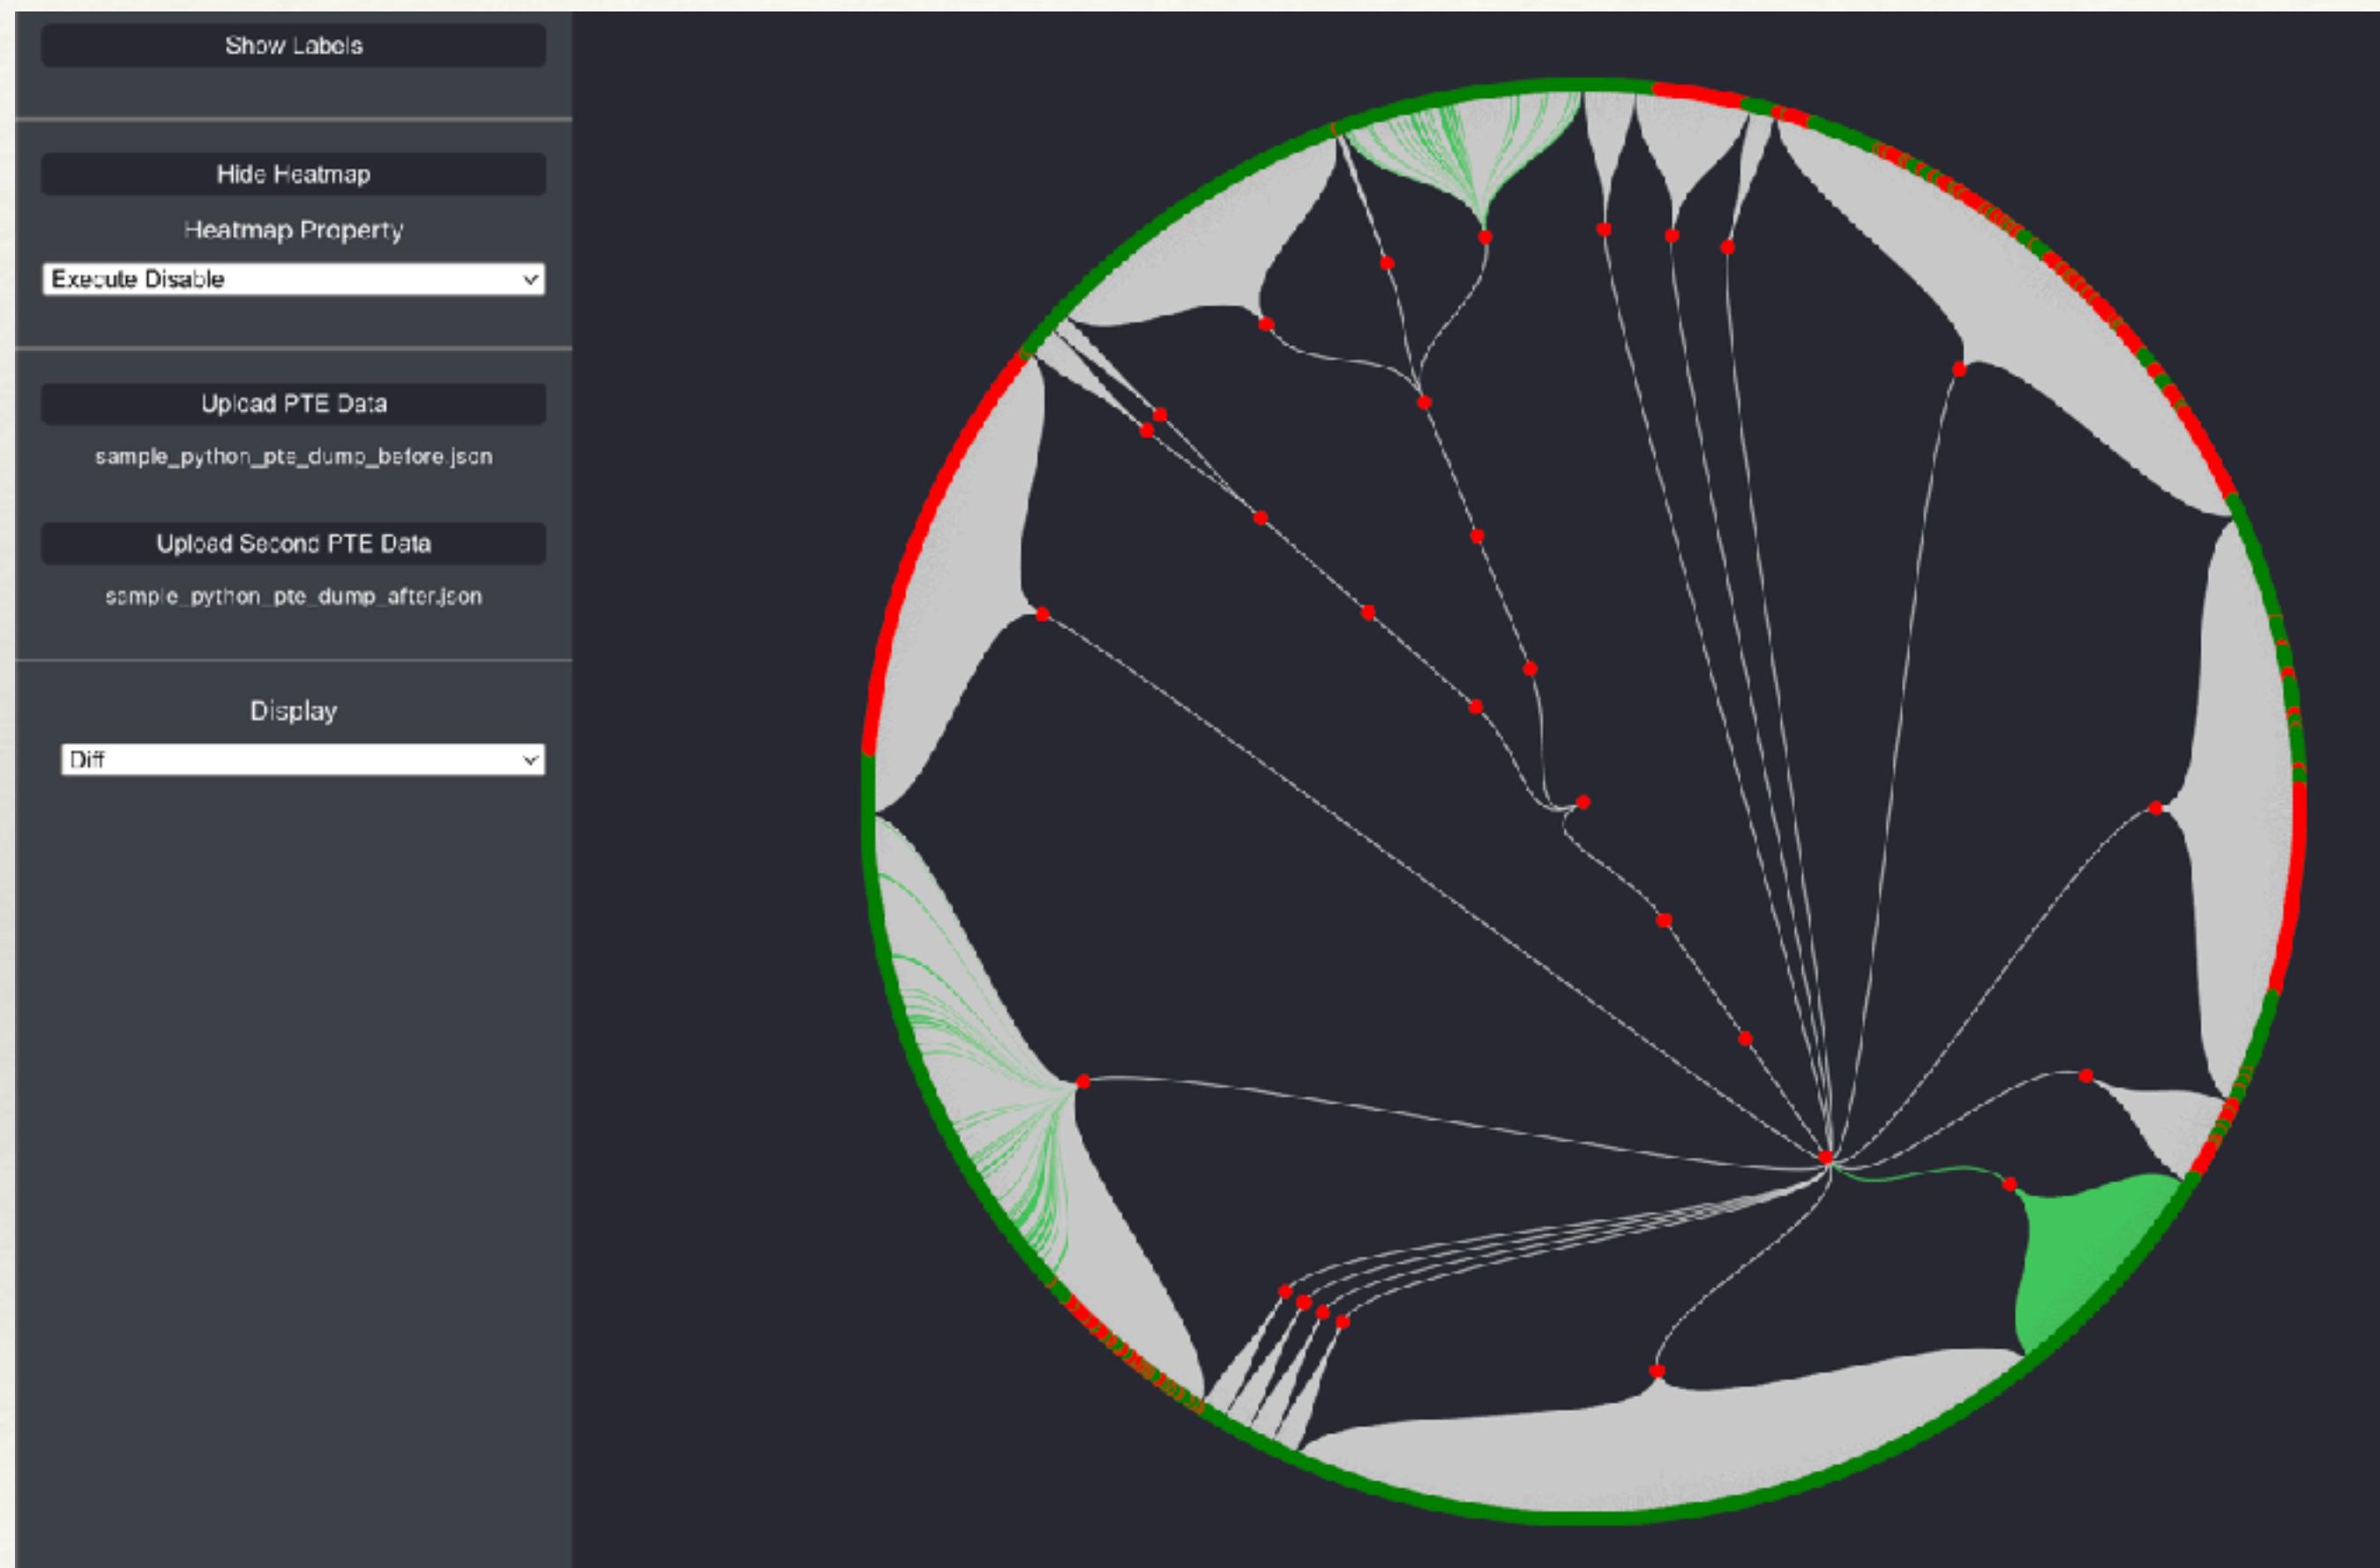
- ❖ Shallow shortcuts: 9-11%  
xput improvement

```
$ ./sc.sh -s write->__x64_sys_write \
--- ./wrLoop $SZ $ITER $PATH
```

```
$ ./sc.sh -s \
write->tcp_sendmsg:0x42ed60 \
--- ./wrLoop $SZ $ITER $PATH
```



# Case Study: Kernel Code Reuse



Page Map Level 5 Entry

63	62	...	52	51	...	M	M-1	...	12	11	...	8	7	6	5	4	3	2	1	0
X D	AVL		Reserved (0)		Bits 12 - (M-1) of address		AVL		R S V D A C W U R P	A V A D A D T S W	P P D T P U R S W									

Page Map Level 4 Entry

63	62	...	52	51	...	M	M-1	...	12	11	...	8	7	6	5	4	3	2	1	0
X D	AVL		Reserved (0)		Bits 12 - (M-1) of address		AVL		R S V D A C W U R P	A V A D A D T S W	P P D T P U R S W									

Page Directory Pointer Table Entry

63	62	...	52	51	...	M	M-1	...	12	11	...	8	7	6	5	4	3	2	1	0
X D	AVL		Reserved (0)		Bits 12 - (M-1) of address		AVL		P A S V A C W U R P	A V A D A D T S W	P P D T P U R S W									

Page Directory Entry

63	62	...	52	51	...	M	M-1	...	12	11	...	8	7	6	5	4	3	2	1	0
X D	AVL		Reserved (0)		Bits 12 - (M-1) of address		AVL		P A S V A C W U R P	A V A D A D T S W	P P D T P U R S W									

P: Present	PS: Page Size
R/W: Read/Write	AVL: Available
U/S: User/Supervisor	M: Maximum
PWT: Write-Through	Physical Address Bit
PCD: Cache Disable	XD: Execute Disable
A: Accessed	

---

# Concluding

---

---

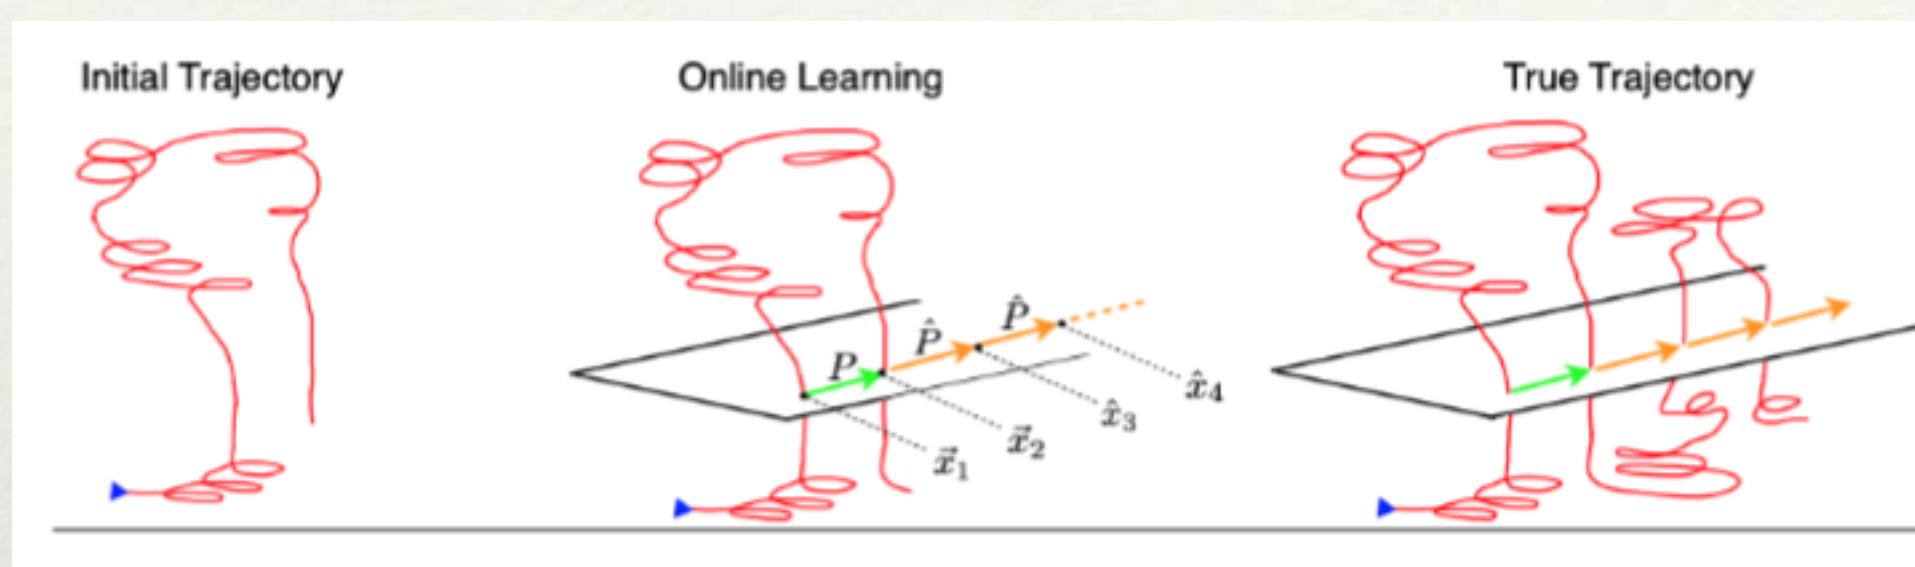
# Contributions Recap

---

- ❖ Defined the OS model Dynamic Privilege; introduced the notion of fine grained HW privilege switching
- ❖ Prototyped a mechanism for extending a GPOS kernel with support for Dynamic Privilege by adding a system call, kElevate
- ❖ Produced a library, kele, which codifies our approach to using kelevate and facilitates future use of the mechanism
- ❖ Explored the use of Dynamic Privilege through case studies

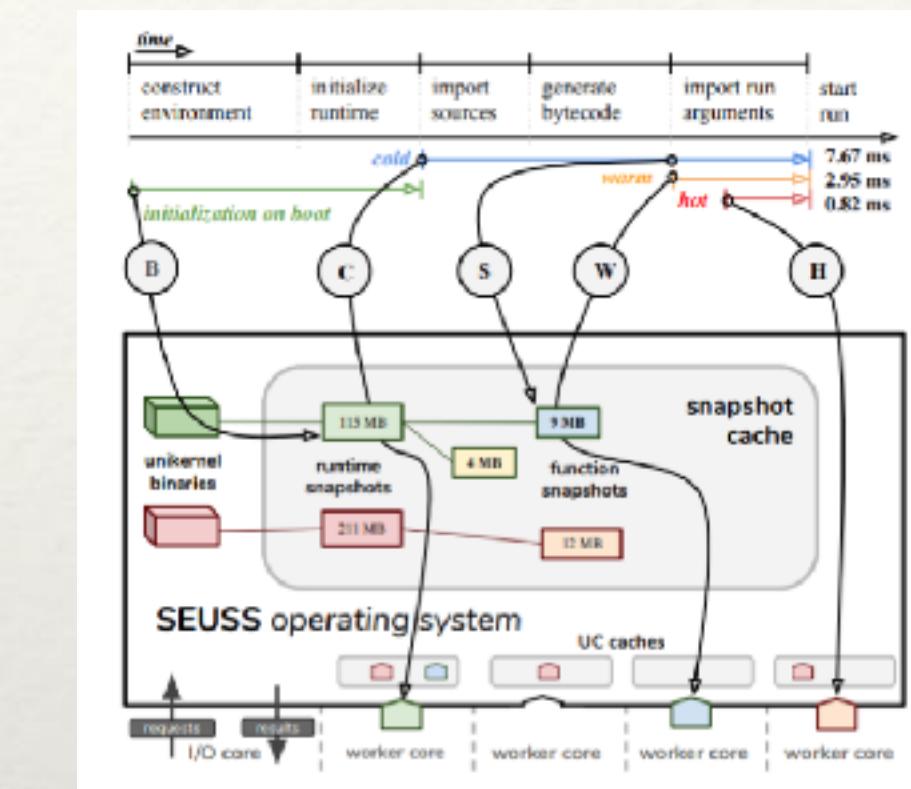
# Prior Work

ASC

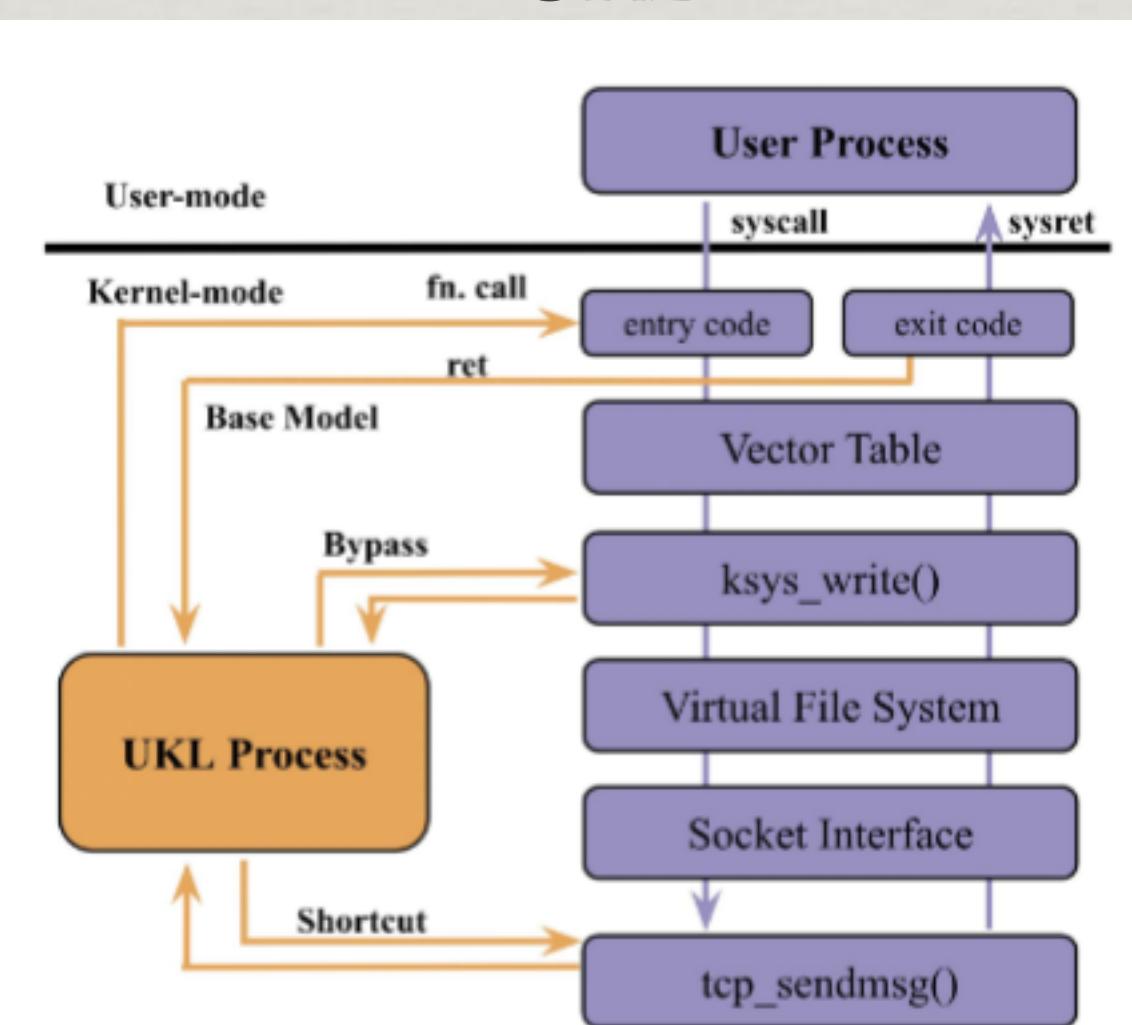


Optimizing  
Systems

SEUSS



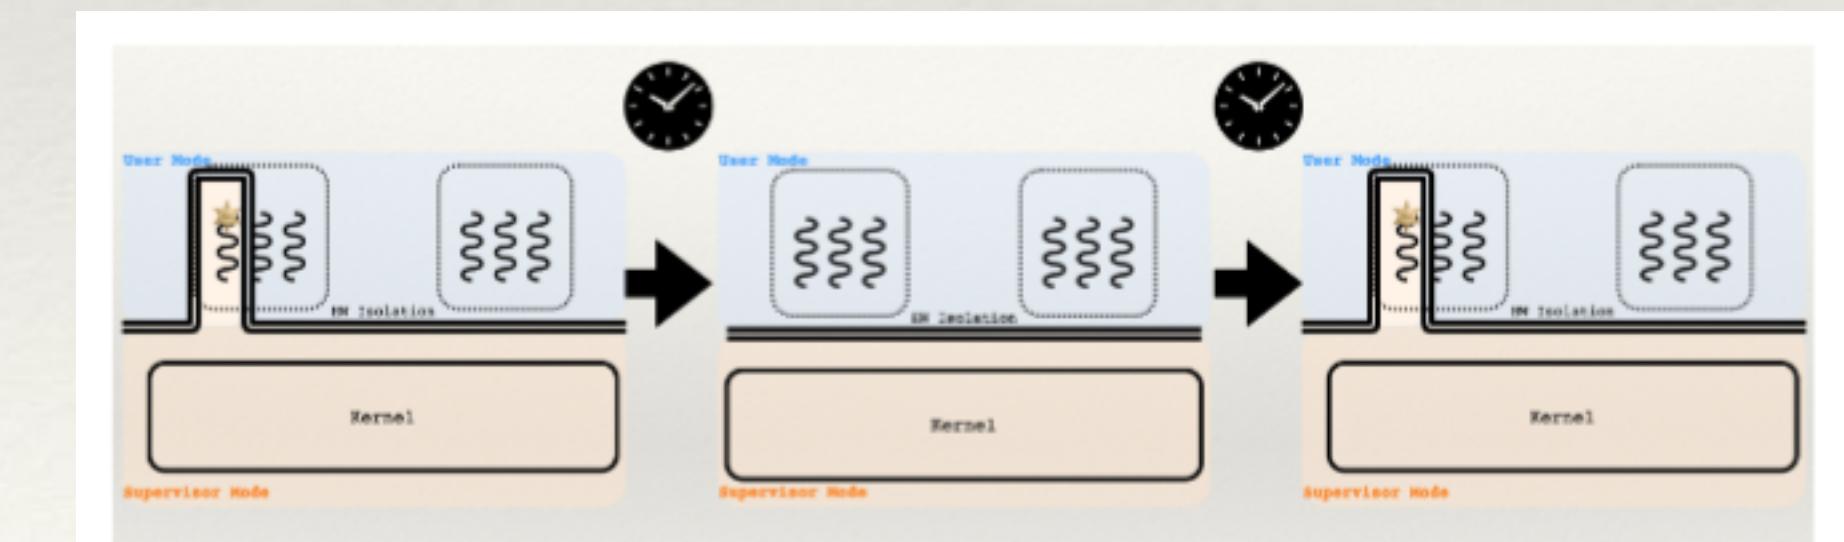
UKL



Realization  
of a known  
alternative  
OS  
Architecture

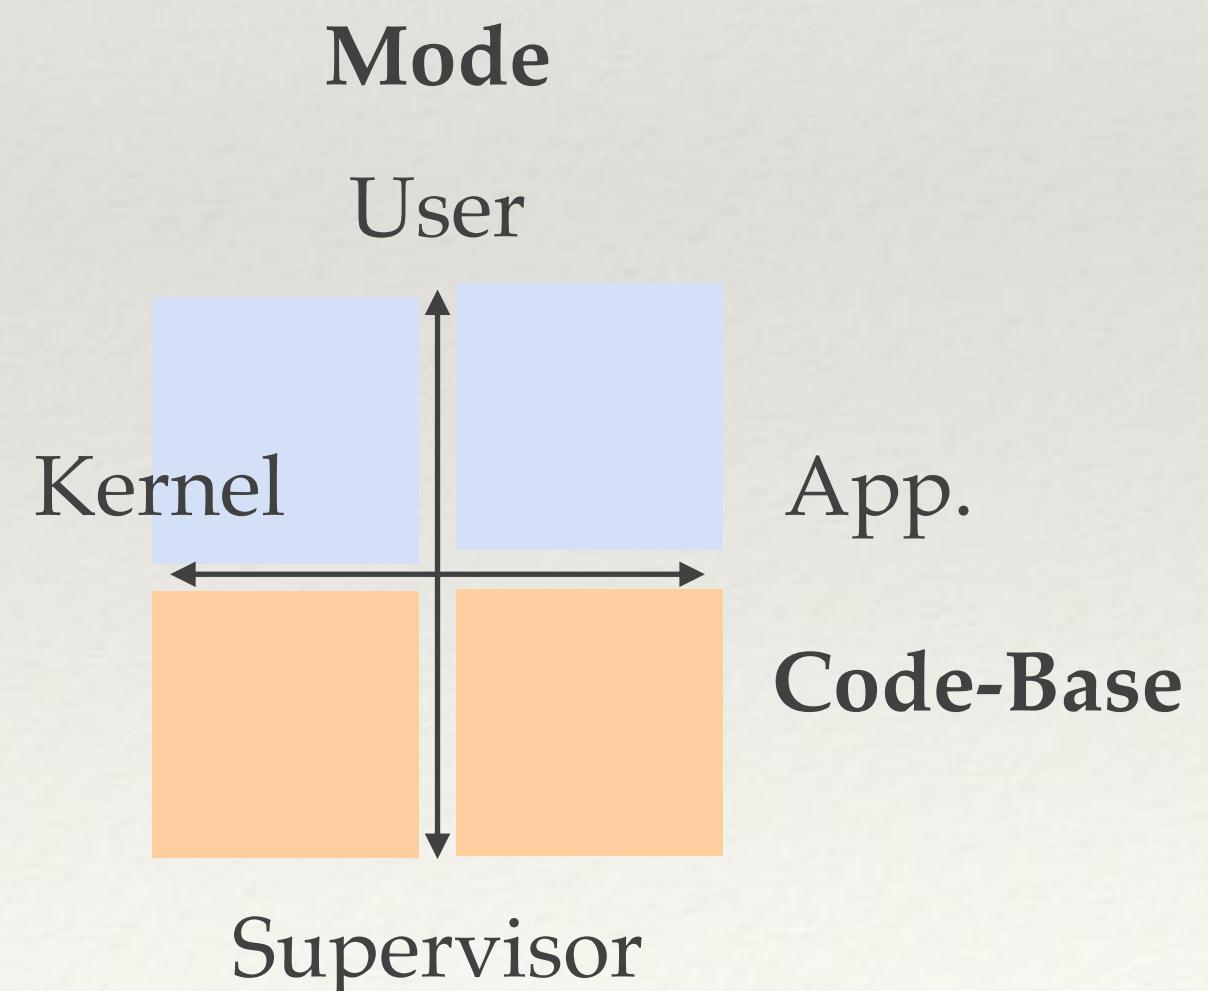
Mechanism

Dynamic Privilege



# Future Work

- ❖ Lowering kernel code
- ❖ Privileged Microkernel Servers
- ❖ Safe Privileged Processes
- ❖ Combination with capabilities



---

# Collaborators

---

- ❖ Mentors:
  - ❖ BU: Jonathan Appavoo / Orran Krieger / Renato Mancuso
  - ❖ RH: Larry Woodman / Uli Drepper / Richard Jones / Daniel Bristot de Oliveira
- ❖ Collaborators:
  - ❖ Arlo Albelli / Eric Munson / Ali Raza / Jim Cadden / Yara Awad
  - ❖ Amos Waterland / Han Dong / Parul Sohal / Dan Schatzberg
  - ❖ Albert Slepak / Ryan Sullivan / Ke Li / Ting Hsu

# Thesis Statement

We introduce Dynamic Privilege, an OS model enabling threads to adjust their hardware privilege levels on the fly, challenging the traditional separation between kernel and application.

Dynamic Privilege provides granular control and requires only minor modifications to support it in an existing OS.

Further, Dynamic Privilege enables new system evolution and optimization while maintaining compatibility with existing software ecosystems.

